



Network Security

Aggelos Kiayias




Networks...

- Computers connected to each other.
- Each machine has a unique address.
- Each message from the sender to receiver may stop at many intermediate hops till it reaches its destination. ...routing...



Networks

A network may extend over various physical mediums





Communication Media

- Wire (copper wire: *cheap, slow*) 10 Mbps, ~100m.
Carries electrical signal.
- Coax Cable (wire+insulation jacket) 100Mbps ~500m.
- Optical fiber (thin strand of glass). Carries pulses of light. 1000Mbps. ~4km.
- Radio signals. *Wireless*. good for short connections.
- microwave, infrared, satellite etc...



OSI Model

- Open Systems Interconnection Model

7	Application	User-level data
6	Presentation	data format (ascii etc.)
5	Session	sequencing
4	Transport	Flow control (acks, retransmissions errors)
3	Network	Routing (where to send)
2	Data Link	Local delivery
1	Physical	bit level representation



Example : e-mail

7	Application	e-mail composition
6	Presentation	text based transliteration, compression
5	Session	-
4	Transport	error-correcting codes, logical connection
3	Network	chop in packets - put addresses
2	Data Link	chop in frames - add MAC addresses
1	Physical	chop in bits - transmit



TCP/IP

- Transmission Control Protocol/ Internet Protocol.
- Four layers:
 - Application.
 - Host-to-Host Transport.
 - Internet.
 - Physical



TCP/IP

Application Layer	Prepare messages from user	Addressing/ Interaction
Transport Layer (e.g., TCP)	Packets are made	Sequencing, Reliability Error Correction
Network Layer (IP)	Into Datagrams	Routing
Data Link Layer	Connection between adjacent hosts - Bits	
Physical	Bit representation	



Data Link Layer Frames

- Source and destination Physical Addresses
- Encoding of bits
- Physical layer aspects (e.g., modulation).



IP Datagrams

- Contain *time to live* information (# of hops).
- Source and Destination IP addresses.
- Information about the encapsulated protocol.



TCP Packets

- Source / Destination ports.
- Acknowledgment number for connecting packets of a session.
- Sequence numbers.
- Integrity (checksums).



Application Data

- Depending on the application layer protocol used.



Example

Physical Layer: eth2
the 2 MAC addresses
+ IP indication

Network layer: IP
IP addresses, TTL,
checksum, fragmentation

0000	00	0f	db	4d	77	95	00	0d	93	b0	a3	24	08	00	45	00
0010	01	75	c8	de	40	00	40	06	44	dd	c0	a8	01	2e	40	ec
0020	29	05	e6	10	00	50	15	86	10	4d	25	b6	67	ed	80	18
0030	ff	ff	2d	2f	00	00	01	01	08	0a	2f	41	cf	64	62	38
0040	81	9a	47	45	54	20	2f	63	6e	6e	2f	32	30	30	36	2f
0050	55	53	2f	30	32	2f	32	37	2f	6b	61	74	72	69	6e	61
0060	2e	70	6f	6c	6c	2f	74	31	2e	32	31	33	35	2e	6d	6f
0070	6e	2e	62	65	61	64	73	2e	61	70	2e	6a	70	67	20	48
0080	54	54	50	2f	31	2e	31	0d	0a	41	63	63	65	70	74	3a
0090	20	2a	2f	2a	0d	0a	41	63	63	65	70	74	2d	4c	61	6e
00a0	67	75	61	67	65	3a	20	65	6e	0d	0a	41	63	63	65	70
00b0	74	2d	45	6e	63	6f	64	69	6e	67	3a	20	67	7a	69	70
00c0	2c	20	64	65	66	6c	61	74	65	0d	0a	52	65	66	65	72
00d0	65	72	3a	20	68	74	74	70	3a	2f	2f	77	77	77	2e	63
00e0	6e	6e	2e	63	6f	6d	2f	0d	0a	55	73	65	72	2d	41	67
00f0	65	6e	74	3a	20	4d	6f	7a	69	6c	6c	61	2f	35	2e	30
0100	20	28	4d	61	63	69	6e	74	6f	73	68	3b	20	55	3b	20
0110	50	50	43	20	4d	61	63	20	4f	53	20	58	3b	20	65	6e
0120	29	20	41	70	70	6c	65	57	65	62	4b	69	74	2f	34	31
0130	37	2e	39	20	28	4b	48	54	4d	4c	2c	20	6c	69	6b	65
0140	20	47	65	63	6b	6f	29	20	53	61	66	61	72	69	2f	34
0150	31	37	2e	38	0d	0a	43	6f	6e	6e	65	63	74	69	6f	6e
0160	3a	20	6b	65	65	70	2d	61	6c	69	76	65	0d	0a	48	6f
0170	73	74	3a	20	69	2e	61	2e	63	6e	6e	2e	6e	65	74	0d
0180	0a	0d	0a													

...Mw.....\$.E.
.u..@.@.D.....@.
)....P...M%.q...
..-/....../A.db8
..GET /cnn/2006/
US/02/27/katrina
.poll/t1.2135.mo
n.beads.ap.jpg H
TTP/1.1..Accept:
/..Accept-Lan
guage: en..Accep
t-Encoding: gzip
, deflate..Refer
er: http://www.c
nn.com/..User-Ag
ent: Mozilla/5.0
(Macintosh; U;
PPC Mac OS X; en
) AppleWebKit/41
7.9 (KHTML, like
Gecko) Safari/4
17.8..Connection
: keep-alive..Ho
st: i.a.cnn.net.
...

Transport Layer: TCP
Ports, Seq Ack numbers,
checksum, timestamps

Application Layer: HTTP
Request: GET
Request URI
Referrer
User-agent info
Connection info



Internet Protocols

- Data link Layer: ethernet, wi-fi etc.
- Network Layer: ICMP (Internet control message), IP etc.
- Transport Layer: UDP (user datagram protocol), TCP etc.
- Application Layer: Finger, FTP (file transfer), HTTP (hypertext transfer), IMAP (internet message access), IRC (internet relay chat), POP (post office), SMTP (simple mail transfer), TELNET (terminal emulation), X-window, etc.



UDP Protocol

- user datagram protocol.
- lightweight alternative to TCP.
- Faster, lighter - adds 8 bytes for control.
- stateless sending and no ordering.
- used for application layer protocols as SNMP (simple network monitoring), Syslog (system audit log), Time etc.



How do data find their way?

- *Application.* HTTP request to www.website.com (DNS resolution)
- *Transport.* which in turn will result to some packets directed to a certain **port #**.
- *Internet.* which in turn will result to some frames directed to a **IP address**.
- *Physical.* which in turn will result to some actual bits being sent to **MAC address**



Addressing

- Two mappings are necessary:
 - From host name to IP address.
 - From IP address to MAC address.
- Host name will be mapped to an IP address through a protocol called DNS
- MAC address will be obtained from an address resolution table.



Transmitting a packet

- A packet needs to be directed to a certain IP address.
- To figure out where to send it next, a routing table is consulted. Example:

Routing table:

Destination	Gateway	Flags	Refs	Use	Netif	Expire
default	137.99.11.1	UGSc	16	1863	en0	
...						



Transmitting a packet, II

- Once the (intermediate hop) IP address is determined the packet must be split into frames and directed to the right MAC address.
- Internet to Ethernet address translation:

Address	HWtype	HWaddress	Flags	Mask	Iface
137.99.11.1	ether	00:0B:46:9A:1B:3F	C		eth0



Receiving a packet

- Keep it or forward it.
- Based on destination address (and perhaps other parameters).
- Forward it using the routing table as before.
- Routing and Address Resolution tables are dynamically updated.



IP Addresses and DNS

- 32-bit (IPv6 will offer 128-bits).
- IP address correspond to names according to the Domain Name Service (DNS).
- Given a certain name at the application layer a query will be transmitted to a *Name Server* to resolve it for the corresponding IP address.



Trace route

```
1  192.168.63.11 (192.168.63.11)  3.896 ms  2.122 ms  1.511 ms
2  195.134.67.1 (195.134.67.1)  1.794 ms  2.839 ms  1.784 ms
3  grnetRouter.L1.uoa.athens-3.access-link.grnet.gr (194.177.209.97)  1.984 ms  2.262 ms  3.360 ms
4  eie2-to-koletti1.backbone.grnet.gr (195.251.27.46)  2.196 ms  4.345 ms  3.539 ms
5  core1.ams.net.google.com (195.69.144.247)  72.084 ms  72.003 ms  72.743 ms
6  209.85.248.88 (209.85.248.88)  74.182 ms *  74.916 ms
7  64.233.175.246 (64.233.175.246)  75.377 ms  75.800 ms  74.950 ms
8  209.85.255.143 (209.85.255.143)  77.161 ms  72.14.239.197 (72.14.239.197)  79.314 ms  209.85.255.166
   (209.85.255.166)  90.792 ms
9  72.14.232.37 (72.14.232.37)  81.473 ms  72.14.232.41 (72.14.232.41)  84.950 ms  72.14.232.37
   (72.14.232.37)  84.960 ms
10 ez-in-f104.1e100.net (66.102.13.104)  79.149 ms  78.015 ms  76.316 ms
```



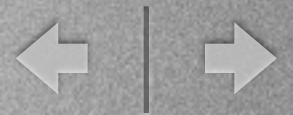

How traceroute works

- Using a special IP header field called TTL : time to live.
- TTL = number of hops a packet is allowed to make. Each router decreases by one.
- When TTL reaches 0 then a router discards the packet and notifies originator.
- For traceroute: send repeatedly packets and calibrate TTL as 1, 2, 3, 4, 5, ...
- not all routers necessarily respond (* * *)



Client Server Model

- Application protocols, FTP, HTTP, Telnet etc.
- Server listens to port for client requests.
- Client initiates protocol



Talking over TCP/IP

Client



state:closed

SYN

state:listen

state:syn-sent

SYN+ACK

state:syn-received

state:established

ACK

state:established

Application layer data

ACK

Application layer data

ACK

...

FIN+ACK

FIN+ACK

Server





Packet Sniffing

- Every computer in the Internet sends and receives packets.
- Anyone with the appropriate privileges in a certain host can “sniff” the packets that are being forwarded by the host (and not only the packets that are directed to the host).
- In this case the host is said to be in **promiscuous** mode.




Packet Sniffing, II

- If you have privileges for promiscuous mode then you can capture all traffic within the sub-network the host belongs to.
- You cannot capture traffic outside your sub-network (e.g., traffic that is not directed towards your gateway).
- When you use your computer do always ponder what is your sub-network (consider: sitting at a cafe connected to a wireless access point)



Wireshark

- Wireshark is a powerful “network protocol analyzer”
- Not only it sniffs data when put on promiscuous mode but also “knows” the protocols and structures the packets in the appropriate format.



Wireshark, II

screen dump of capture window after an FTP connection

No.	Time	Source	Destination	Protocol	Info
1	0.000000	192.168.1.46	192.168.1.1	DNS	Standard query A ftp.debian.org
2	0.156872	192.168.1.1	192.168.1.46	DNS	Standard query response A 128.101.240.212
3	0.203708	192.168.1.46	128.101.240.212	TCP	58408 > ftp [SYN] Seq=0 Ack=0 Win=65535 Len=0 MSS=1
4	0.311009	128.101.240.212	192.168.1.46	TCP	ftp > 58408 [SYN, ACK] Seq=0 Ack=1 Win=5792 Len=0 M
5	0.311128	192.168.1.46	128.101.240.212	TCP	58408 > ftp [ACK] Seq=1 Ack=1 Win=65535 Len=0 TSV=79
6	0.427572	128.101.240.212	192.168.1.46	FTP	Response: 220 saens.debian.org FTP server (vsftpd)
7	0.457218	192.168.1.46	128.101.240.212	TCP	58408 > ftp [ACK] Seq=1 Ack=43 Win=65535 Len=0 TSV=7
8	3.908879	192.168.1.46	128.101.240.212	FTP	Request: USER anonymous
9	3.995051	128.101.240.212	192.168.1.46	TCP	ftp > 58408 [ACK] Seq=43 Ack=17 Win=6144 Len=0 TSV=4
10	3.995621	128.101.240.212	192.168.1.46	FTP	Response: 331 Please specify the password.
11	4.058261	192.168.1.46	128.101.240.212	TCP	58408 > ftp [ACK] Seq=17 Ack=77 Win=65535 Len=0 TSV=
12	8.388059	192.168.1.46	128.101.240.212	FTP	Request: PASS ak@ak.org
13	8.473188	128.101.240.212	192.168.1.46	FTP	Response: 230-
14	8.473824	128.101.240.212	192.168.1.46	FTP	Response: 230-This site is just another one in a worldwid
15	8.659296	192.168.1.46	128.101.240.212	TCP	58408 > ftp [ACK] Seq=33 Ack=158 Win=65535 Len=0 TSV=
16	8.751453	128.101.240.212	192.168.1.46	FTP	Response: 230-It is not the "primary Debian FTP site" - it
17	8.758911	192.168.1.46	128.101.240.212	FTP	Request: SYST

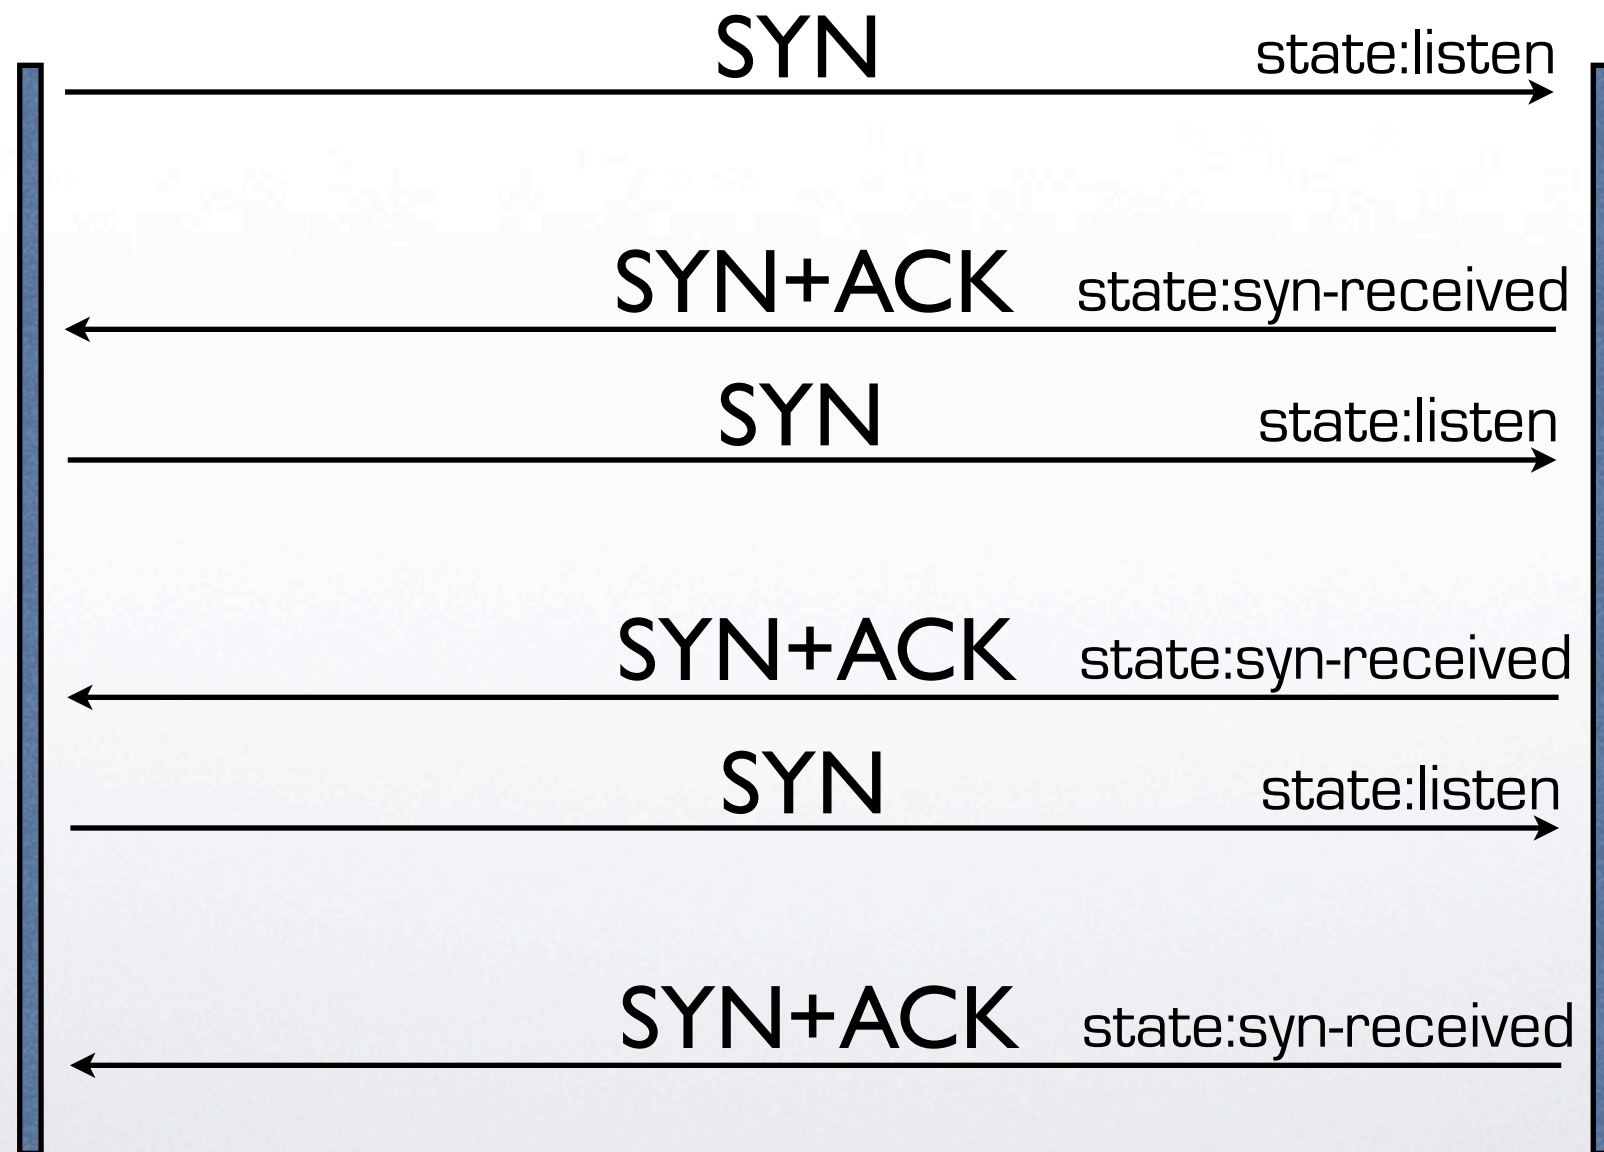
the three-way handshake

Observe the password
and username



A malicious client

Client




Server



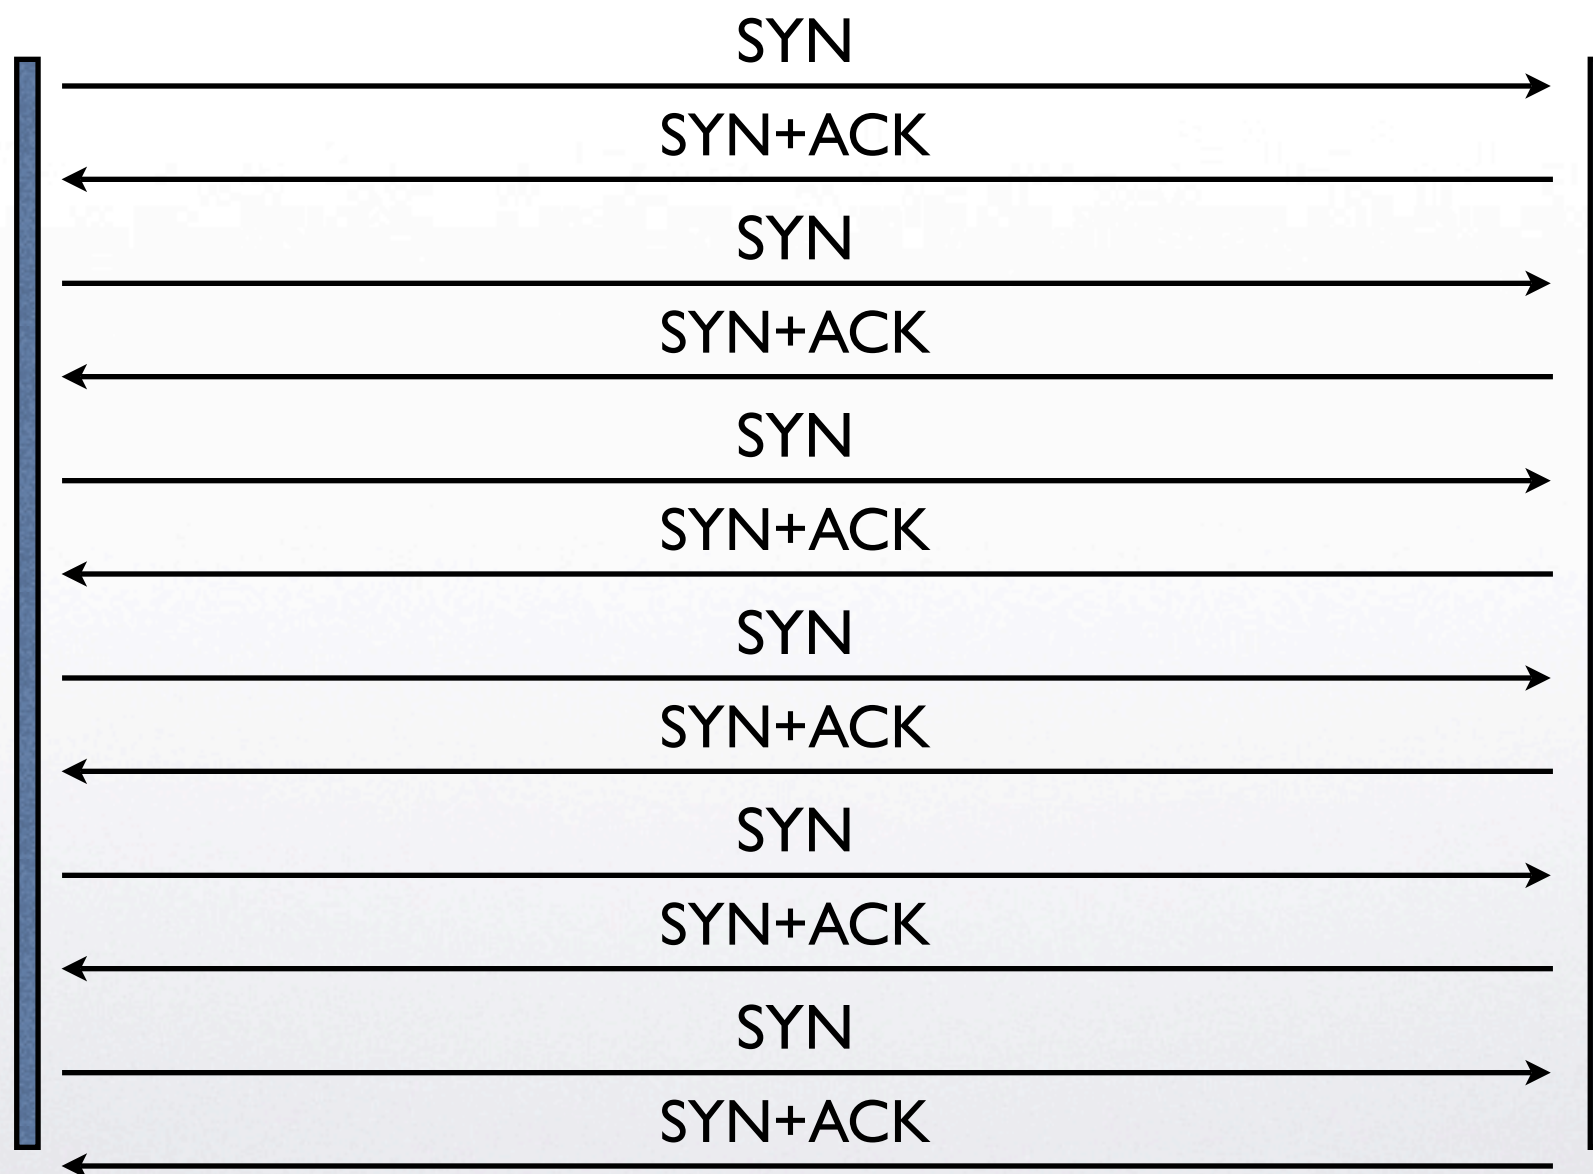


A malicious client, II

Client



IP
Spoofing



Server





SYN Flooding

- Client bombards Server with SYN packets that are never ACK'ed.
- IP spoofing can be used to make packets look their coming from other places.
- Physical limitation: *bandwidth*.
- If bandwidth on client is substantial compared to server there is serious potential for a *Denial of Service (DoS) Attack*



DoS Attacks

- Deplete / misconfigure / misallocate the resources on a target server host so that it cannot serve its clients.
- resources:
 - bandwidth.
 - memory.
 - cpu
 - ...



DoS By Flooding

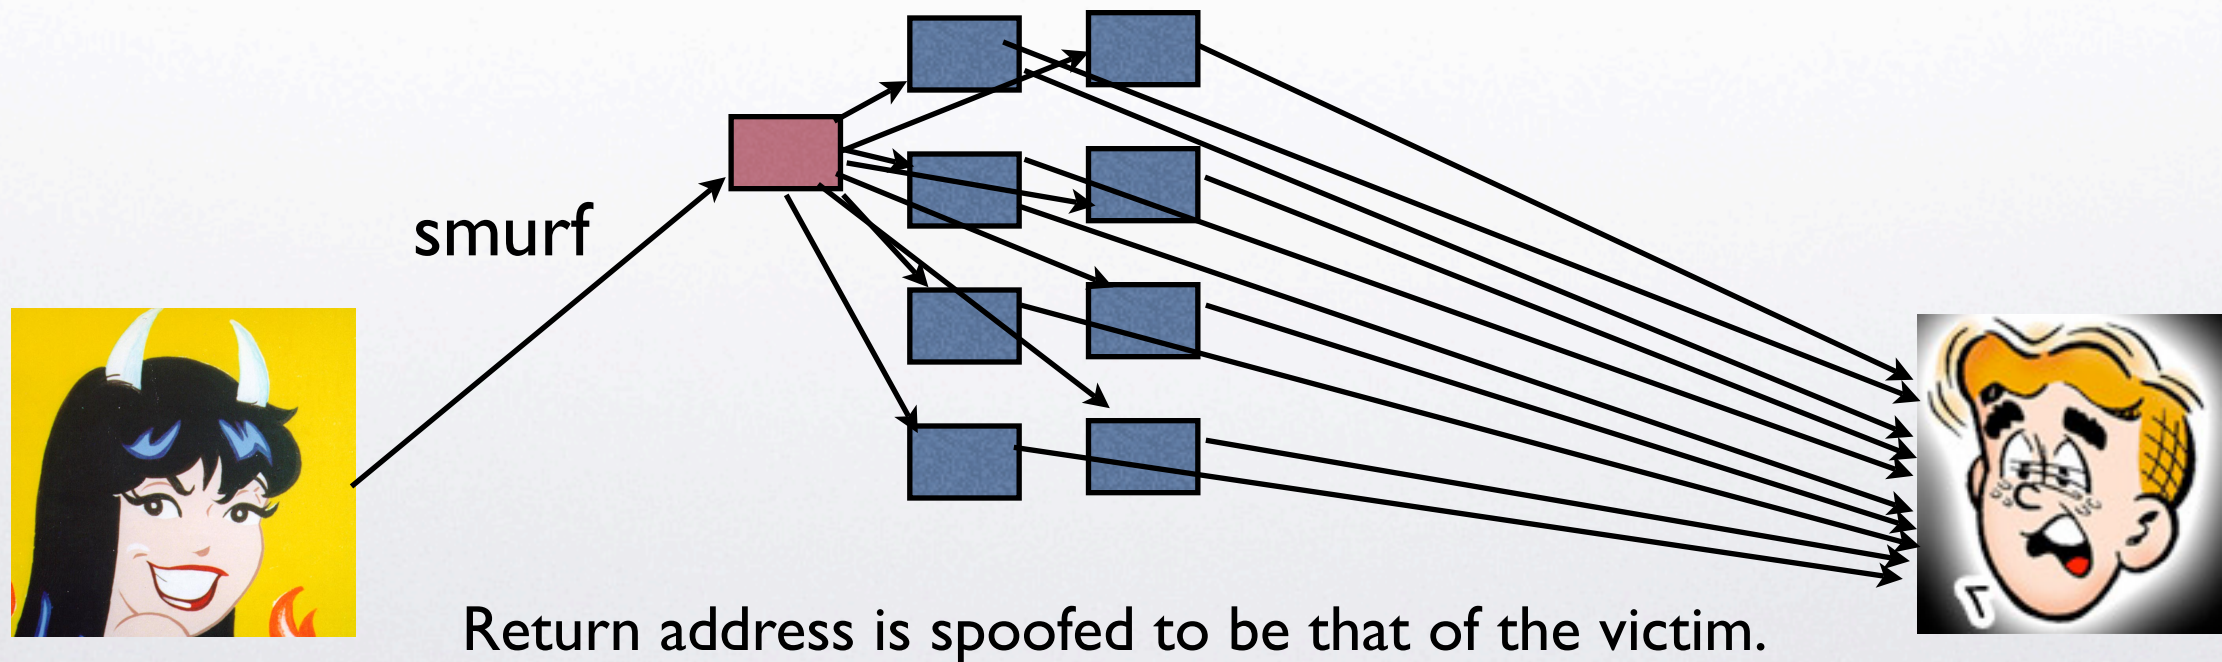
- SYN Flooding - we just saw it.
- Ping Floods: this is a flooding of ICMP Echo Request packets.

```
aggelos@grub:~$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1): 56 data bytes
64 bytes from 192.168.1.1: icmp_seq=0 ttl=64 time=0.686 ms
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=0.611 ms
64 bytes from 192.168.1.1: icmp_seq=2 ttl=64 time=0.617 ms
^C
--- 192.168.1.1 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max/stddev = 0.611/0.638/0.686/0.034 ms
```




Smurf Attack

- An enhanced Ping flood attack that utilizes IP Broadcast:
- it is possible to specify the destination IP address as a broadcast to all hosts in a subnetwork.






DoS by Crashing

- Any action of a client that may crash a server results in a DoS.
- Buffer overflows.
- Once upon a time: **Ping of Death** delivering a ping packet of size greater than 64Kb used to crash systems in good old times 1997-98.



Distributed DoS

- The real deal!





Zombification

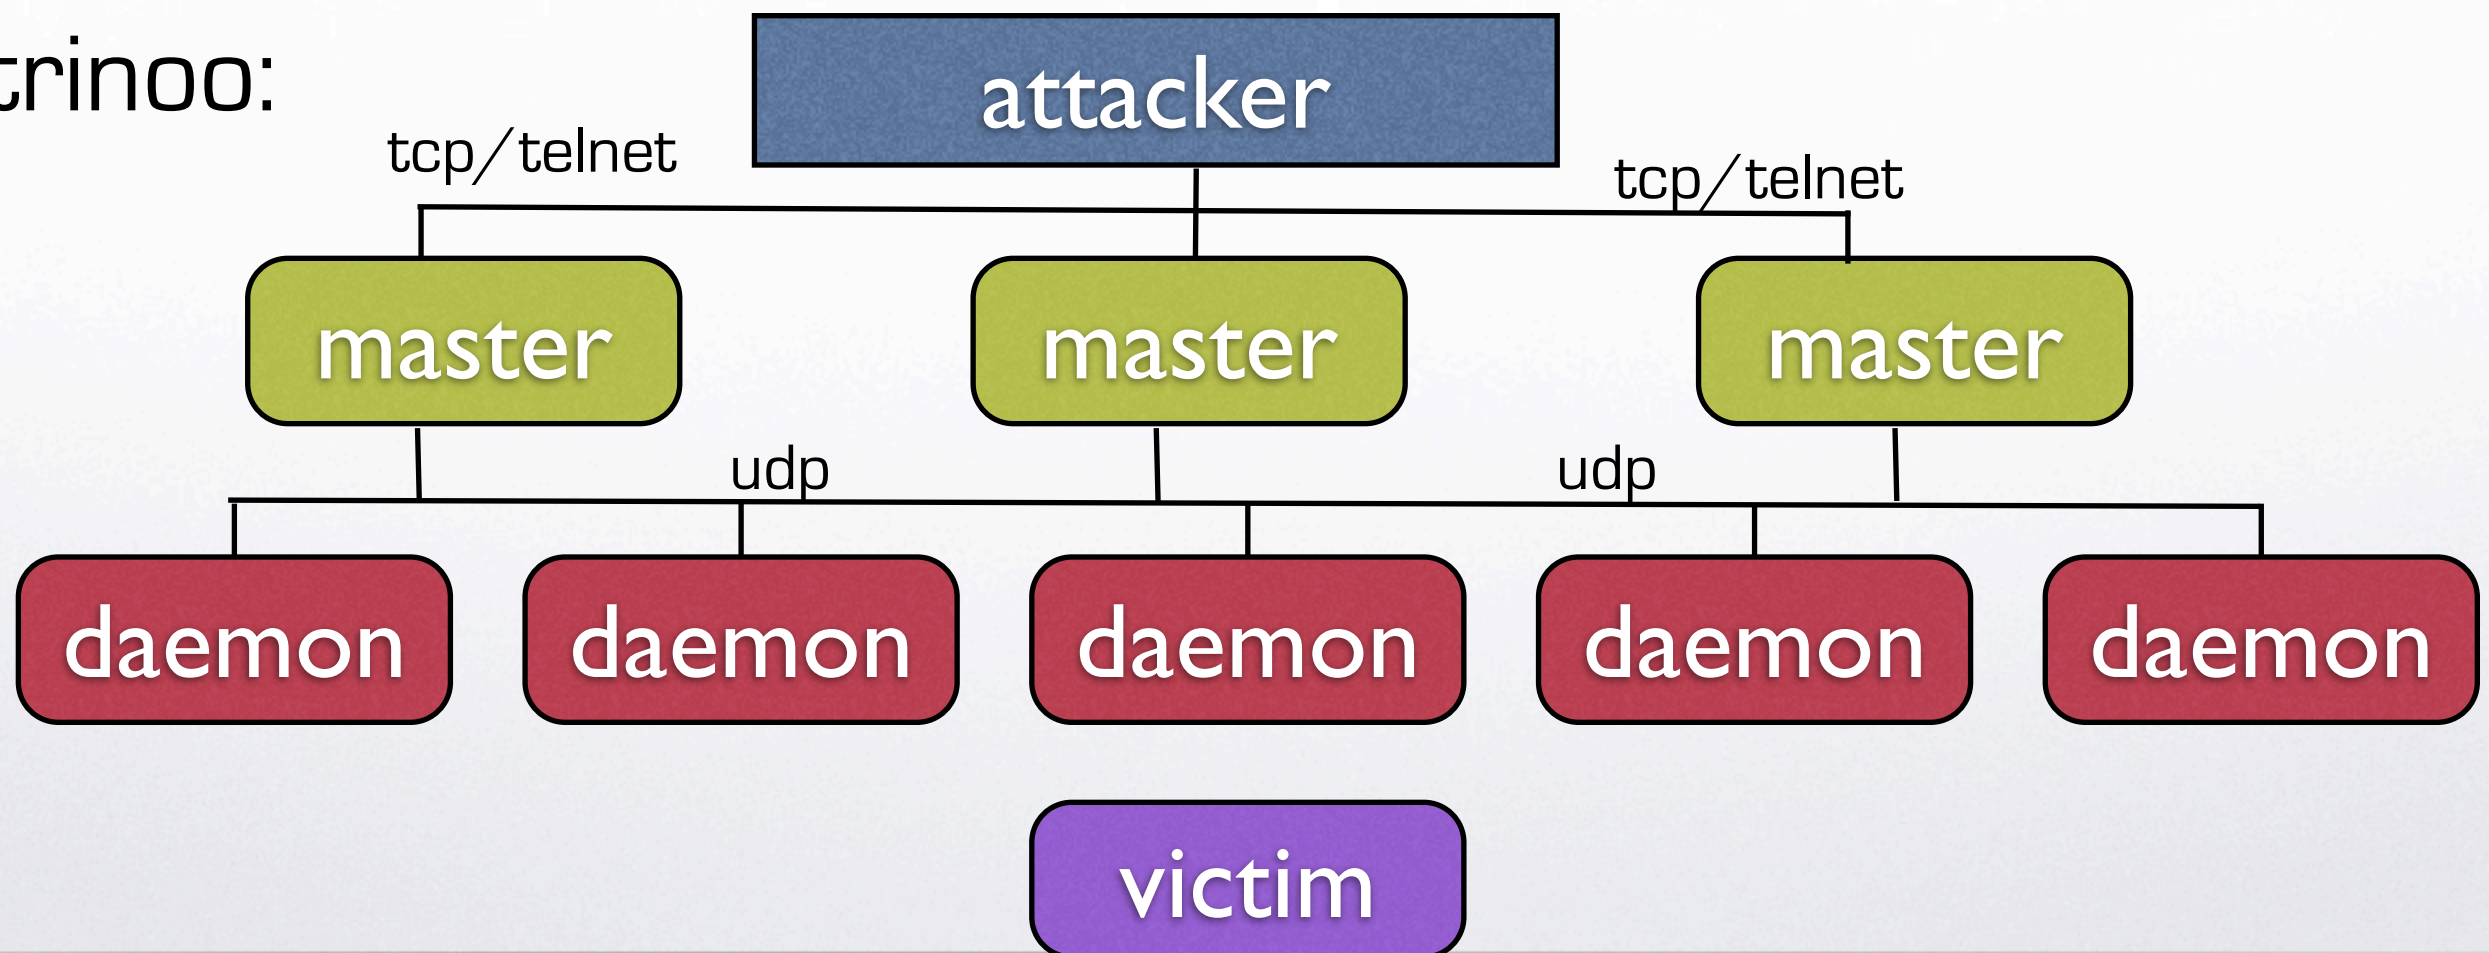
- Host compromised by virus, worm, trojan.
- Runs rootkit remote administration tool.
- When commanded it launches DoS attack against victim.
- Attack seems to be coming from everywhere!



Tools of the trade.

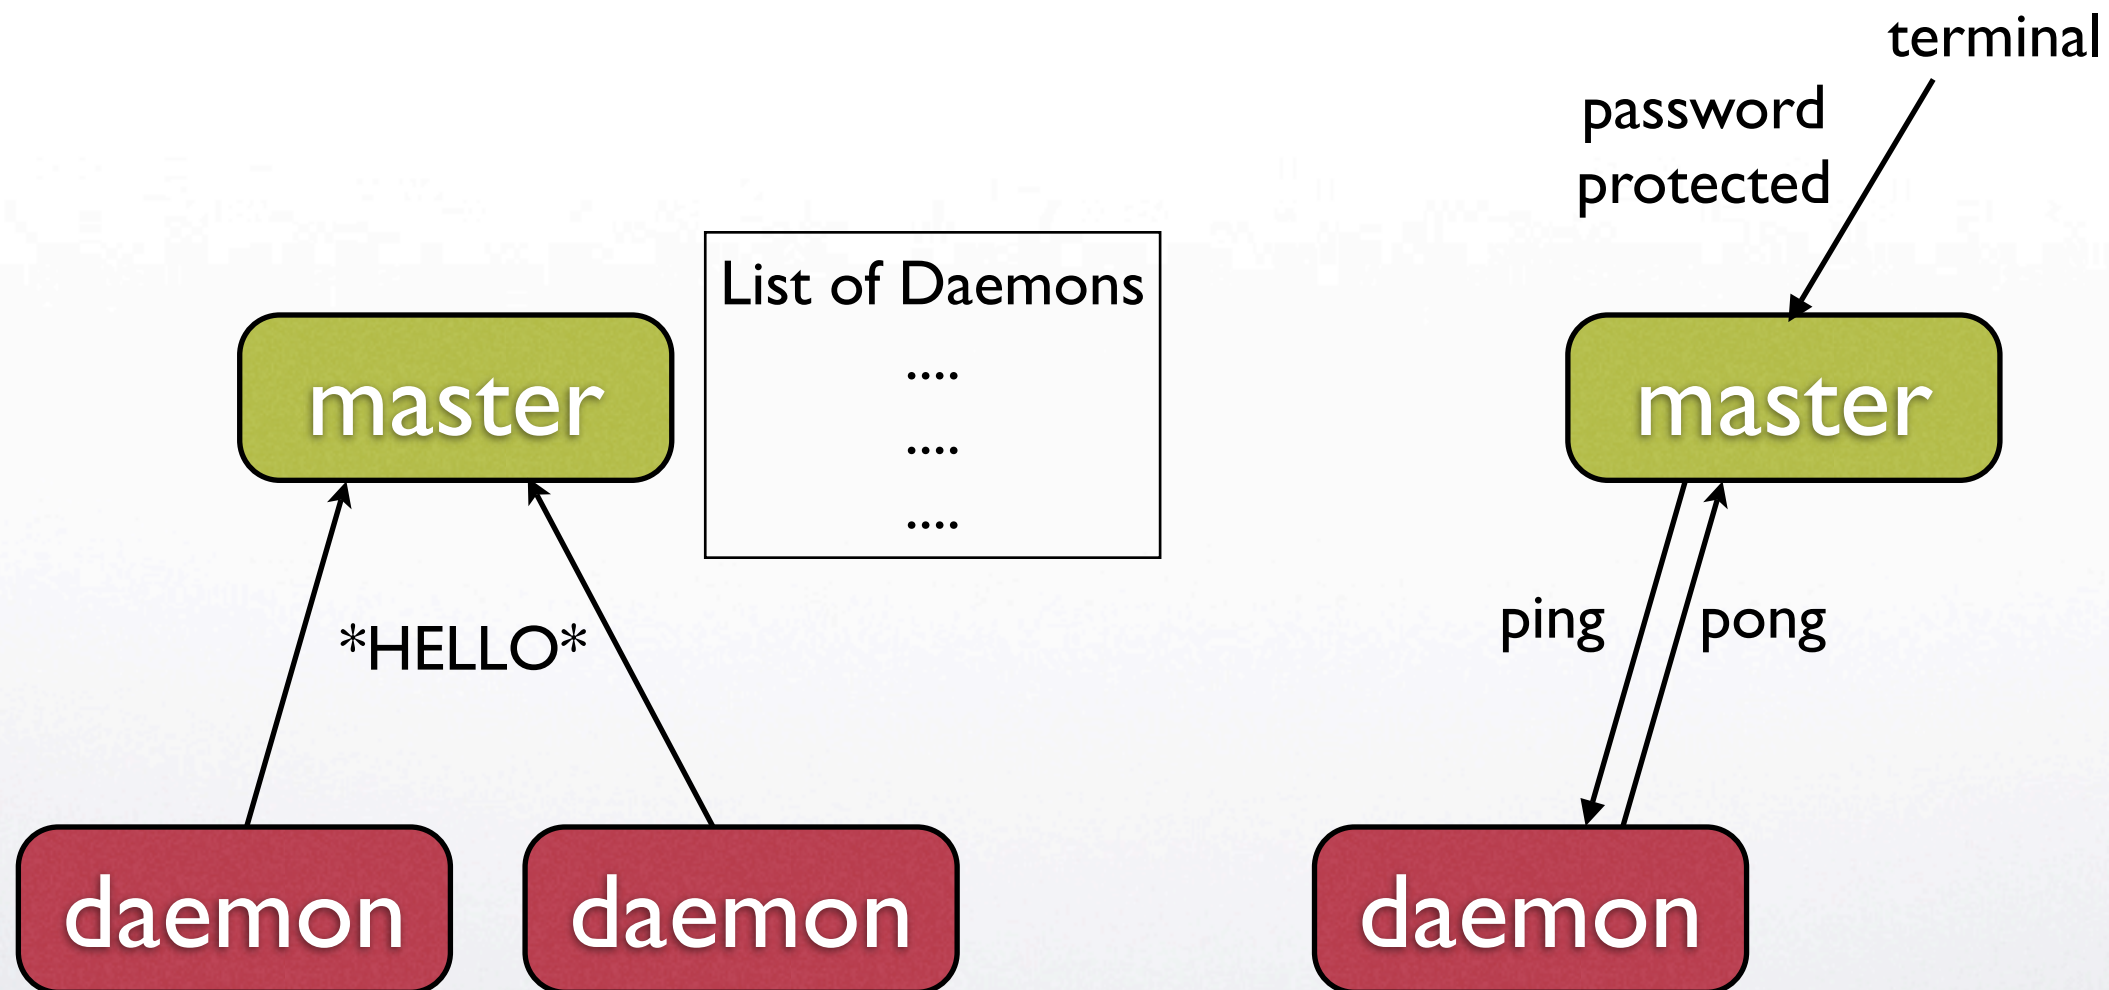
- trinoo (trin00)/ wintrinoo, TFN2K (tribal flood network 2K), Stacheldraht, etc.

trinoo:





trinoo networks



<http://staff.washington.edu/dittrich/misc/trinoo.analysis>



Logging in to a trinoo master

```
attacker$ telnet 10.0.0.1 27665
Trying 10.0.0.1
Connected to 10.0.0.1
Escape character is '^]'.
kwijibo
Connection closed by foreign host.
. . .
```

special port

```
attacker$ telnet 10.0.0.1 27665
Trying 10.0.0.1
Connected to 10.0.0.1
Escape character is '^]'.
betaalmostdone
trinoo v1.07d2+f3+c..[rpm8d/cb4Sx/ ]
```

correct password

```
trinoo>
```



trinoo commands

dos IP DoS the IP address specified. A command ("aaa 144adsl IP") is sent to each Bcast host (i.e., trinoo daemons) telling them to DoS the specified IP address.

mdos <ip1:ip2:ip3>

Multiple DoS. Sends a multiple DoS command ("xyz 144adsl 123:ip1:ip2:ip3") to each Bcast host.

bcast List all active Bcast hosts.

mdie pass Disable all Bcast hosts, if the correct password is specified. A command is sent ("dle 144adsl") to each Bcast host telling them to shut down. A separate password is required for this command.



DDoS Tools

- TFN2k:
 - like trinoo but with more attack capabilities. spoofed IP addresses.
- Stacheldraht
 - like trinoo and TFN2k but more capabilities, encrypted communication.



Command & Control

- Hard code address of master process into zombie code (obvious disadvantage)
- **Internet Relay Chat (IRC)**. An IRC server hosts a number of channels.
- a bot can incorporate a stripped IRC client to login to a certain channel.
- Once logged on, wait for commands.



Botnet statistics

- 1000 Bots with 128KBit/s upstream > 100 MBits/s [for such kinds of bandwidth a company may pay >>\$10K per month]
- Several large botnets of > 50,000 bots have been observed in the wild. Recently even larger sizes.
- November 2004 - January 2005, specific study (see below) observed 226 DDoS-attacks against 99 unique targets.
- Some botnet operators are very unskilled (destroying their botnets by simple mistakes, lacking basic programming skills etc.)


<http://www.honeynet.org/papers/bots/>



In the news...

- January 24, 2006:
 - J.J. Ancheta, 20 yr. pleads guilty to charges of conspiracy, damaging computers used by the US government and fraud.
 - faces max. sentence of 25 years.
 - controlled > 100,000 bots, used for various illegal activities.
 - As part of the plea he forfeits the proceeds of his illicit activities that include \$60,000 cash, a BMW, computer equipment.

http://news.zdnet.com/2100-1009_22-6030270.html



In the news...

03-18-05

NEWARK, N.J. - A Michigan man was arrested today on a federal charge that he hired a New Jersey juvenile to conduct highly destructive computer attacks on competitors of his online sportswear business, including a web-based New Jersey company.

...

According to the complaint, the computer attacks were conducted by the New Jersey juvenile from the juvenile's home computer. The juvenile secretly infected thousands of computers with copies of a computer program known as a "bot" (short for "robot"). As described in the complaint, a "bot" can have legitimate functions, but can also be used to gain unauthorized access to and control over computers that they infect, and can thus cause the infected computers to attack other computers. "Bots" used for such illicit purposes are frequently disguised as MP3 music files or photographs that unsuspecting computer users download from public Internet sites. Having downloaded an infected file, a computer user is usually unaware of the presence of a "bot" on his or her computer.

http://www.usdoj.gov/usao/nj/publicaffairs/NJ_Press/files/arab0318_r.htm

11-10-2004

...

You have 2 choices," Card Services International was told via e-mail earlier this year. "You can ignore this email and try to keep your site up, which will cost you tens of thousands of dollars ... or you can send us \$10K by Western Union to make sure your site experiences no problem. If you choose not to pay for our help, then you will probably not be in business much longer, as you will be under attack each weekend for the next 20 weeks."

It wasn't a bluff. The Kentucky-based credit card processing firm suffered about a week's worth of outages before blocking the attack, according to president Jay Broder. The firm didn't pay, and the FBI is investigating the incident.

<http://www.msnbc.msn.com/id/6436834/>



Botnet Uses Beyond DDoS

- The convenience of separation between ‘hacking a computer’ and ‘committing a crime.’
- Sending Spam.
- Click-Fraud.
- Identity Theft.
- Stealing files: e.g., game Diablo-2 items were stolen and sold to EBay..

we will return to those later



Other DoS attacks

- **The Fork Bomb:** any program that constantly forks by creating child processes that do the same.
- Any time a process is called the O/S allocates memory for the process's requirements and enters the process specifics in a data structure.

Example:

```
int main(void) {  
    while(1) {  
        fork();  
    }  
    return 0;  
}
```

this little program may
crash your linux-box

You can write fork-bombs
for all major languages.



Defending against DOS, I

- Attacks like Smurf, Fork Bombs etc. result from the ability of an entity to allocate more resources than necessary in normal operation.
- Restricting such capability will thwart the related attack, e.g.,
 - restrict the use of IP broadcasting.
 - restrict the number of processes a user may create.



Defending against DOS, II

- Being proactive also helps:
 - A network router may detect outbound traffic that serves a DoS attack (e.g., outbound traffic with randomly spoofed source IP address. Such traffic can be dropped before leaving the network).



Defending against DDoS

- Harder since an attack may look as quite legitimate traffic: (e.g., HTTP)
- Challenge: distinguish between good traffic and DDoS traffic.
- We will examine some general approaches for DDoS defense next.



Simple Measures

- Try to use various parameters to make a system understand it is being attacked:
 - counting number of half-open connections (syn-rcvd).
 - counting number of connections refused.
- Once some **red-flags** are observed the system may shorten the time that it keeps half-open connections active.



SYN Cookies

Client



state:closed

SYN

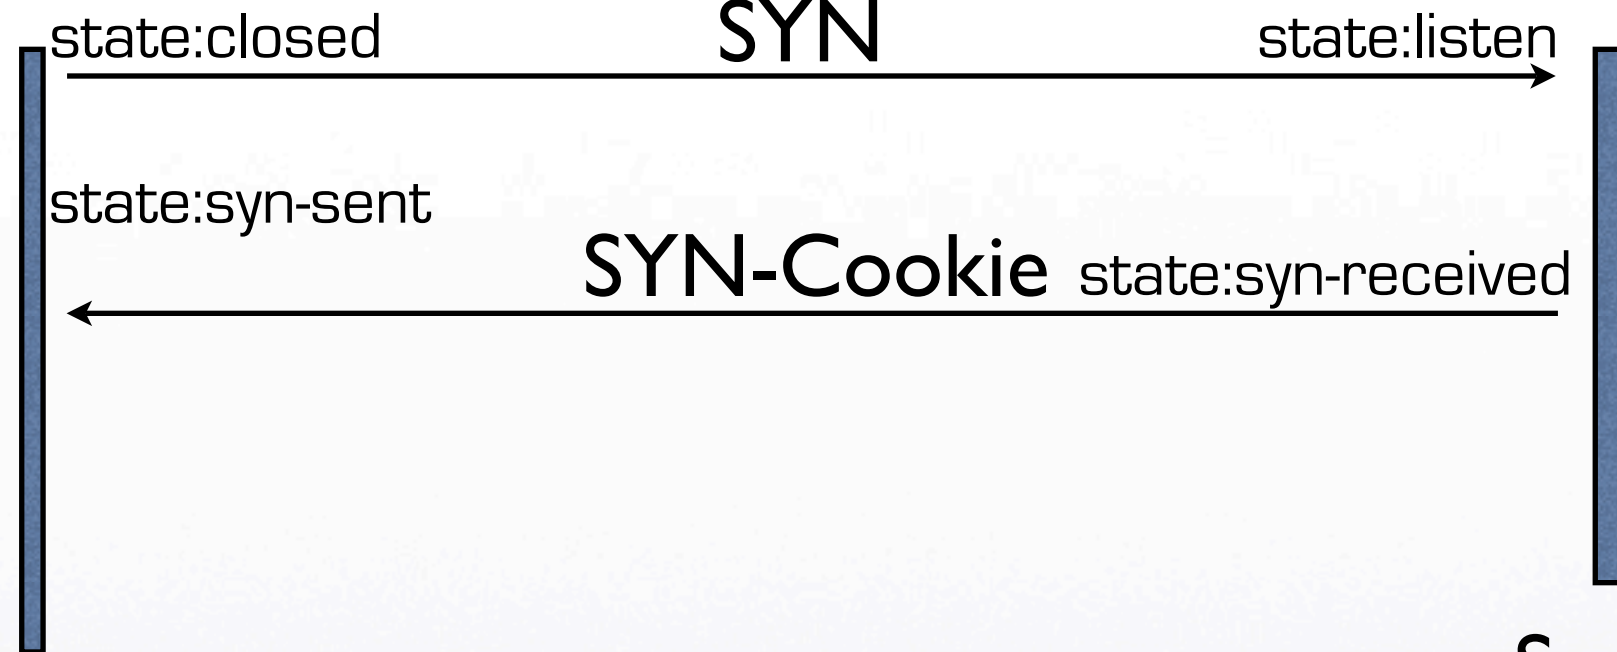
state:listen

state:syn-sent

SYN-Cookie

state:syn-received

Server



A SYN-Cookie contains info derived by the source-address port, destination address and port etc.

Server can forget everything about the client till the ACK message that will include the SYN-Cookie

implemented in BSD



Proofs of Work

- Client is requested to solve a puzzle in order to establish a connection.
- **Pros:** works in any client-server system. can be tuned to slow down malicious systems.
- **Cons:** needs changes in the basic infrastructure (e.g., changes in both client and server). Will not eliminate bandwidth emaciation.




Example.

From Juels/Brainard (RSA)

Hash function

$$\mathcal{H} : \{0, 1\}^* \rightarrow \{0, 1\}^k$$

secret / time / request


$$\mathcal{H} \longrightarrow x \in \{0, 1\}^k \xrightarrow{\text{trunc}} x' \in \{0, 1\}^{k-w}$$

Server to client: sends x' , $y = \mathcal{H}(x)$ stores (x, y)

Client needs to respond with $x'' : \mathcal{H}(x' || x'') = y$

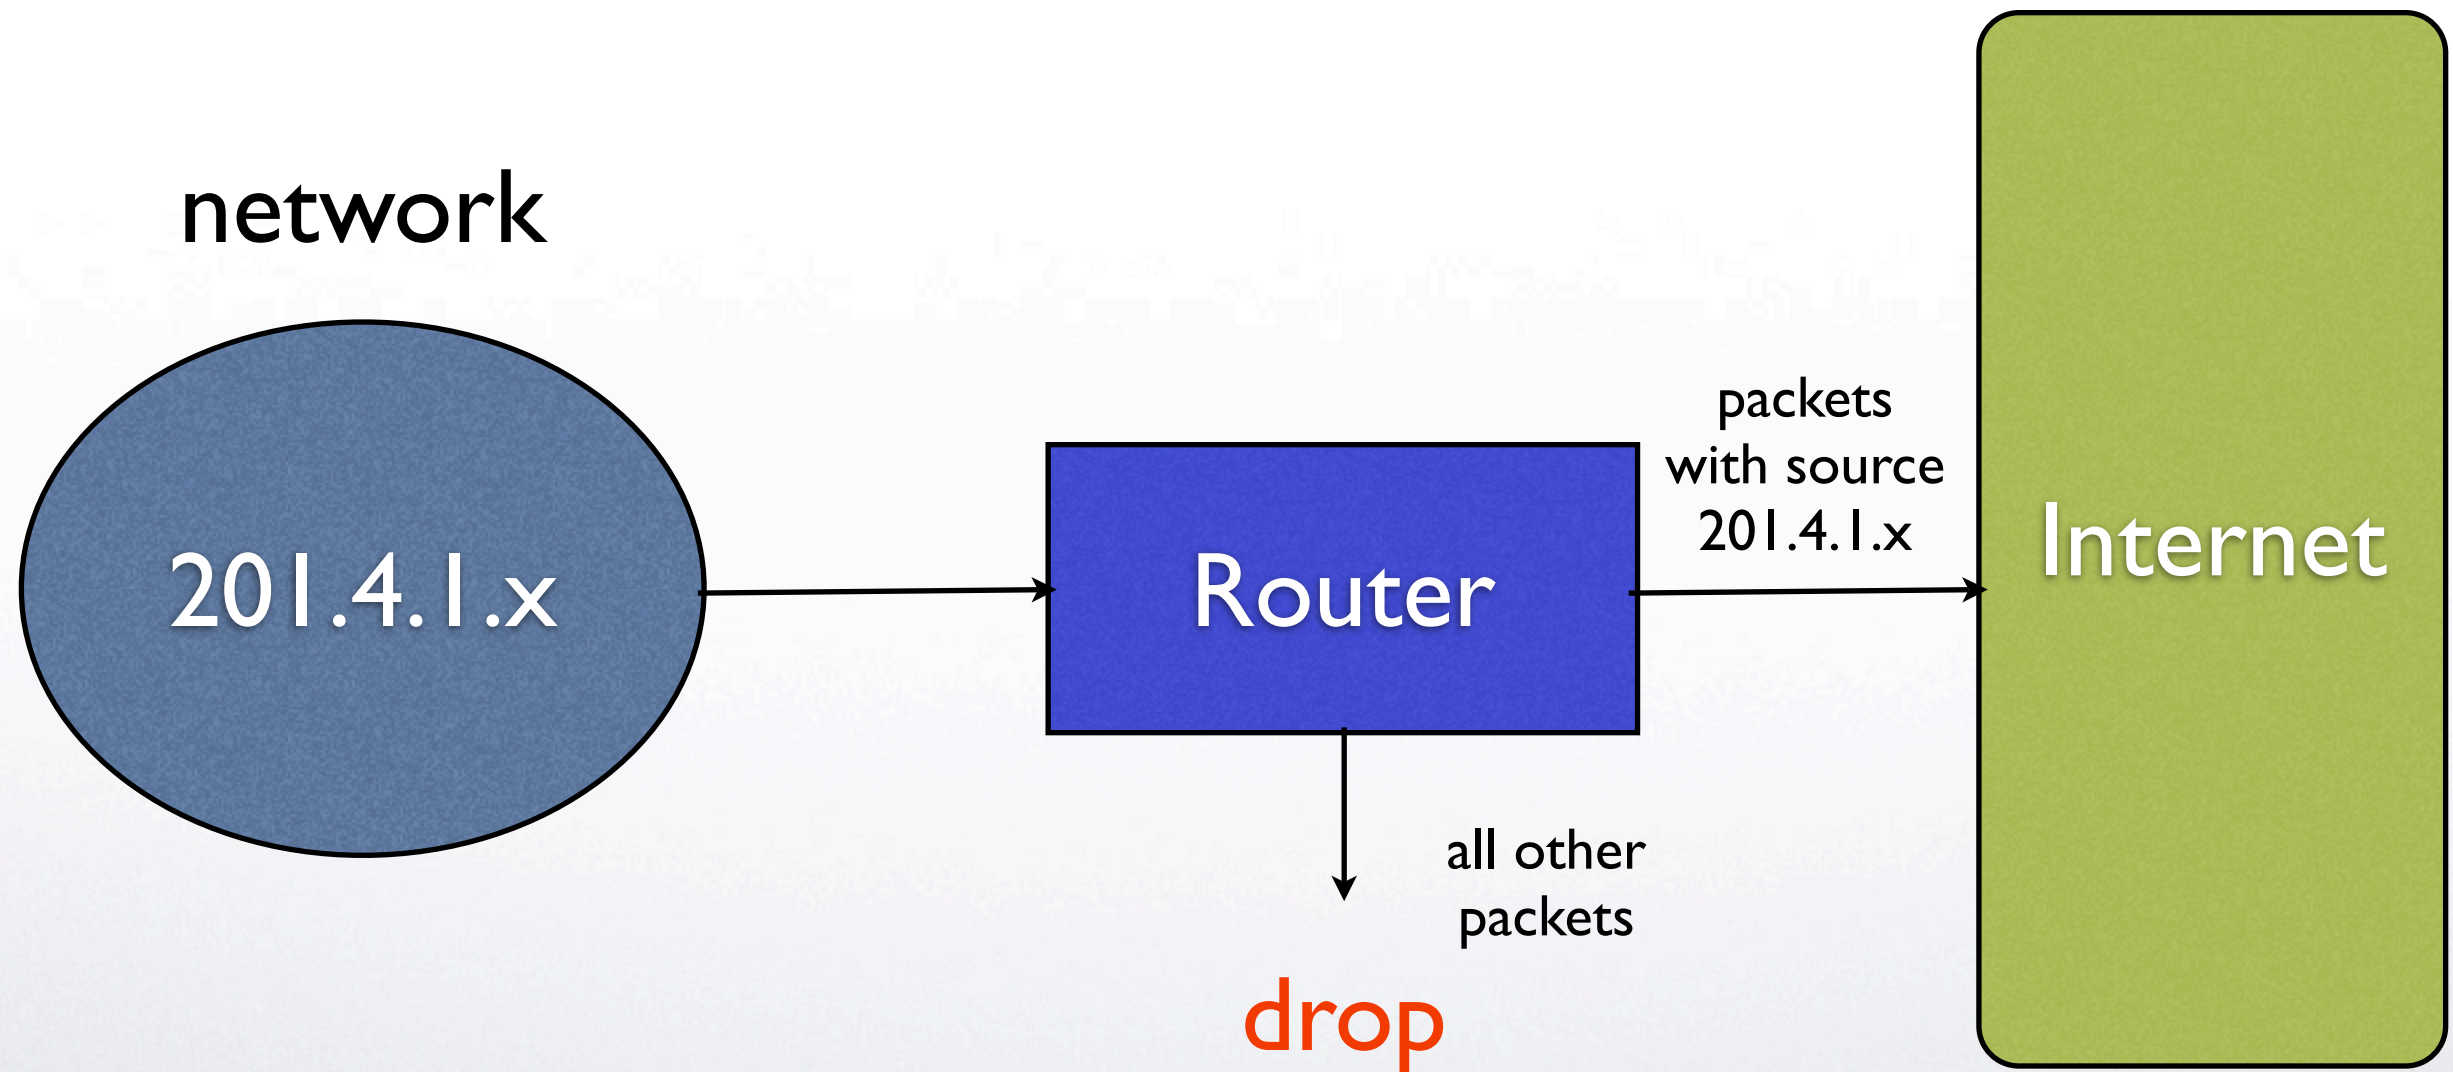


IP Throttling

- Any IP address that issues a big upstream will have its incoming packets being dropped at a certain rate.
- **Advantage:** can be configured in your router. No need to modify client/server.
- **Disadvantage:** what are the right settings?



Ingress Filtering



<http://www.faqs.org/rfcs/rfc2827.html>



IP Traceback

- A “forensic” technique:
 - try to identify the ‘real’ source of a packet.
 - Have routers mark packets with their IP address. => problem: not enough space in a packet.



Edge Marking.

- With some low probability q a router **marks** a packet. **Marking**:
 - If packet unmarked: enter your IP address as the start IP address.
 - If **distance** = 0: enter your IP address as the end IP address.
 - **Otherwise** increase the **distance**.



Encoding IP addresses

- Use only the 16 bits of the IP identification field used for fragmentation.

- How to pack 64+5 bit information into 16 bits?

$$\mathcal{H}_1, \mathcal{H}_2 : \{0, 1\}^{32} \rightarrow \{0, 1\}^{11}$$

- Use

first router	second router
$\mathcal{H}_1(\text{startIP})$	$\mathcal{H}_1(\text{startIP}) \oplus \mathcal{H}_2(\text{endIP})$



Reconstruct Path

