



Network Attacks

Aggelos Kiayias




Networks...

- Computers connected to each other.
- Each machine has a unique address.
- Each message from the sender to receiver may stop at many intermediate hops till it reaches its destination. (networks are not complete graphs ...routing...)



Networks

A network may extend over various physical mediums





Communication Media

- Wire (copper wire: *cheap, slow*) 10 Mbps, ~100m.
Carries electrical signal.
- Coax Cable (wire+insulation jacket) 100Mbps ~500m.
- Optical fiber (thin strand of glass). Carries pulses of light. 1000Mbps. ~4km.
- Wireless: WiFi, Radio signals.
- infrared, satellite etc...



OSI Model

- Open Systems Interconnection Model

7	Application	User-level data
6	Presentation	Data format (ascii etc.)
5	Session	Sequencing
4	Transport	Flow control (acks, retransmissions errors)
3	Network	Routing (where to send)
2	Data Link	Local delivery
1	Physical	Bit level representation



Example : e-mail

7	Application	e-mail composition
6	Presentation	text based transliteration, compression
5	Session	-
4	Transport	error-correcting codes, logical connection
3	Network	chop in packets - put addresses
2	Data Link	chop in frames - add MAC addresses
1	Physical	chop in bits - transmit



TCP/IP

- Transmission Control Protocol/ Internet Protocol.
- Four layers:
 - Application.
 - Host-to-Host Transport.
 - Internet.
 - Physical.



TCP/IP

Application Layer	Prepare messages from user	Addressing/ Interaction
Transport Layer (e.g., TCP)	Packets are made	Sequencing, Reliability Error Correction
Network Layer (IP)	Into Datagrams	Routing
Data Link Layer	Connection between adjacent hosts - Bits	
Physical	Bit representation	e.g. radio modulation



Data Link Layer Frames

- Source and destination Physical Addresses
- Encoding of bits
- Physical layer aspects (e.g., modulation).



IP Datagrams

- Contain *time to live (TTL)* information (# of hops).
- Source and Destination IP addresses.
- Information about the encapsulated protocol.



TCP Packets

- Source / Destination ports.
- Acknowledgment number for connecting packets of a session.
- Sequence numbers.
- Integrity (checksums).



Application Data

- Depends on the application layer protocol used.
- Example:



Example

Physical Layer: eth
the 2 MAC addresses
+ IP indication

Network Layer: IP
IP addresses, TTL,
checksum, fragmentation

0000	00	0f	db	4d	77	95	00	0d	93	b0	a3	24	08	00	45	00
0010	01	75	c8	de	40	00	40	06	44	dd	c0	a8	01	2e	40	ec
0020	29	05	e6	10	00	50	15	86	10	4d	25	b6	67	ed	80	18
0030	ff	ff	2d	2f	00	00	01	01	08	0a	2f	41	cf	64	62	38
0040	81	9a	47	45	54	20	2f	63	6e	6e	2f	32	30	30	36	2f
0050	55	53	2f	30	32	2f	32	37	2f	6b	61	74	72	69	6e	61
0060	2e	70	6f	6c	6c	2f	74	31	2e	32	31	33	35	2e	6d	6f
0070	6e	2e	62	65	61	64	73	2e	61	70	2e	6a	70	67	20	48
0080	54	54	50	2f	31	2e	31	0d	0a	41	63	63	65	70	74	3a
0090	20	2a	2f	2a	0d	0a	41	63	63	65	70	74	2d	4c	61	6e
00a0	67	75	61	67	65	3a	20	65	6e	0d	0a	41	63	63	65	70
00b0	74	2d	45	6e	63	6f	64	69	6e	67	3a	20	67	7a	69	70
00c0	2c	20	64	65	66	6c	61	74	65	0d	0a	52	65	66	65	72
00d0	65	72	3a	20	68	74	74	70	3a	2f	2f	77	77	77	2e	63
00e0	6e	6e	2e	63	6f	6d	2f	0d	0a	55	73	65	72	2d	41	67
00f0	65	6e	74	3a	20	4d	6f	7a	69	6c	6c	61	2f	35	2e	30
0100	20	28	4d	61	63	69	6e	74	6f	73	68	3b	20	55	3b	20
0110	50	50	43	20	4d	61	63	20	4f	53	20	58	3b	20	65	6e
0120	29	20	41	70	70	6c	65	57	65	62	4b	69	74	2f	34	31
0130	37	2e	39	20	28	4b	48	54	4d	4c	2c	20	6c	69	6b	65
0140	20	47	65	63	6b	6f	29	20	53	61	66	61	72	69	2f	34
0150	31	37	2e	38	0d	0a	43	6f	6e	6e	65	63	74	69	6f	6e
0160	3a	20	6b	65	65	70	2d	61	6c	69	76	65	0d	0a	48	6f
0170	73	74	3a	20	69	2e	61	2e	63	6e	6e	2e	6e	65	74	0d
0180	0a	0d	0a													

...Mw.....\$.E.
.u..@.@.D.....@.
)....P...M%.q...
..-/. /A.db8
..GET /cnn/2006/
US/02/27/katrina
.poll/t1.2135.mo
n.beads.ap.jpg H
TTP/1.1..Accept:
/..Accept-Lan
guage: en..Accep
t-Encoding: gzip
, deflate..Refer
er: http://www.c
nn.com/..User-Ag
ent: Mozilla/5.0
(Macintosh; U;
PPC Mac OS X; en
) AppleWebKit/41
7.9 (KHTML, like
Gecko) Safari/4
17.8..Connection
: keep-alive..Ho
st: i.a.cnn.net.
...

Transport Layer: TCP
Ports, Seq Ack numbers,
checksum, timestamps

Application Layer: HTTP
Request: GET
Request URI
Referrer
User-agent info
Connection info



Internet Protocols

- Data link Layer: ethernet, wi-fi etc.
- Network Layer: ICMP (Internet control message), IP etc.
- Transport Layer: UDP (user datagram protocol), TCP etc.
- Application Layer: Finger, FTP (file transfer), HTTP (hypertext transfer), IMAP (internet message access), IRC (internet relay chat), POP (post office), SMTP (simple mail transfer), TELNET (terminal emulation), X-window, etc.



UDP Protocol

- user datagram protocol.
- lightweight alternative to TCP.
- Faster, lighter - adds 8 bytes for control.
- stateless sending and no ordering.
- used for application layer protocols as SNMP (simple network monitoring), Syslog (system audit log), Time etc.



How does data find its way?

- *Application.* HTTP request to www.website.com (DNS resolution)
- *Transport.* which in turn will result to some packets directed to a certain **port #**.
- *Internet.* which in turn will result to some frames directed to a **IP address**.
- *Physical.* which in turn will result to some actual bits being sent to **MAC address**



Addressing

- Two mappings are necessary:
 - From host name to IP address.
 - From IP address to MAC address.
- Host name will be mapped to an IP address through a protocol called DNS
- MAC address will be obtained from an address resolution table. (ARP)



Transmitting a packet

- A packet needs to be directed to a certain IP address.
- To figure out where to send it next, a routing table is consulted. Example:

Routing table:

Destination	Gateway	Flags	Refs	Use	Netif	Expire
default	137.99.11.1	UGSc	16	1863	en0	
...						



Transmitting a packet, II

- Once the (intermediate hop) IP address is determined the packet must be split into frames and directed to the right MAC address.
- Internet to Ethernet address translation:

Address	HWtype	HWaddress	Flags	Mask	Iface
137.99.11.1	ether	00:0B:46:9A:1B:3F	C		eth0



Receiving a packet

- Keep it or forward it.
- Based on destination address (and perhaps other parameters).
- Forward it using the routing table as before.
- Routing and Address Resolution tables are dynamically updated.



IP Addresses and DNS

- 32-bit (IPv6 offers 128-bits).
- IP addresses are assigned to names according to Domain Name Service (DNS).
- Given a certain name at the application layer a query will be transmitted to a *Name Server* to resolve it for the corresponding IP address.



Trace route

```
1  192.168.63.11 (192.168.63.11)  3.896 ms  2.122 ms  1.511 ms
2  195.134.67.1 (195.134.67.1)  1.794 ms  2.839 ms  1.784 ms
3  grnetRouter.L1.uoa.athens-3.access-link.grnet.gr (194.177.209.97)  1.984 ms  2.262 ms  3.360 ms
4  eie2-to-koletti1.backbone.grnet.gr (195.251.27.46)  2.196 ms  4.345 ms  3.539 ms
5  core1.ams.net.google.com (195.69.144.247)  72.084 ms  72.003 ms  72.743 ms
6  209.85.248.88 (209.85.248.88)  74.182 ms *  74.916 ms
7  64.233.175.246 (64.233.175.246)  75.377 ms  75.800 ms  74.950 ms
8  209.85.255.143 (209.85.255.143)  77.161 ms  72.14.239.197 (72.14.239.197)  79.314 ms  209.85.255.166
   (209.85.255.166)  90.792 ms
9  72.14.232.37 (72.14.232.37)  81.473 ms  72.14.232.41 (72.14.232.41)  84.950 ms  72.14.232.37
   (72.14.232.37)  84.960 ms
10 ez-in-f104.1e100.net (66.102.13.104)  79.149 ms  78.015 ms  76.316 ms
```



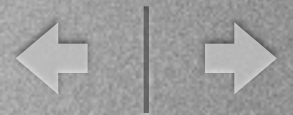

How traceroute works

- Using a special IP header field called TTL : time to live.
- TTL = number of hops a packet is allowed to make. Each router decreases by one.
- When TTL reaches 0 then a router discards the packet and notifies originator.
- For traceroute: send repeatedly packets and calibrate TTL as 1, 2, 3, 4, 5, ...
- not all routers necessarily respond (* * *)



Client Server Model

- Application protocols, FTP, HTTP, Telnet etc.
- Server listens to port for client requests.
- Client initiates protocol



Talking over TCP/IP

Client



state:closed

SYN

state:listen

state:syn-sent

SYN+ACK

state:syn-received

state:established

ACK

state:established

Application layer data

ACK

Application layer data

ACK

...

FIN+ACK

FIN+ACK

Server





Packet Sniffing

- Every computer in the Internet sends and receives packets.
- Anyone with the appropriate privileges in a certain host can “sniff” the packets that are being forwarded by the host (and not only the packets that are directed to the host).
- In this case the host is said to be in **promiscuous** mode.



Packet Sniffing, II

- If you have privileges for promiscuous mode then you can capture all traffic within the sub-network the host belongs to.
- You cannot capture traffic outside your sub-network (e.g., traffic that is not directed towards your gateway).
- When you use your computer do always ponder what is your sub-network (consider: sitting at a cafe connected to a wireless access point)



Wireshark

- Wireshark is a powerful “network protocol analyzer”
- It not only sniffs data when put on promiscuous mode but also “knows” the protocols and structures the packets in the appropriate format.



Wireshark, II

screen dump of capture window after an FTP connection

No.	Time	Source	Destination	Protocol	Info
1	0.000000	192.168.1.46	192.168.1.1	DNS	Standard query A ftp.debian.org
2	0.156872	192.168.1.1	192.168.1.46	DNS	Standard query response A 128.101.240.212
3	0.203708	192.168.1.46	128.101.240.212	TCP	58408 > ftp [SYN] Seq=0 Ack=0 Win=65535 Len=0 MSS=1
4	0.311009	128.101.240.212	192.168.1.46	TCP	ftp > 58408 [SYN, ACK] Seq=0 Ack=1 Win=5792 Len=0 M
5	0.311128	192.168.1.46	128.101.240.212	TCP	58408 > ftp [ACK] Seq=1 Ack=1 Win=65535 Len=0 TSV=79
6	0.427572	128.101.240.212	192.168.1.46	FTP	Response: 220 saens.debian.org FTP server (vsftpd)
7	0.457218	192.168.1.46	128.101.240.212	TCP	58408 > ftp [ACK] Seq=1 Ack=43 Win=65535 Len=0 TSV=7
8	3.908879	192.168.1.46	128.101.240.212	FTP	Request: USER anonymous
9	3.995051	128.101.240.212	192.168.1.46	TCP	ftp > 58408 [ACK] Seq=43 Ack=17 Win=6144 Len=0 TSV=4
10	3.995621	128.101.240.212	192.168.1.46	FTP	Response: 331 Please specify the password.
11	4.058261	192.168.1.46	128.101.240.212	TCP	58408 > ftp [ACK] Seq=17 Ack=77 Win=65535 Len=0 TSV=
12	8.388059	192.168.1.46	128.101.240.212	FTP	Request: PASS ak@ak.org
13	8.473188	128.101.240.212	192.168.1.46	FTP	Response: 230-
14	8.473824	128.101.240.212	192.168.1.46	FTP	Response: 230-This site is just another one in a worldwid
15	8.659296	192.168.1.46	128.101.240.212	TCP	58408 > ftp [ACK] Seq=33 Ack=158 Win=65535 Len=0 TSV=
16	8.751453	128.101.240.212	192.168.1.46	FTP	Response: 230-It is not the "primary Debian FTP site" - it
17	8.758911	192.168.1.46	128.101.240.212	FTP	Request: SYST

the three-way handshake

Observe the password
and username

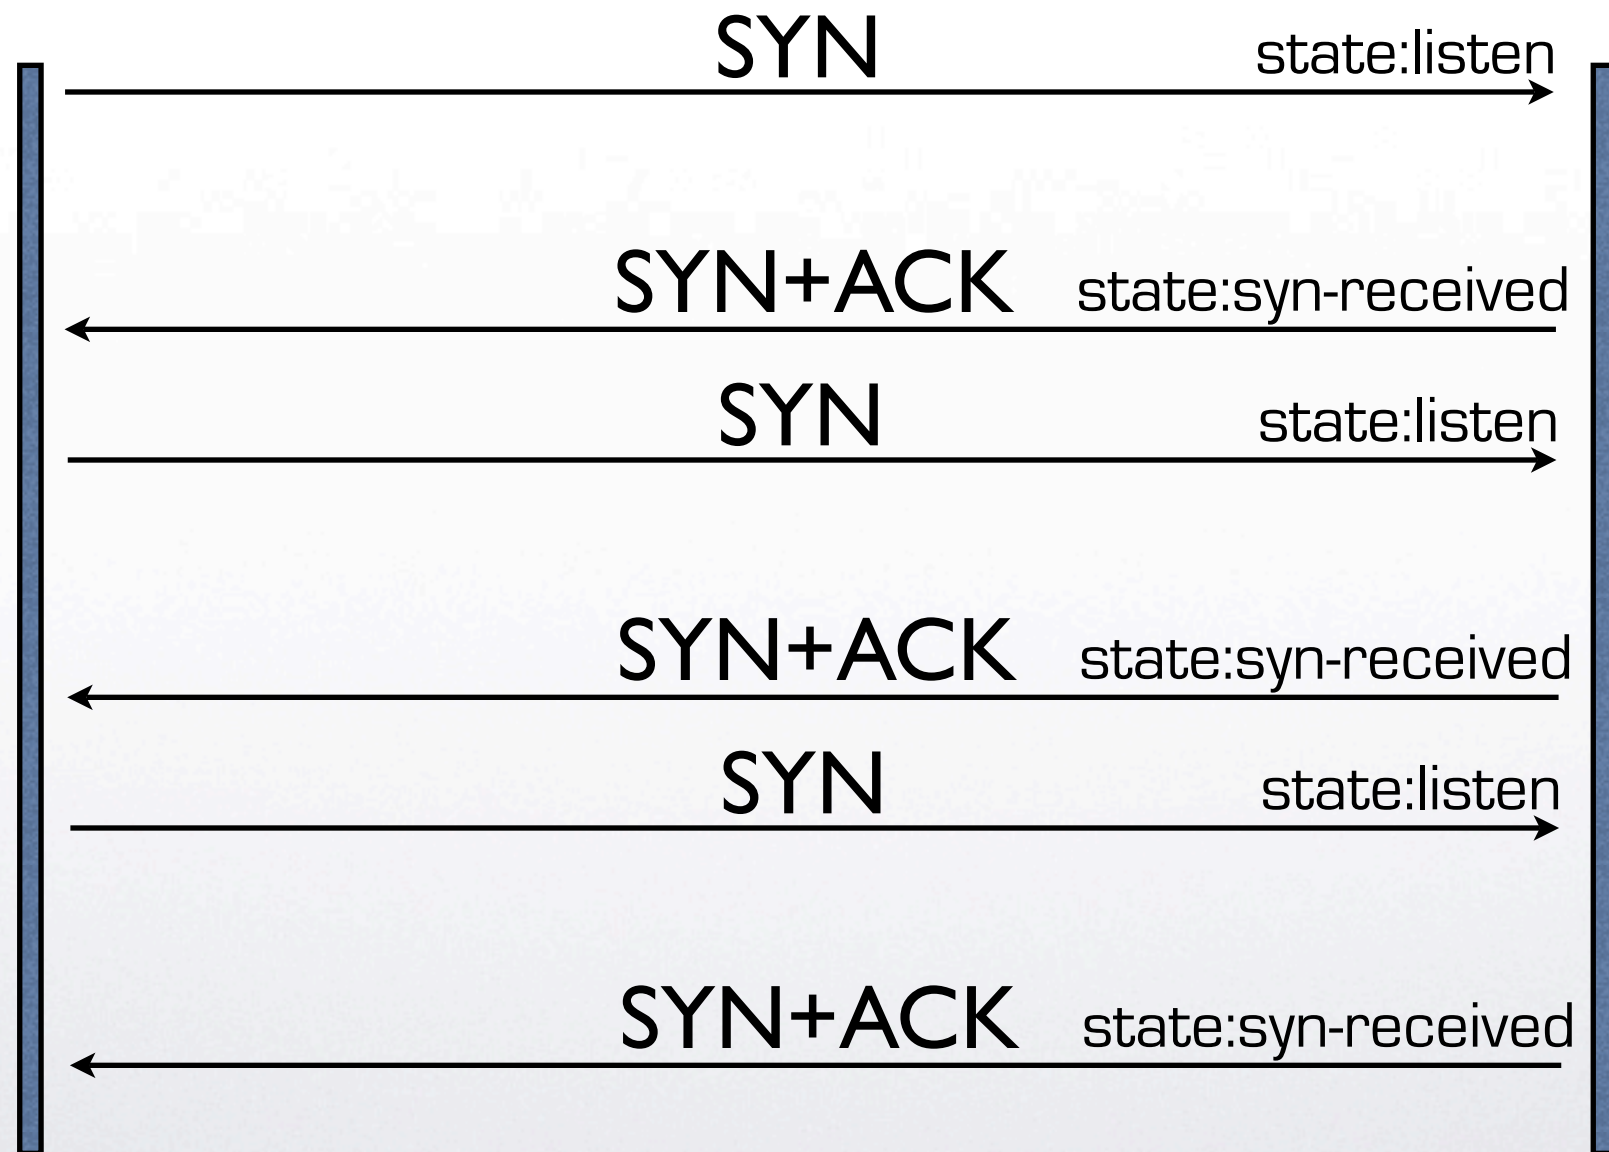



A malicious client

Client



Server



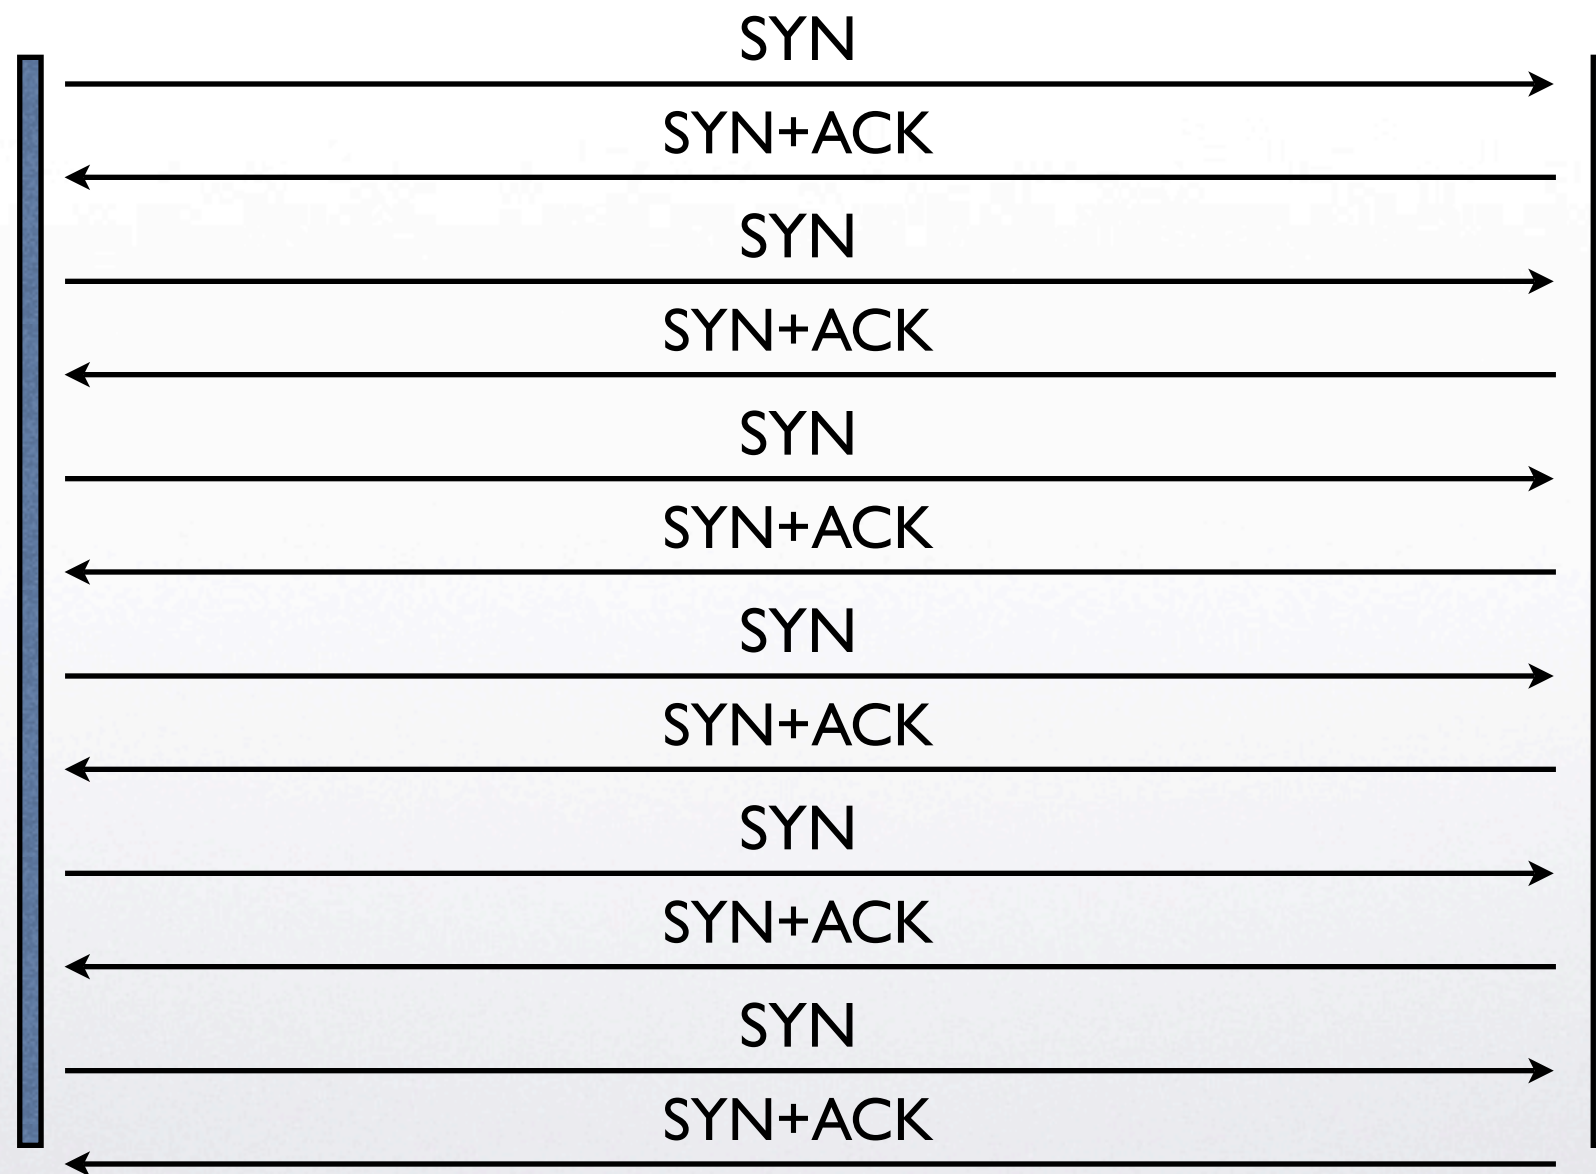


A malicious client, II

Client



IP
Spoofing



Server





SYN Flooding

- Client bombards Server with SYN packets that are never ACK'ed.
- *IP spoofing* can be used to make packets look like they are coming from other places.
- Physical limitation: *bandwidth*.
- If bandwidth on client is substantial compared to server there is serious potential for a *Denial of Service (DoS) Attack*



DoS Attacks

- Deplete / misconfigure / misallocate the resources on a target server host so that it cannot serve its clients.
- resources:
 - bandwidth.
 - memory.
 - cpu.
 - ...



DoS By Flooding

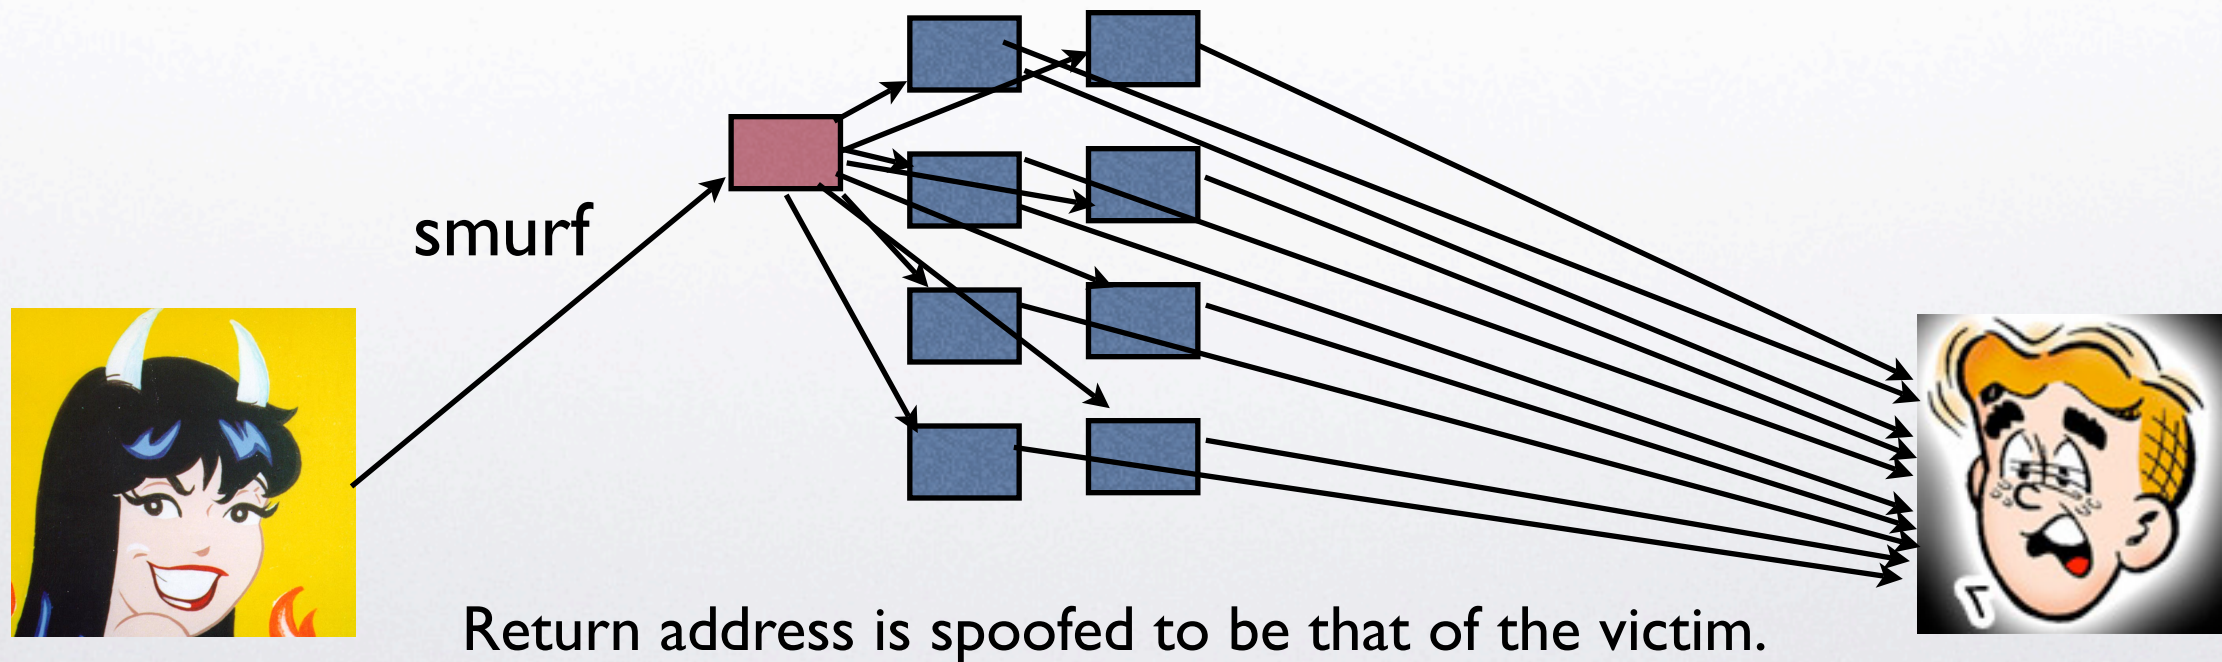
- SYN Flooding - we just saw it.
- Ping Floods: this is a flooding of ICMP Echo Request packets.

```
aggelos@grub:~$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1): 56 data bytes
64 bytes from 192.168.1.1: icmp_seq=0 ttl=64 time=0.686 ms
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=0.611 ms
64 bytes from 192.168.1.1: icmp_seq=2 ttl=64 time=0.617 ms
^C
--- 192.168.1.1 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max/stddev = 0.611/0.638/0.686/0.034 ms
```




Smurf Attack


- An enhanced Ping flood attack that utilizes IP Broadcast:
- it is possible to specify the destination IP address as a broadcast to all hosts in a subnetwork.





Distributed DoS

- The real deal!





Botnet Creation

- Host compromised by virus, worm, trojan.
- Runs rootkit remote administration tool.
- When commanded it launches DoS attack against victim.
- Attack seems to be coming from everywhere!



Current Botnet Uses

- The convenience of separation between ‘hacking a computer’ and ‘committing a crime.’
- Sending Spam.
- Click-Fraud.
- Identity Theft.
- Stealing files: e.g., game Diablo-2 items were stolen and sold on EBay...



Botnet Design Challenges

- Choosing botnet topology. (centralized, hierarchical, P2P).
- Command&Control flow (pull, push, ongoing).
- Command&Control protocols (old-use: IRC, IM, HTTP, specialized protocols).



Rallying Mechanism

- Botnet assembly:
 - hardcoded IP addresses.
 - list of DNS names.
 - Distributed DNS services implemented by the botnet itself.
 - DNS Fast Flux



Design challenges

- Bot should try to avoid being detected due its C&C communication or O/S *footprint*.
- Botnet may experience *Bot loss*.
- Botnet may experience *Bot stealing*.
- Botnet should detect fake bots that try to infiltrate.



Other DoS attacks

- **The Fork Bomb**: any program that constantly forks by creating child processes that do the same.
- Any time a process is called the O/S allocates memory for the process's requirements and enters the process specifics in a data structure.

Example:

```
int main(void) {  
    while(1) {  
        fork();  
    }  
    return 0;  
}
```

this little program may
crash your PC

You can write fork-bombs
for all major languages.

bash example `:(){ :|:& }::`



Defending against DOS

- Attacks like Smurf, Fork Bombs etc. result from the ability of an entity to allocate more resources than necessary in normal operation.
- Restricting such capability will thwart the related attack, e.g.,
 - restrict the use of IP broadcasting.
 - restrict the number of processes a user may create.

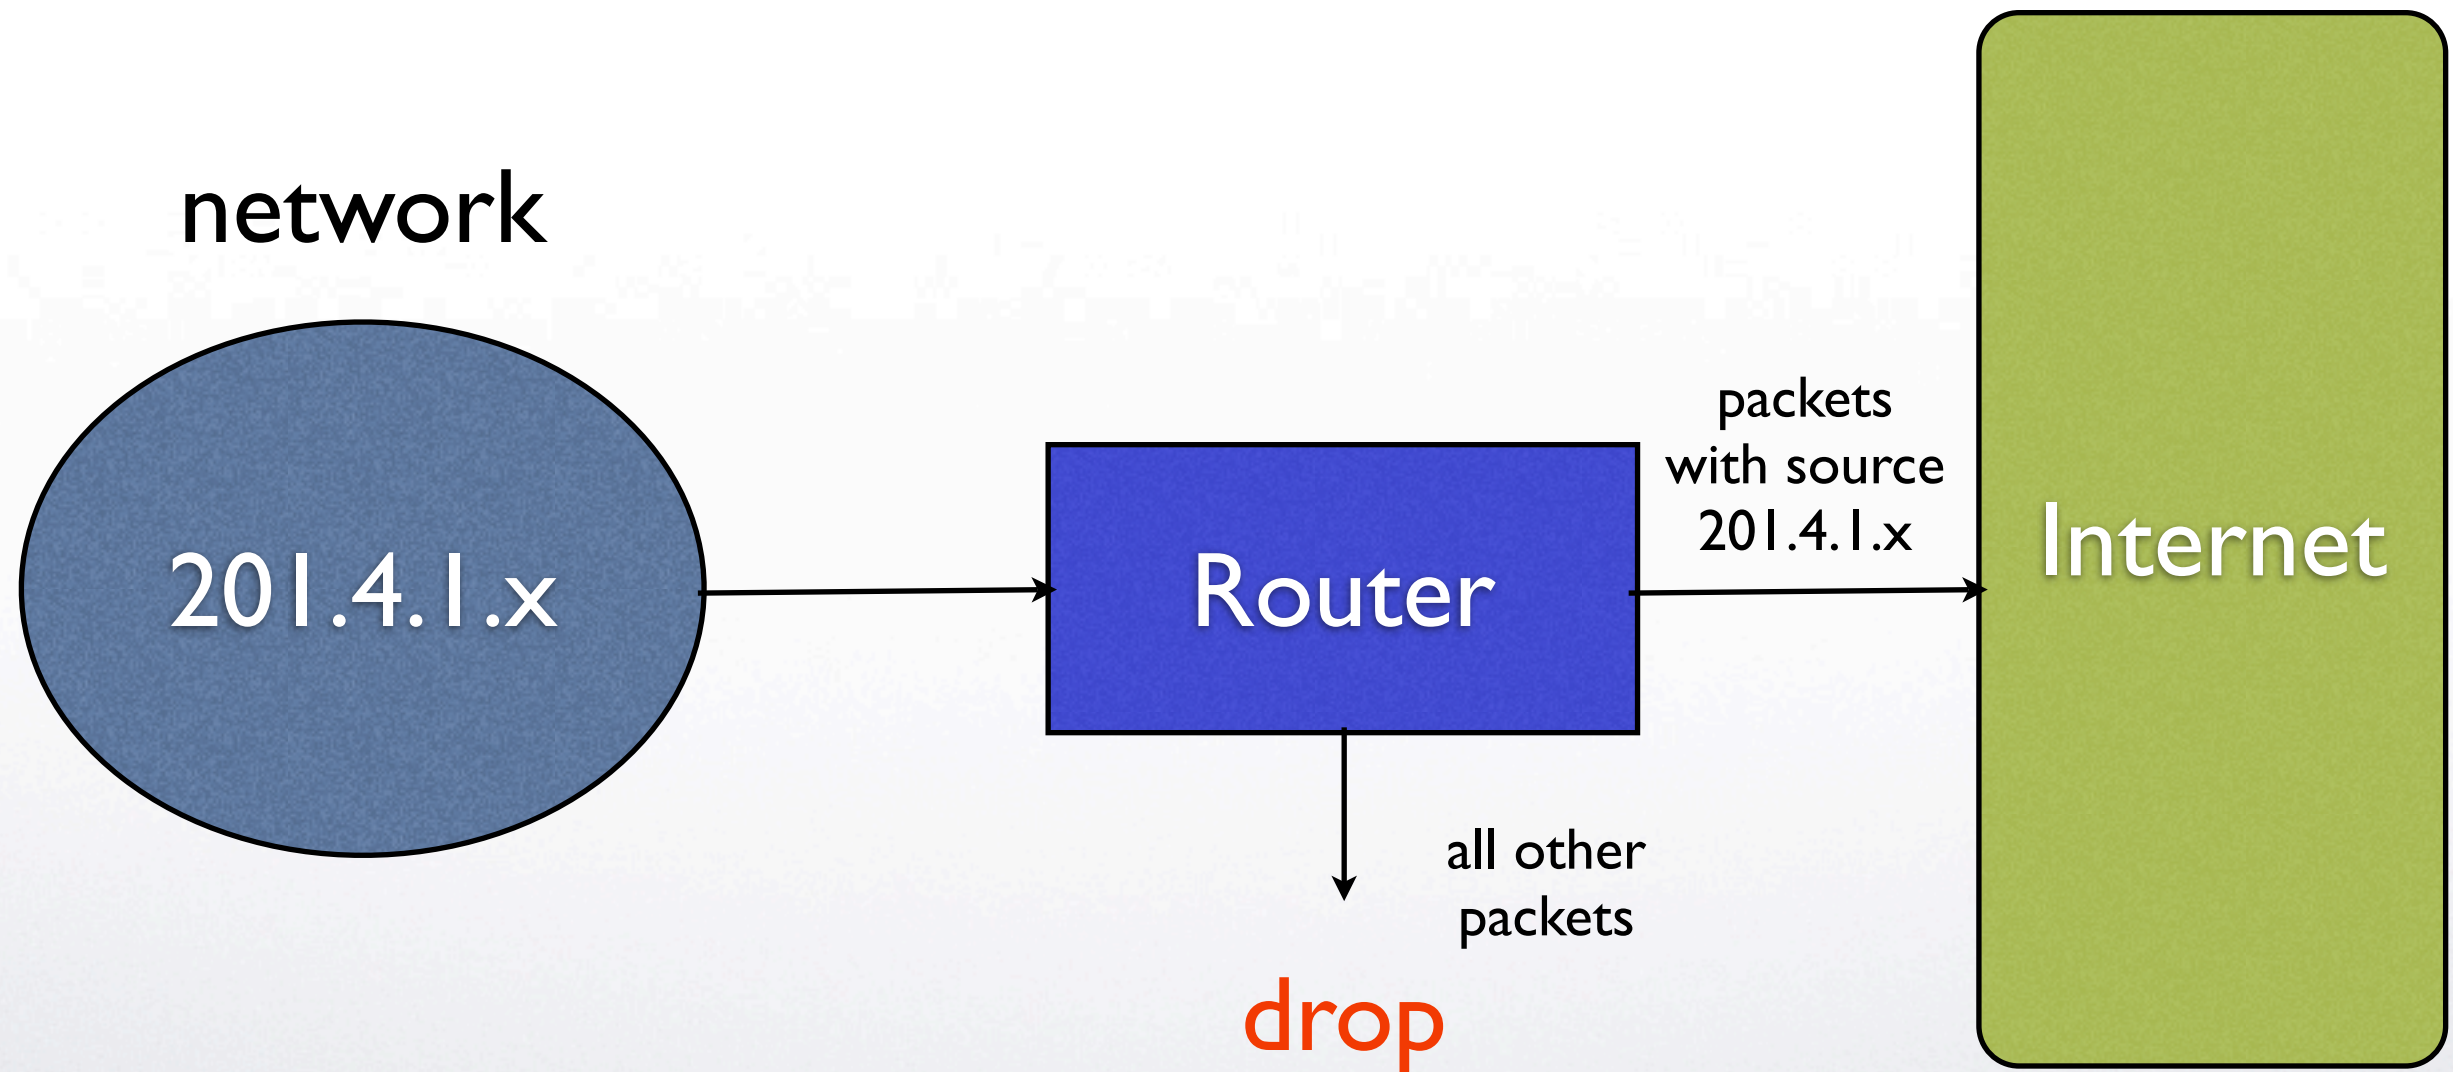


IP Throttling

- Any IP address that issues a big upstream will have its incoming packets being dropped at a certain rate.
- **Advantage:** can be configured in your router. No need to modify client/server.
- **Disadvantage:** what are the right settings?



Ingress Filtering



<http://www.faqs.org/rfcs/rfc2827.html>



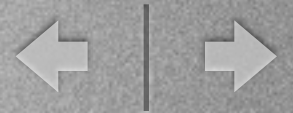
Defending against DDoS

- Harder since an attack may look as quite legitimate traffic: (e.g., HTTP)
- Challenge: distinguish between good traffic and DDoS traffic.
- We will examine some general approaches for DDoS defense next.



Simple Measures

- Try to use various parameters to make a system understand it is being attacked:
 - counting number of half-open connections (syn-rcvd).
 - counting number of connections refused.
- Once some **red-flags** are observed the system may shorten the time that it keeps half-open connections active.



SYN Cookies

Client



state:closed

SYN

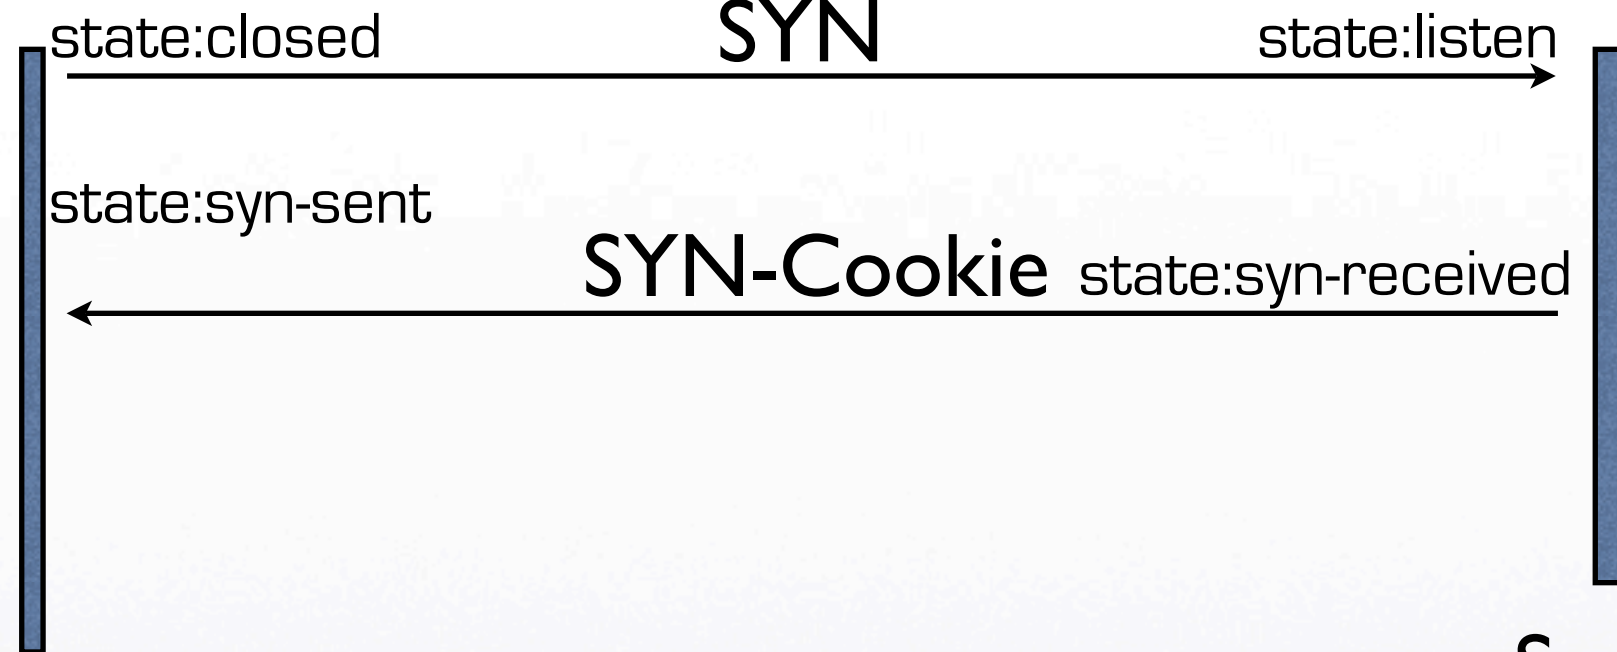
state:listen

state:syn-sent

SYN-Cookie

state:syn-received

Server



A SYN-Cookie contains info derived by the source-address port, destination address and port etc.

Server can forget everything about the client till the ACK message that will include the SYN-Cookie

outsourcing the server state to the client



Proofs of Work

- Client is requested to solve a puzzle in order to establish a connection.
- **Pros:** works in any client-server system. can be tuned to slow down malicious systems.
- **Cons:** needs changes in the basic infrastructure (e.g., changes in both client and server). Will not eliminate bandwidth emaciation.



Example.

From Juels/Brainard (RSA)

Hash function

$$\mathcal{H} : \{0, 1\}^* \rightarrow \{0, 1\}^k$$

$r = \text{secret} \parallel \text{time} \parallel \text{client request}$

r
↓

$$\mathcal{H} \longrightarrow x = \mathcal{H}(r) \in \{0, 1\}^k \xrightarrow{\text{trunc}} x' \in \{0, 1\}^{k-w}$$

Server to client: sends $x', y = \mathcal{H}(x)$ stores (x, y)

Client needs to respond with $x'' : \mathcal{H}(x' \parallel x'') = y$
 $x'' \in \{0, 1\}^w$



IP Traceback

- A “forensic” technique:
 - try to identify the real path of a packet.
 - Have routers mark packets with their IP address. => problem: not enough space in a packet.



Edge Marking.

- With some low probability q a router **marks** a packet. **Marking is defined as:**
 - IF packet unmarked:
 - enter your IP address as the start IP address of **the edge**. Set **distance** = 0
 - ELSE: If **distance** = 0, enter your IP address as the end IP address of **the edge**. increment the **distance**.
- **Otherwise** increase the **distance**.



Encoding IP addresses

- Use only the 16 bits of the IP identification field used for fragmentation.
- How to pack 64+5 bit information into 16 bits?

$$\mathcal{H}_1, \mathcal{H}_2 : \{0, 1\}^{32} \rightarrow \{0, 1\}^{11}$$

- Use
first router $\mathcal{H}_1(\text{startIP})$ second router $\mathcal{H}_1(\text{startIP}) \oplus \mathcal{H}_2(\text{endIP})$

birthday paradox says that there will be 50/50 odds of a collision at around $\sqrt{2^{11}}$ which is about $2^6 = 64$ hops. This means that we are only able to hash from 32 bit IPs down to about 11 bits before noise overwhelms the hashing



Reconstruct Path

sender

server



$$\mathcal{H}_1(\text{startIP}_1) \oplus \mathcal{H}_2(\text{endIP}_1)$$

$$\mathcal{H}_1(\text{startIP}_2) \oplus \mathcal{H}_2(\text{endIP}_2)$$

$$\mathcal{H}_1(\text{startIP}_3) \oplus \mathcal{H}_2(\text{endIP}_3)$$

obtained
from different
packets
originating
from the same
source

reconstruct

given endIP_3
you recover

$$\mathcal{H}_1(\text{startIP}_3)$$

from which

you recover
 startIP_3

which equals
 endIP_2

etc.

<http://www.ece.cmu.edu/~dawnsong/papers/iptrace.ps>