



O/S & Access Control

Aggelos Kiayias



One system Many users

- Objects that require protection
 - memory
 - I/O devices (disks, printers)
 - programs and processes
 - networks
 - stored data in general



Separation Layers

- keeping one user's objects separate from others.
 - *Physically* - Separate devices
 - *Logically* - e.g. Access Control Lists
 - *Cryptographically* - Encryption

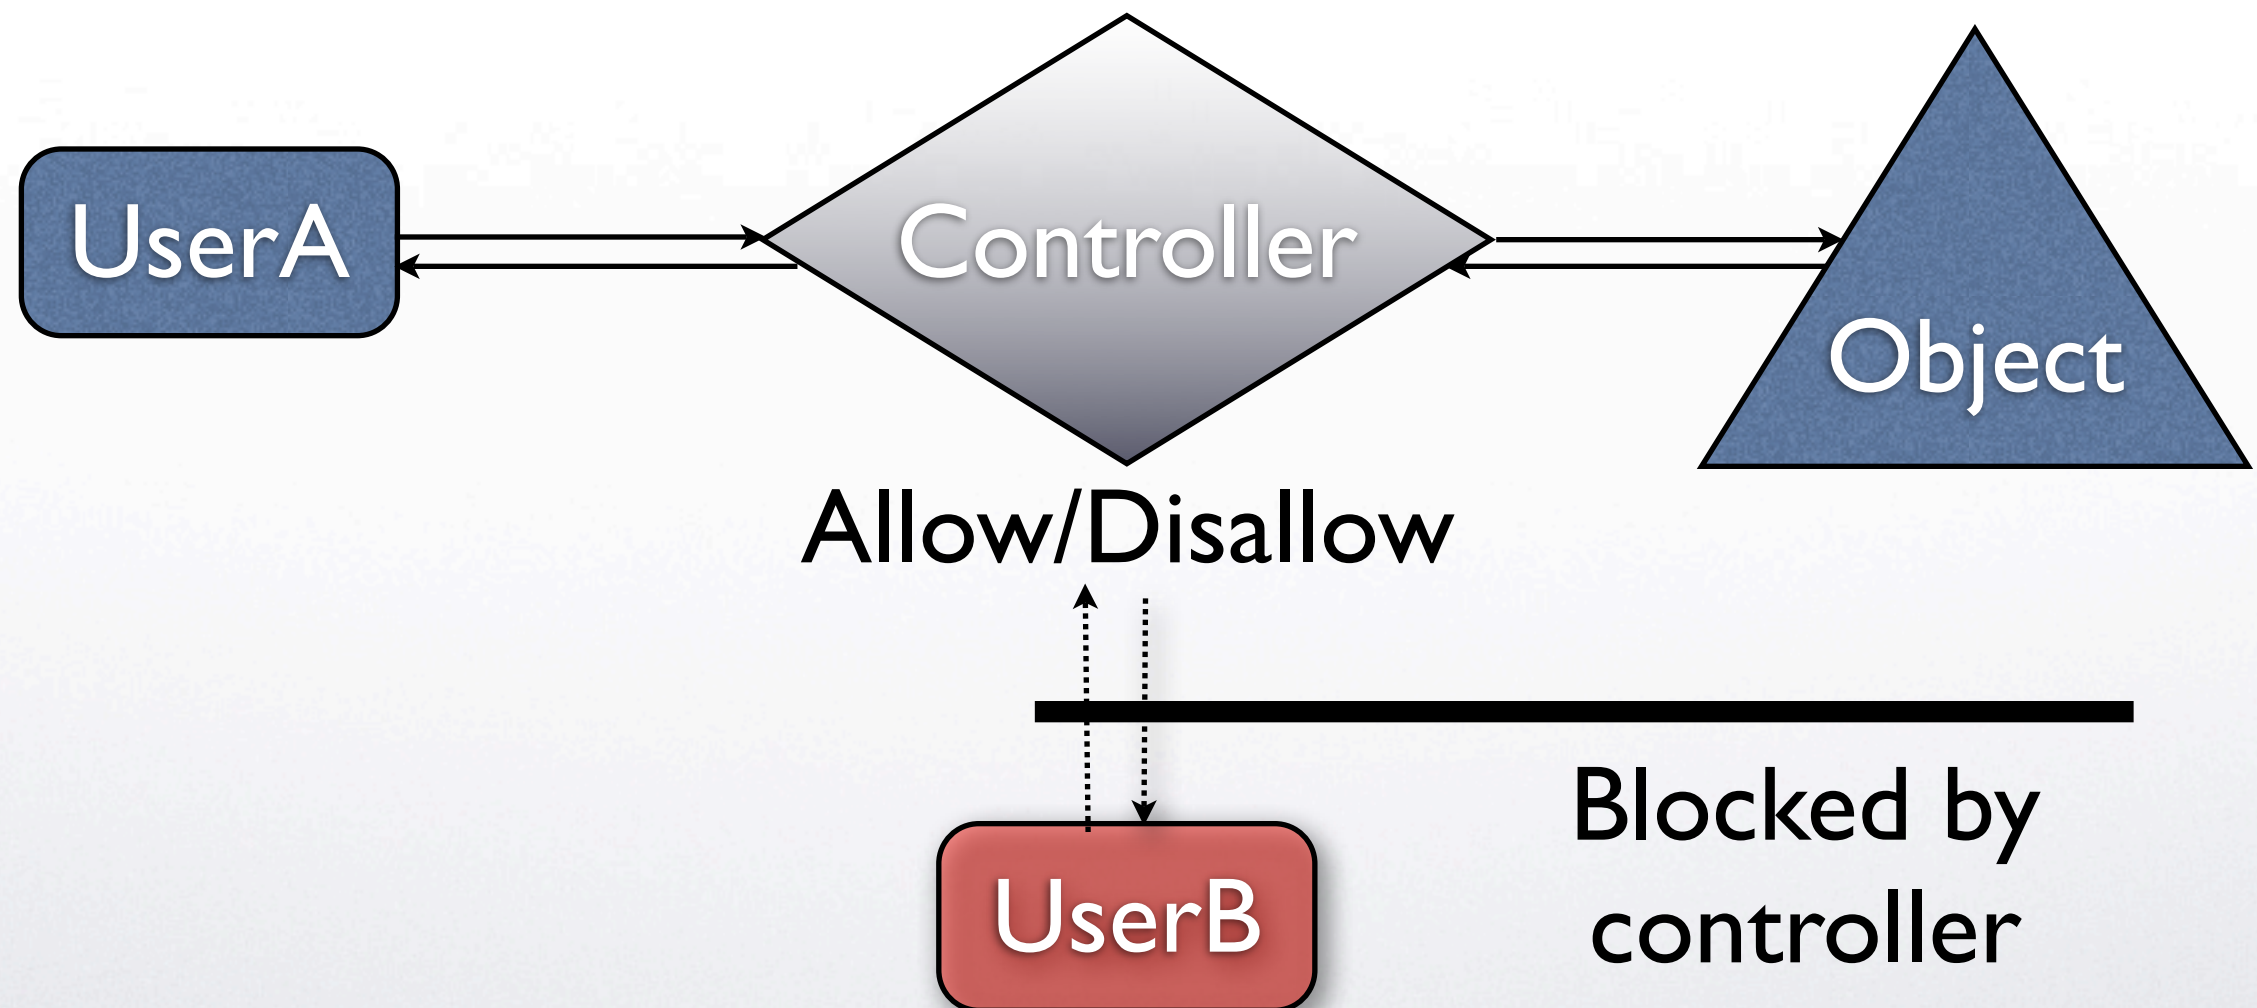


Sharing is Inevitable

- Users need to coexist on a single system:
- Two extremes:
 - Monolithic*: all users share a single account
 - Isolation*: every resource is assigned to strictly one user
- In real world systems, a blend of the above is typically used



Logical Access Control





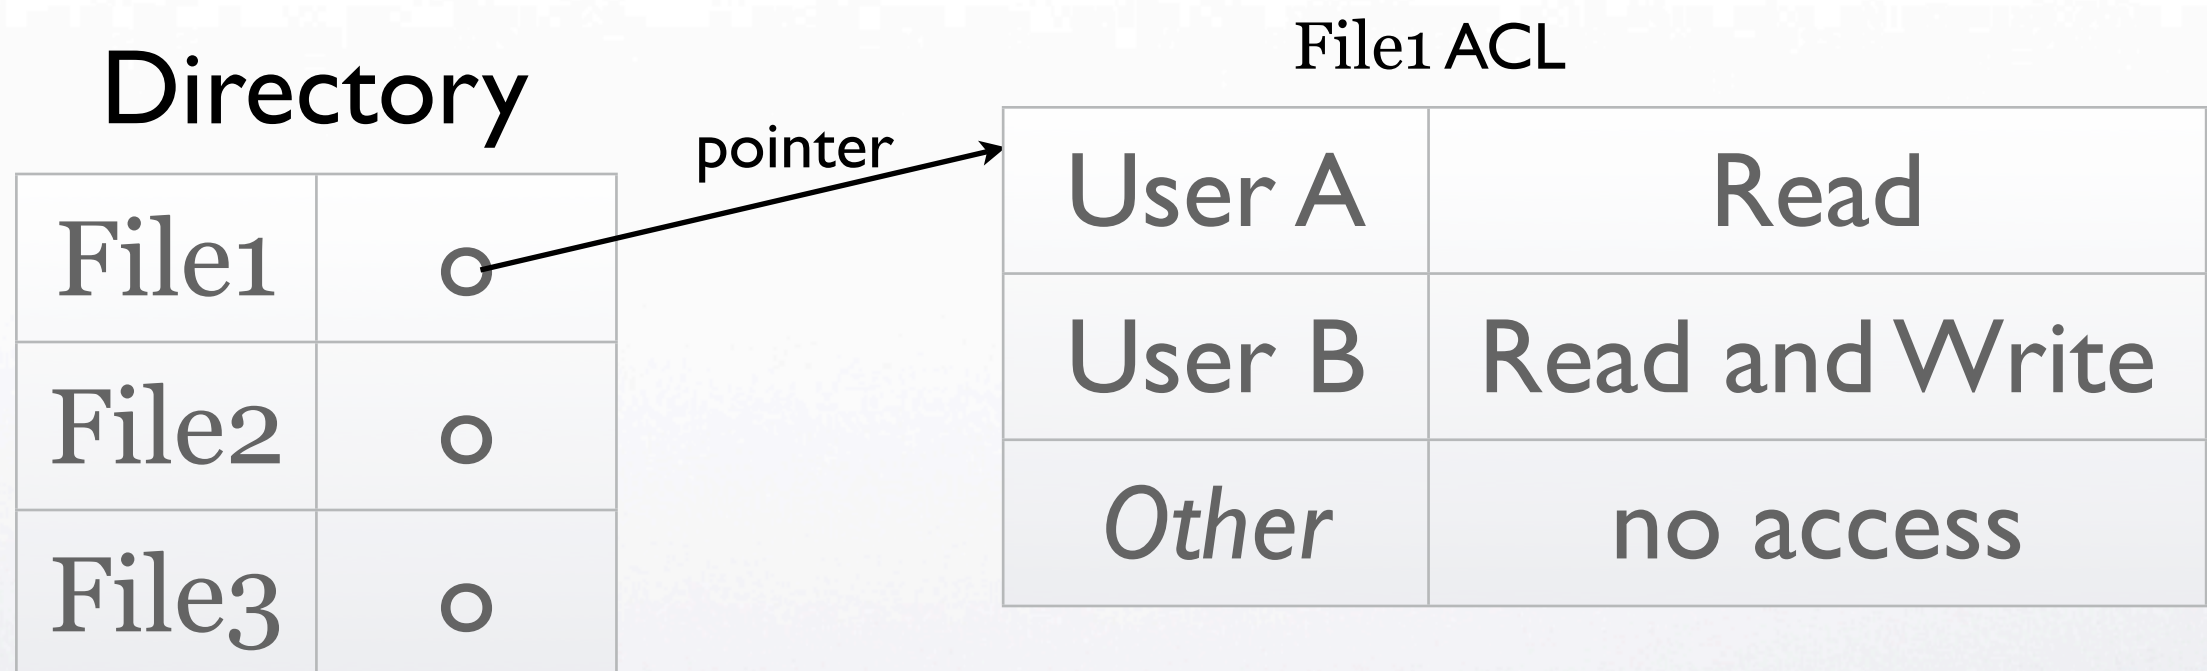
Focus: file system

- A paradigm for access control.
- All objects can be thought as files (*NIX).



Access Control List

- Each object has an ACL.





Windows NTFS (5+)

- ACL is stored with every file.
- Contains users and groups and corresponding permissions for each.
- Folder permissions:
Read, Write, List, Read & Execute, Modify, Full Control



In Unix/Linux

- Processes make requests to access resources.
- Each process is associated with a **uid**.
- Each file has an ACL that contains a triple of
rwX rwX rwX
user group other
- The ACL contains both *user* and *group* info.
- **x** is *execute* for files and *access* for dirs.



Processes

- Objects are accessed by processes.
- How is a process assigned a **uid**?
 - It is created by a parent process and inherits the **uid** of the parent process.
 - But how does a user access a system?



Login Process (console)

- Prompt process: *invites user* (cf., getty running as root)
- and *challenges user to authenticate*.
- if login is unsuccessful, restart the prompt
- if successful an interface process is spawned that inherits the **uid** and **gid** of the authenticated user.



Temp Acquired Permission: *suid bit*

- How is it possible to allow a certain **uid** to peep into a higher access level via an executable?
- When an executable has the *suid bit* set, an executed file inherits the **uid** of its *owner* rather than the **uid** of the *caller*. E.g.,
-rwsr-xr-x 1 root wheel 32680 2013-10-11 12:13 passwd



Separation

- in a multi-user environment
 - Access-control as described so far offers a logical separation; is this foolproof?
 - What would a cryptographic separation offer?



User Authentication

- Can be based on:
 - Something the user *knows*. (e.g. Password)
 - Something the user *has*. (e.g. Physical Token)
 - Something the user *is*. (e.g. Fingerprint)
- Two Factor Authentication (or Multi Factor)



Password-based Auth

- Authentication based on what a user *knows*.
- O/S must keep a database of username/password pairs.
- Where to store it?
- What to store?



The `/etc/passwd` file in UNIX

- **FORMAT**
Name:Password: UserID:PrincipleGroup:Gecos: HomeDirectory:Shell
- `guest:AvCSyg9e75YZM:200:0::/home/guest:/usr/bin/sh`
- One line for each user.
- The file is publicly readable.
 - In current deployments the passwords are *shadowed* in another location (e.g., `/etc/shadow`) -- this file is not publicly readable.

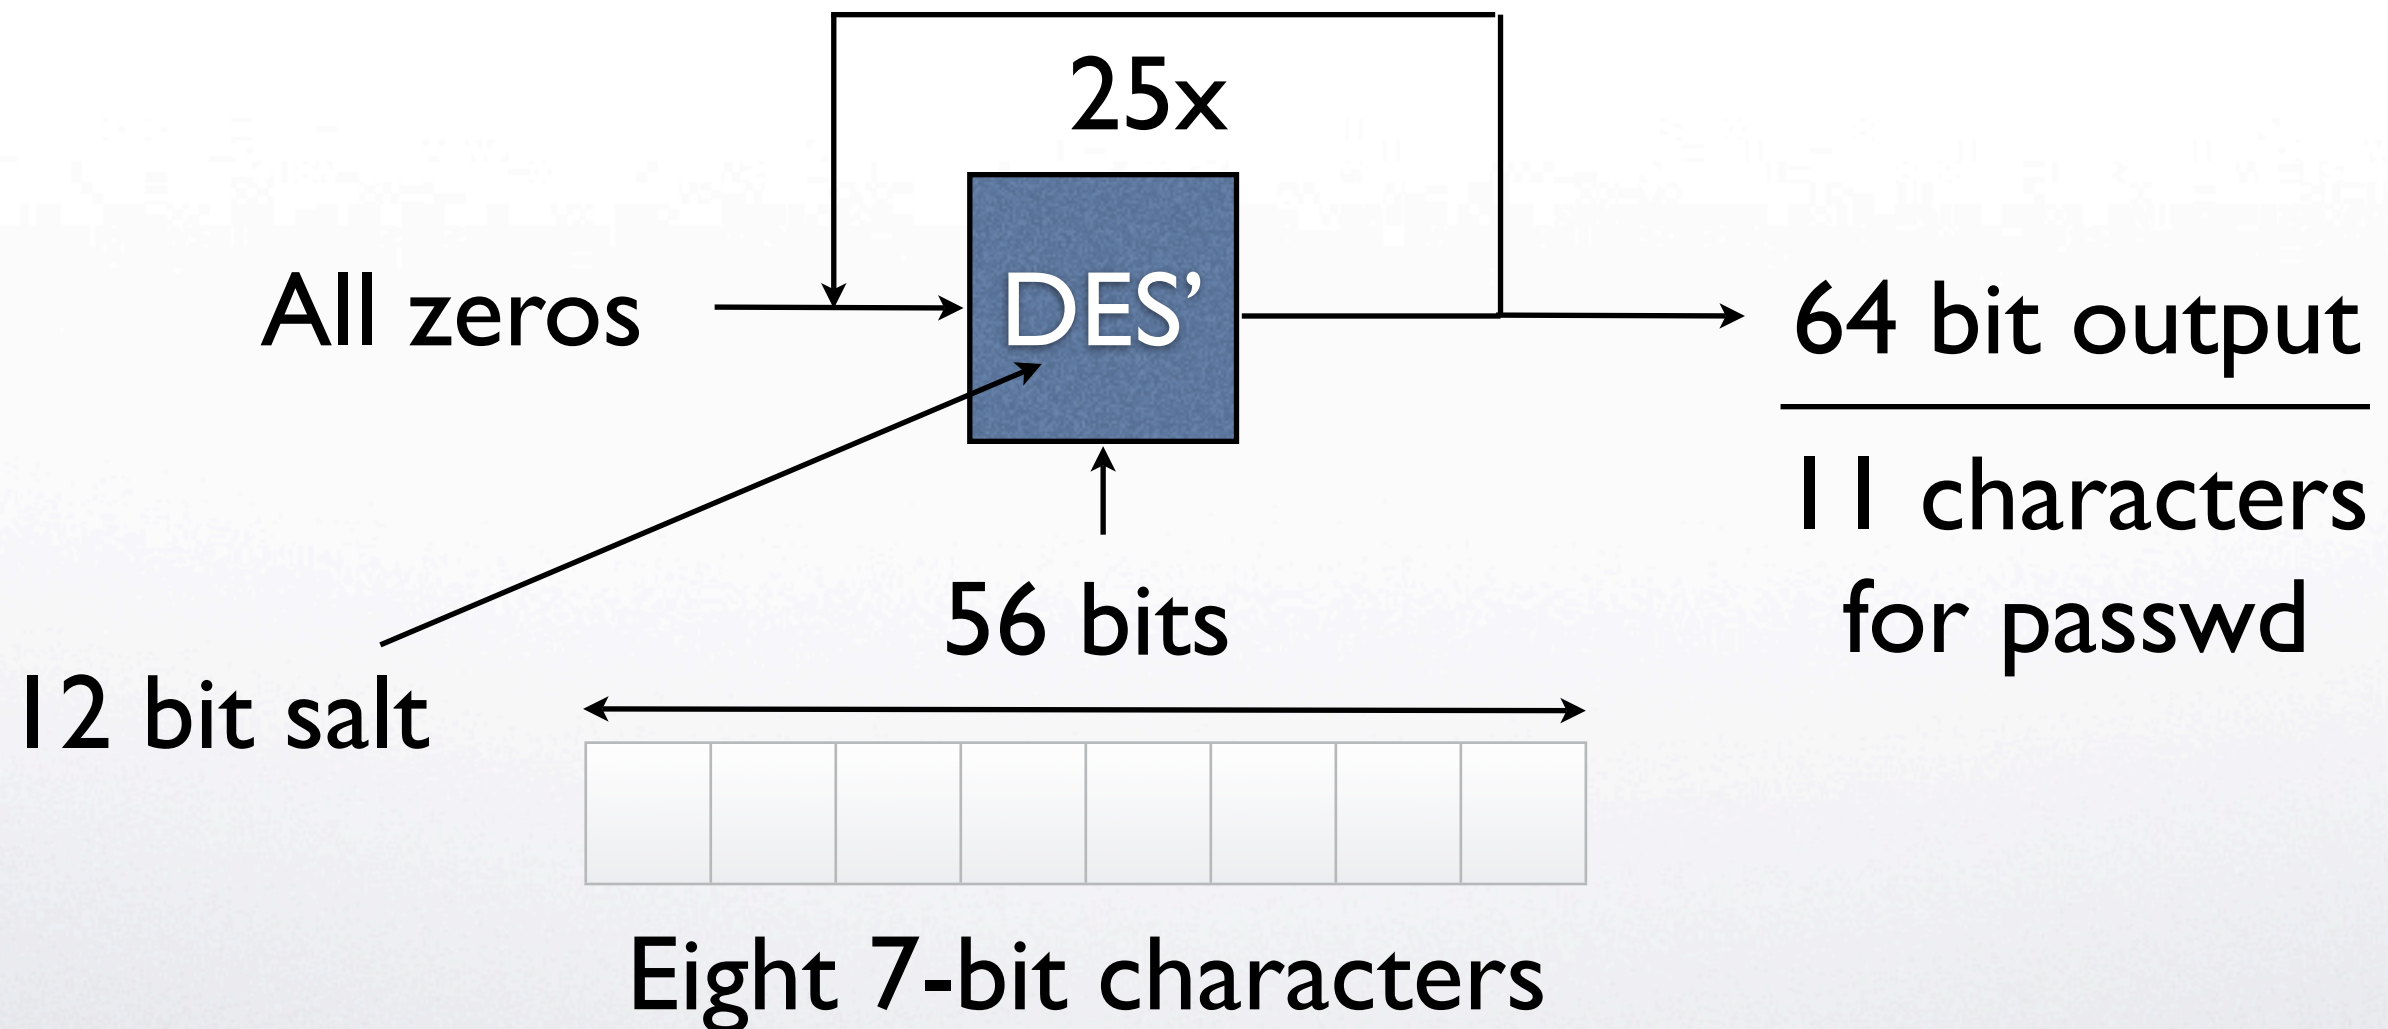



Storing passwords

- Should passwords be stored in the clear?
- No, use a one-way transformation.
- Can be based on a hash function.




The old crypt() function





Examples

- `crypt("password", "Ee") = EeAjqAJ0sluG.`

- `crypt("password", "4!") = 4!wpbYhg6VW8qM`
- `crypt("password is what some people choose but I chose a passphrase!", "4!") = 4!wpbYhg6VW8qM`

this results in a collision since the DES' based crypt function only used the first 8 characters



The glibc2 extension

- If salt starts with \$1\$ followed by at most 8 characters, terminated by \$; then it is not using the DES based algorithm.
- MD5 based algorithm with 22 char output from [a-zA-Z0-9./].
- entire password is now significant.



Examples

- `crypt("password", "IGoodSalt") =`
`IGoodSalt$czxNIPirYBY5pqEI Q98el.`
- `crypt("password is what people choose but`
`I chose a passphrase", "IGoodSalt") =`
`IGoodSalt$Obp/S5k35O0rIymT0v9t./`
- currently `$2y$=Blowfish`, `5=SHA-256`,
`6=SHA-512`
 - *test on linux:* `perl -e 'print(crypt("password", "\$I\$GoodSalt")."\n");'`



In Windows?

- Security Accounts Management Database (SAM) stored in the registry (hive).
- It stores hashed copies of user passwords.
- The database itself is encrypted with a locally stored system key.
- It is possible to store this key elsewhere.
 - Attack against NT4.0, 2000 if SAM was deleted one gets a free login.
 - check: *Offline NT Password and Registry editor* or (*NTPASSWD live linux CD*)



Online Dictionary Attack

- Given a dictionary of possible passwords.
- you have a way to test whether a guessed password is correct.
- e.g., you have access to console login, or you have the password hash and the salt.
- Salting is not intended to protect against this attack.

check: John the Ripper (JTR) password cracker <http://www.openwall.com/john/>



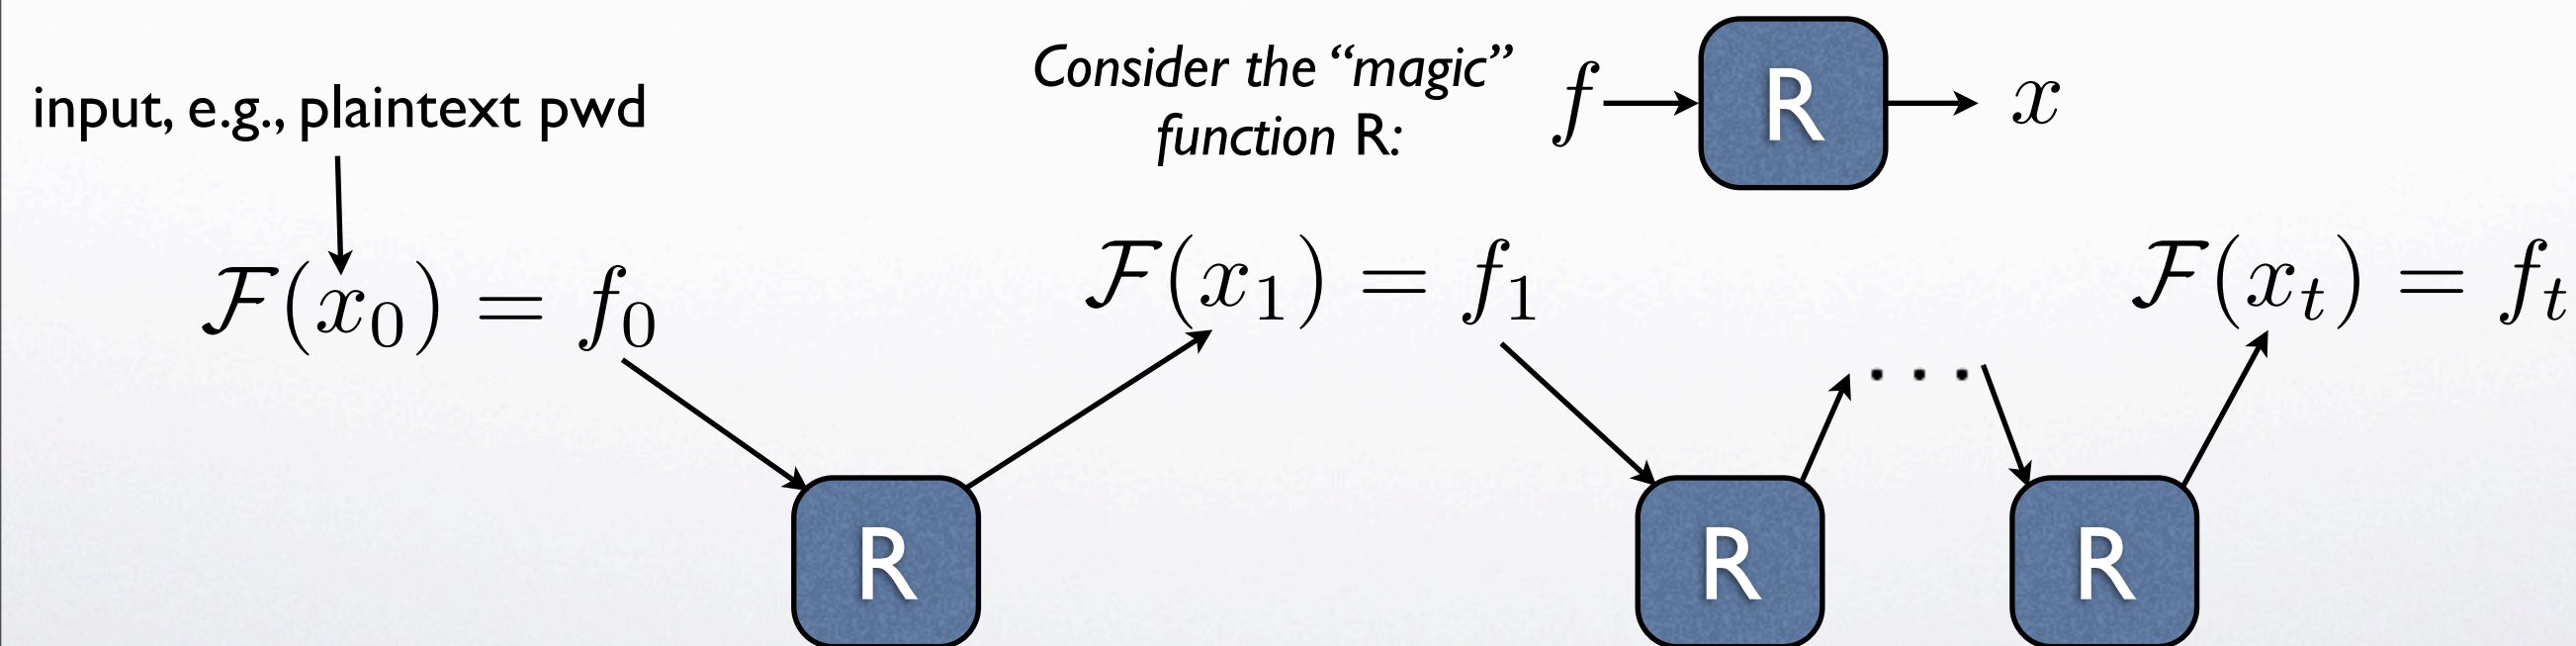
Codebook Dictionaries

<ul style="list-style-type: none">● Produce a “codebook dictionary”<ul style="list-style-type: none">● Apply one-way transformation to each candidate pwd.● Sort according to transformation output	<p><i>linear in dictionary size</i></p> <p>Offline!</p>
<ul style="list-style-type: none">● Given the password hash, binary search through the codebook dictionary.● <u>the online cost is low! but salting can really make a difference against this attack.</u>	<p><i>logarithmic in dictionary size</i></p> <p>Online!</p>



Time-Memory Tradeoff

- Can we do something between the previous approaches? Build *rainbow* tables to invert any function \mathcal{F} (e.g., crypt) with output value f (the obfuscated password)



Storage reduction:

$$(\underline{x_0}, f_0), (x_1, f_1), \dots, (x_t, \underline{f_t})$$



Time-Memory Tradeoff

Create the rainbow table by sorting according to end of chains_{offline}

_{online}

Given $f = f[0]$ calculate $f[1], \dots, f[t]$ applying **R**

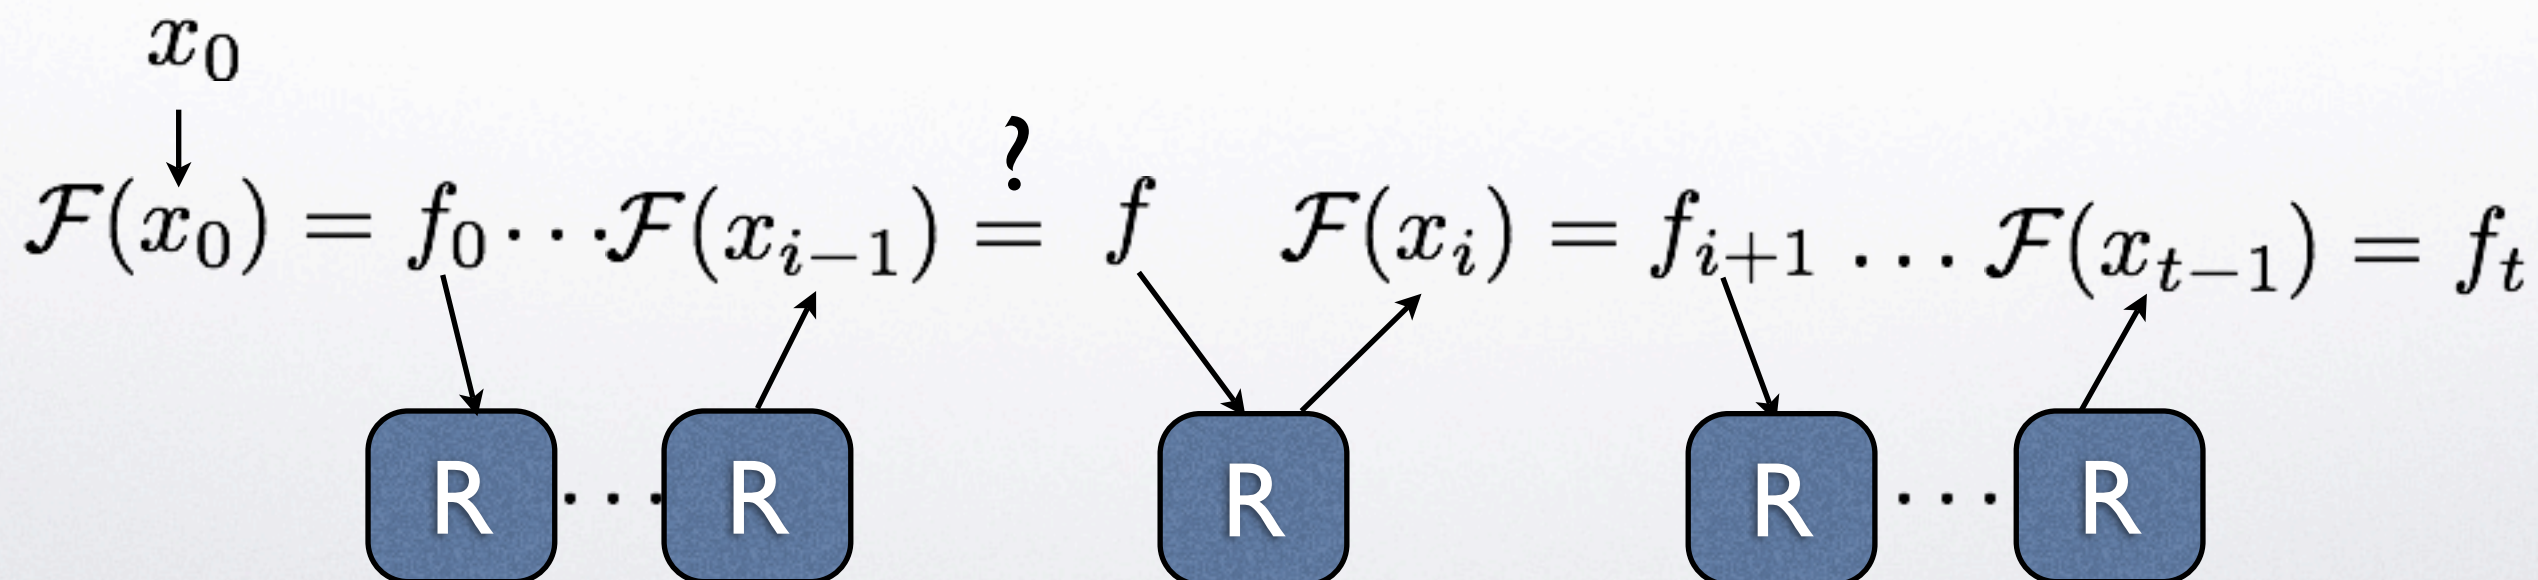
Perform binary search in the codebook dictionary
for each of $f = f[0], f[1], \dots, f[t]$
every chain hit gives a candidate password

Tradeoff: Dictionary size has been reduced by size $\sim t$
searching time has been multiplied by $\sim t$



Time-Memory Tradeoff

- How to recover the password after you hit end of chain?
- Start from the beginning of chain.





Time-Memory Tradeoff

- Tight tradeoff is contingent on a good choice of **R**
- Too few/short chains may not cover the full dictionary.
- Too many/long chains will overlap and waste space/time.



Time-Memory Tradeoff

- Even possible to model **R** to produce “human” passwords, i.e., consider those chains for which it holds that x follows a certain distribution



Time-Memory Tradeoff

- Rainbow tables
 - are a very powerful technique if applied against unsalted hashes - can break any **strong** human memorizable password.
- Implementations: *Ophcrack*, *RainbowCrack*.
- **Random** Windows NT Lan Manager passwords can be broken in 13 seconds with 1.4 GB tables. [*Oeschlin* CRYPTO '03]



Choosing a Dictionary

- Without salting one is totally vulnerable (even with **random** *but of human-memorizable length* passwords).



How does salting help?

- No Salting : you want to invert $\mathcal{F}(\cdot)$
- Salting : you want to invert a member of $\{\mathcal{F}_s(\cdot) \mid s\}$

It is possible to build rainbow tables that simultaneously walk along passwords and members of $\{\mathcal{F}_s(\cdot) \mid s\}$

Complexity : is multiplied by a factor equal to the family size



Language Entropy

http://csrc.nist.gov/publications/nistpubs/800-63/SP800-63v6_3_3.pdf



Language Entropy

- Shannon : *entropy* of English language is 1.5 bits per character for 8-char long words {a...z} *versus* $\lg(26)=4.7$ bits theoretical max for purely random selection.

http://csrc.nist.gov/publications/nistpubs/800-63/SP800-63v6_3_3.pdf



Language Entropy

- Shannon : *entropy* of English language is 1.5 bits per character for 8-char long words {a...z} *versus* $\lg(26)=4.7$ bits theoretical max for purely random selection.
- How much entropy do Human memorizable passwords have?

http://csrc.nist.gov/publications/nistpubs/800-63/SP800-63v6_3_3.pdf



Language Entropy

- Shannon : *entropy* of English language is 1.5 bits per character for 8-char long words {a...z} *versus* $\lg(26)=4.7$ bits theoretical max for purely random selection.
- How much entropy do Human memorizable passwords have?
- NIST: Using an allowed 94 character alphabet.

http://csrc.nist.gov/publications/nistpubs/800-63/SP800-63v6_3_3.pdf



Language Entropy

- Shannon : *entropy* of English language is 1.5 bits per character for 8-char long words {a...z} *versus* $\lg(26)=4.7$ bits theoretical max for purely random selection.
- How much entropy do Human memorizable passwords have?
- NIST: Using an allowed 94 character alphabet.
 - 8 chars: only 18 bits! vs. 52 bits for random

http://csrc.nist.gov/publications/nistpubs/800-63/SP800-63v6_3_3.pdf



Kerberos

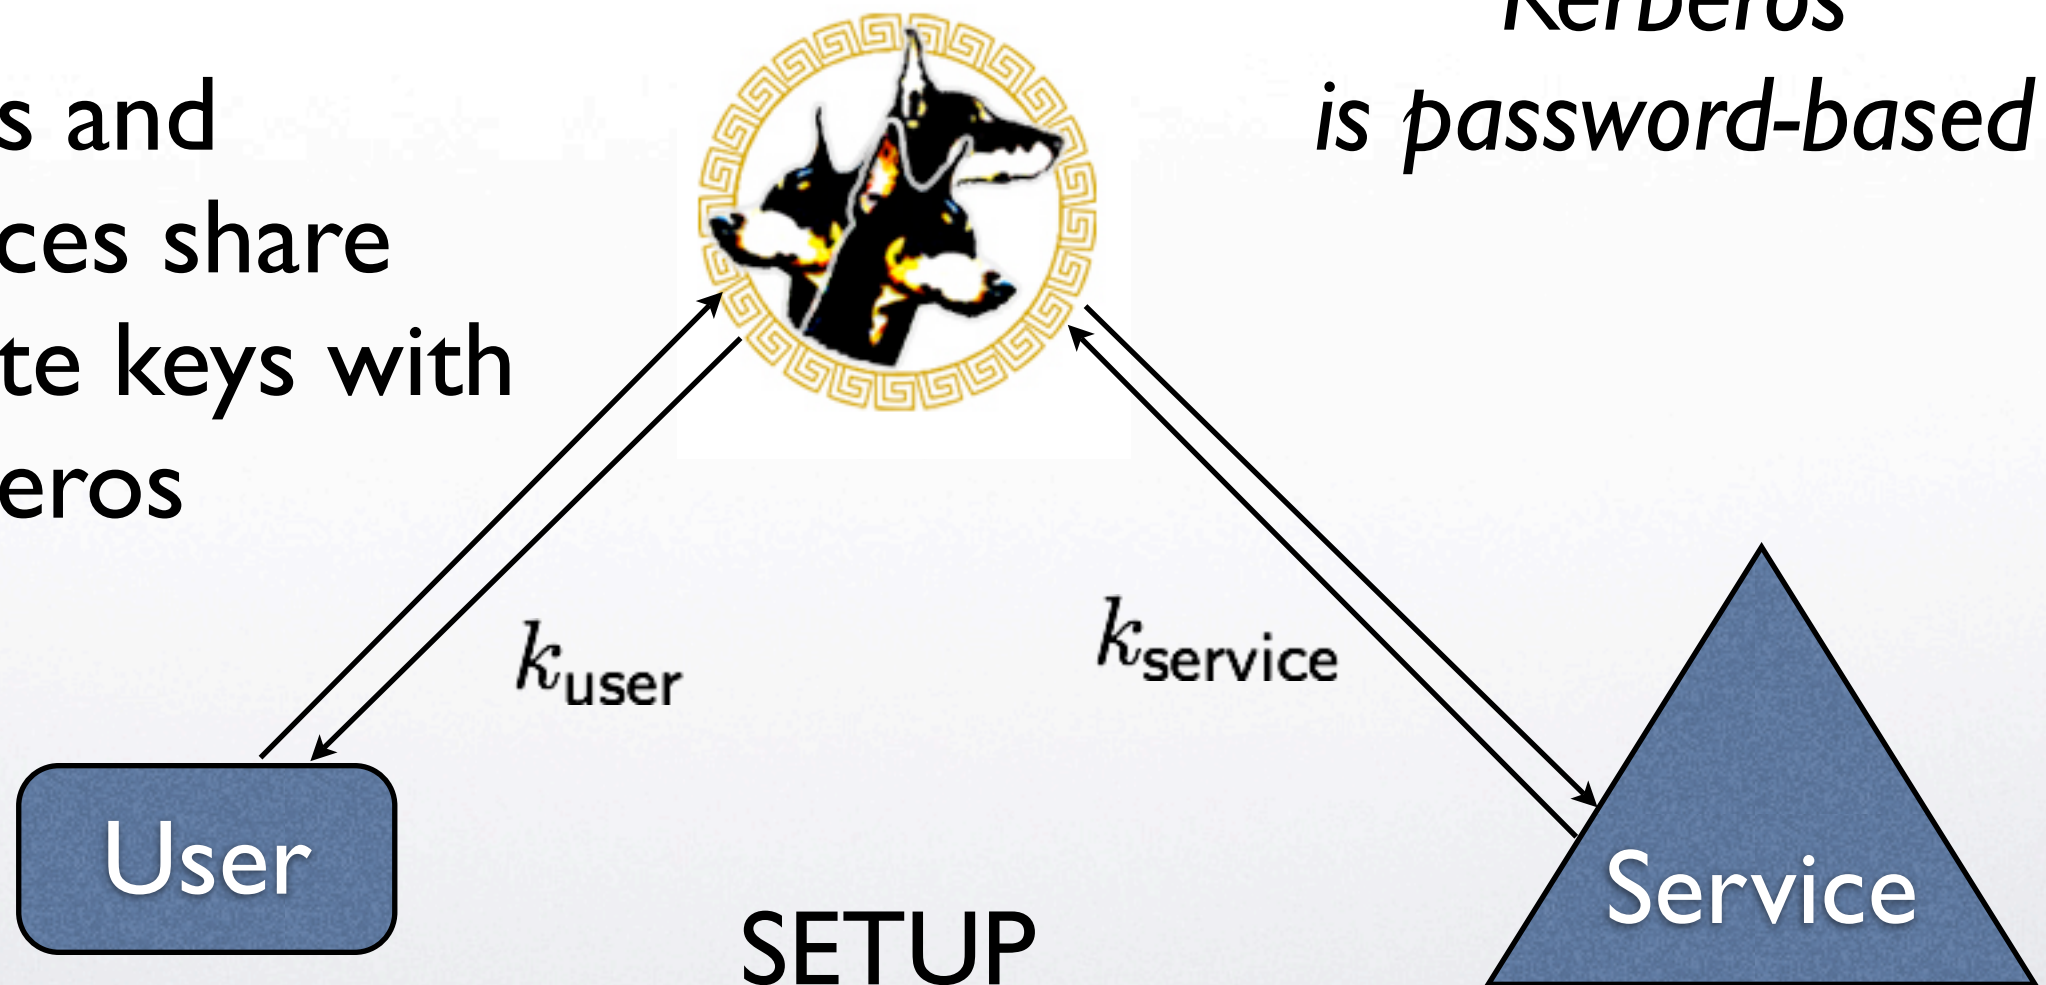




The Kerberos Approach

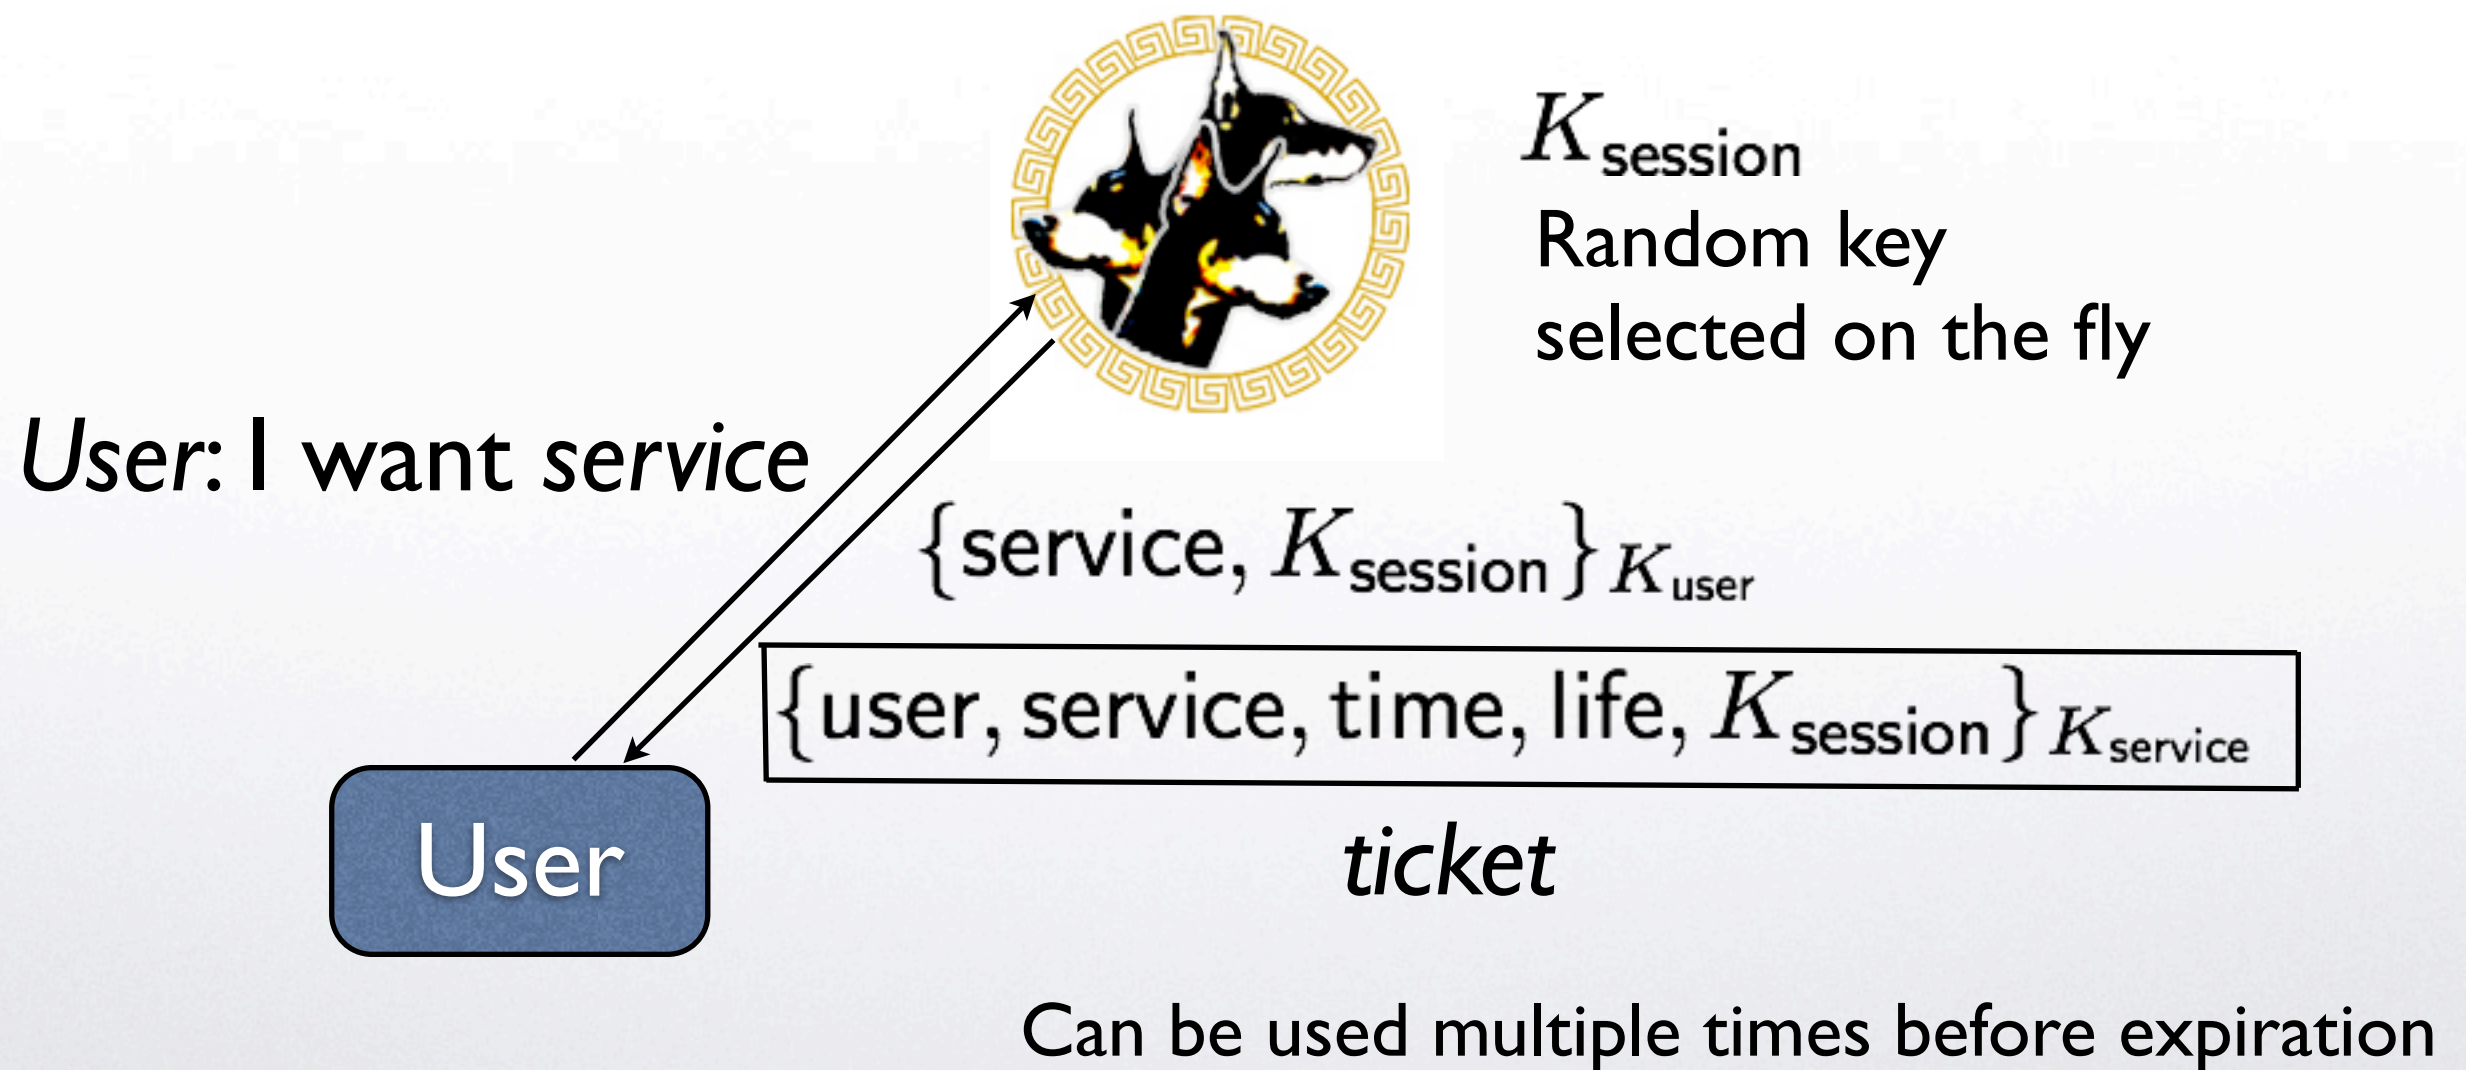
Users and services share private keys with Kerberos

Kerberos is password-based





Kerberos, II

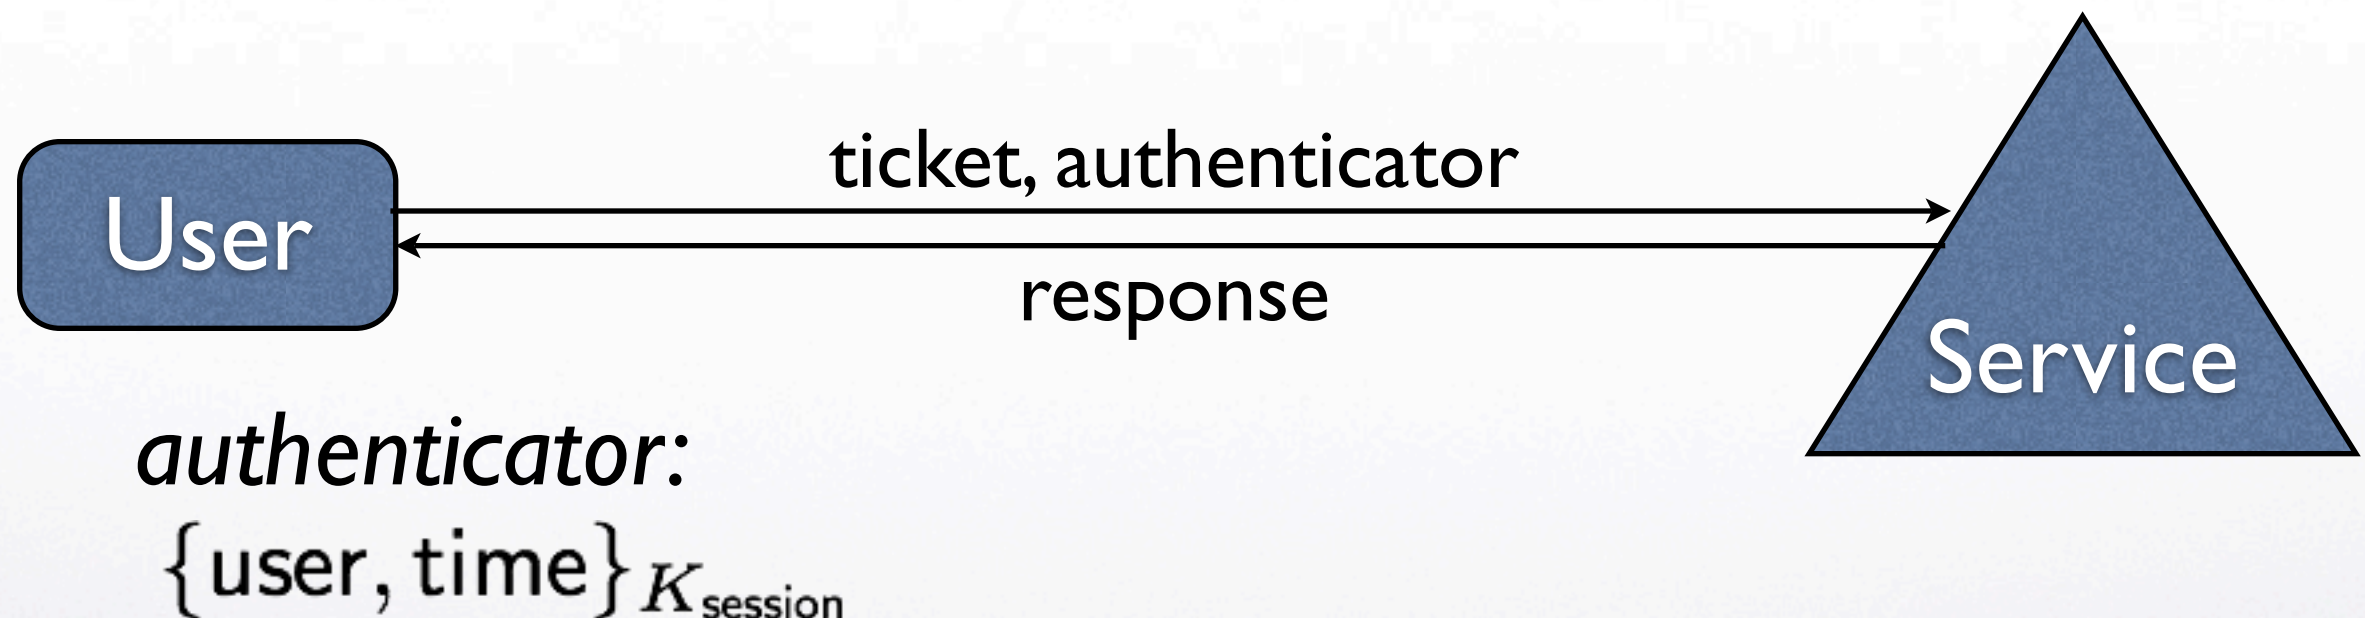




Kerberos, III

ticket:

$\{\text{user, service, time, life, } K_{\text{session}}\} K_{\text{service}}$



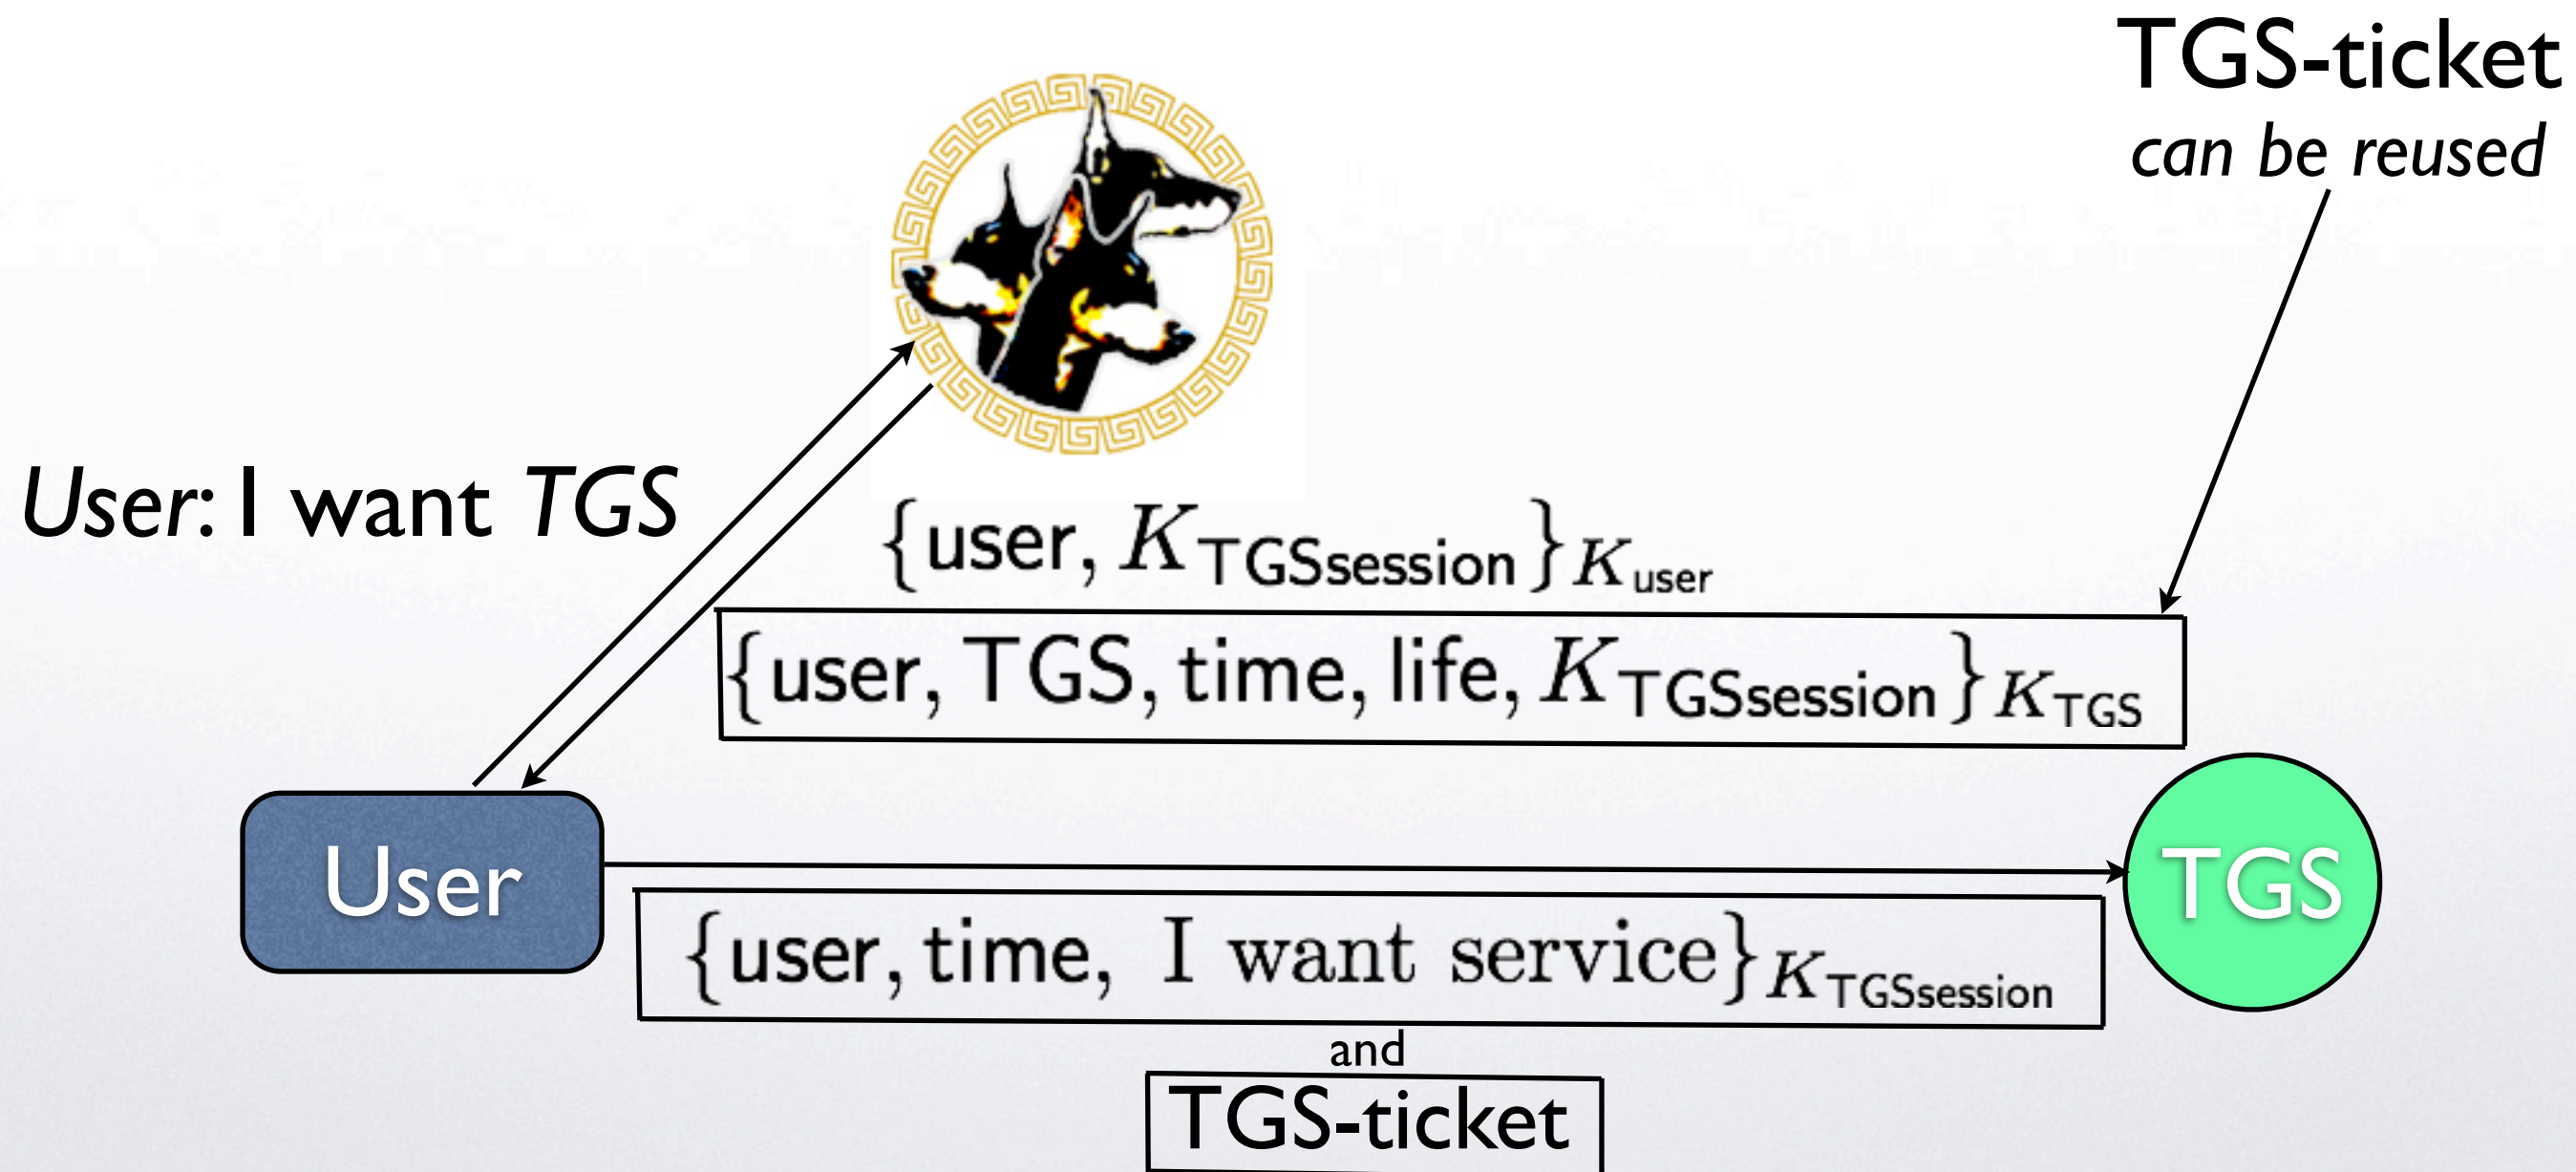


Kerberos, IV

- Above description too stressful for Kerberos.
- Easing Kerberos task:
 - Kerberos will recognize only one service, the *Ticket Granting Service*.
 - Instead of giving tickets for every service it will give tickets only for using the TGS.

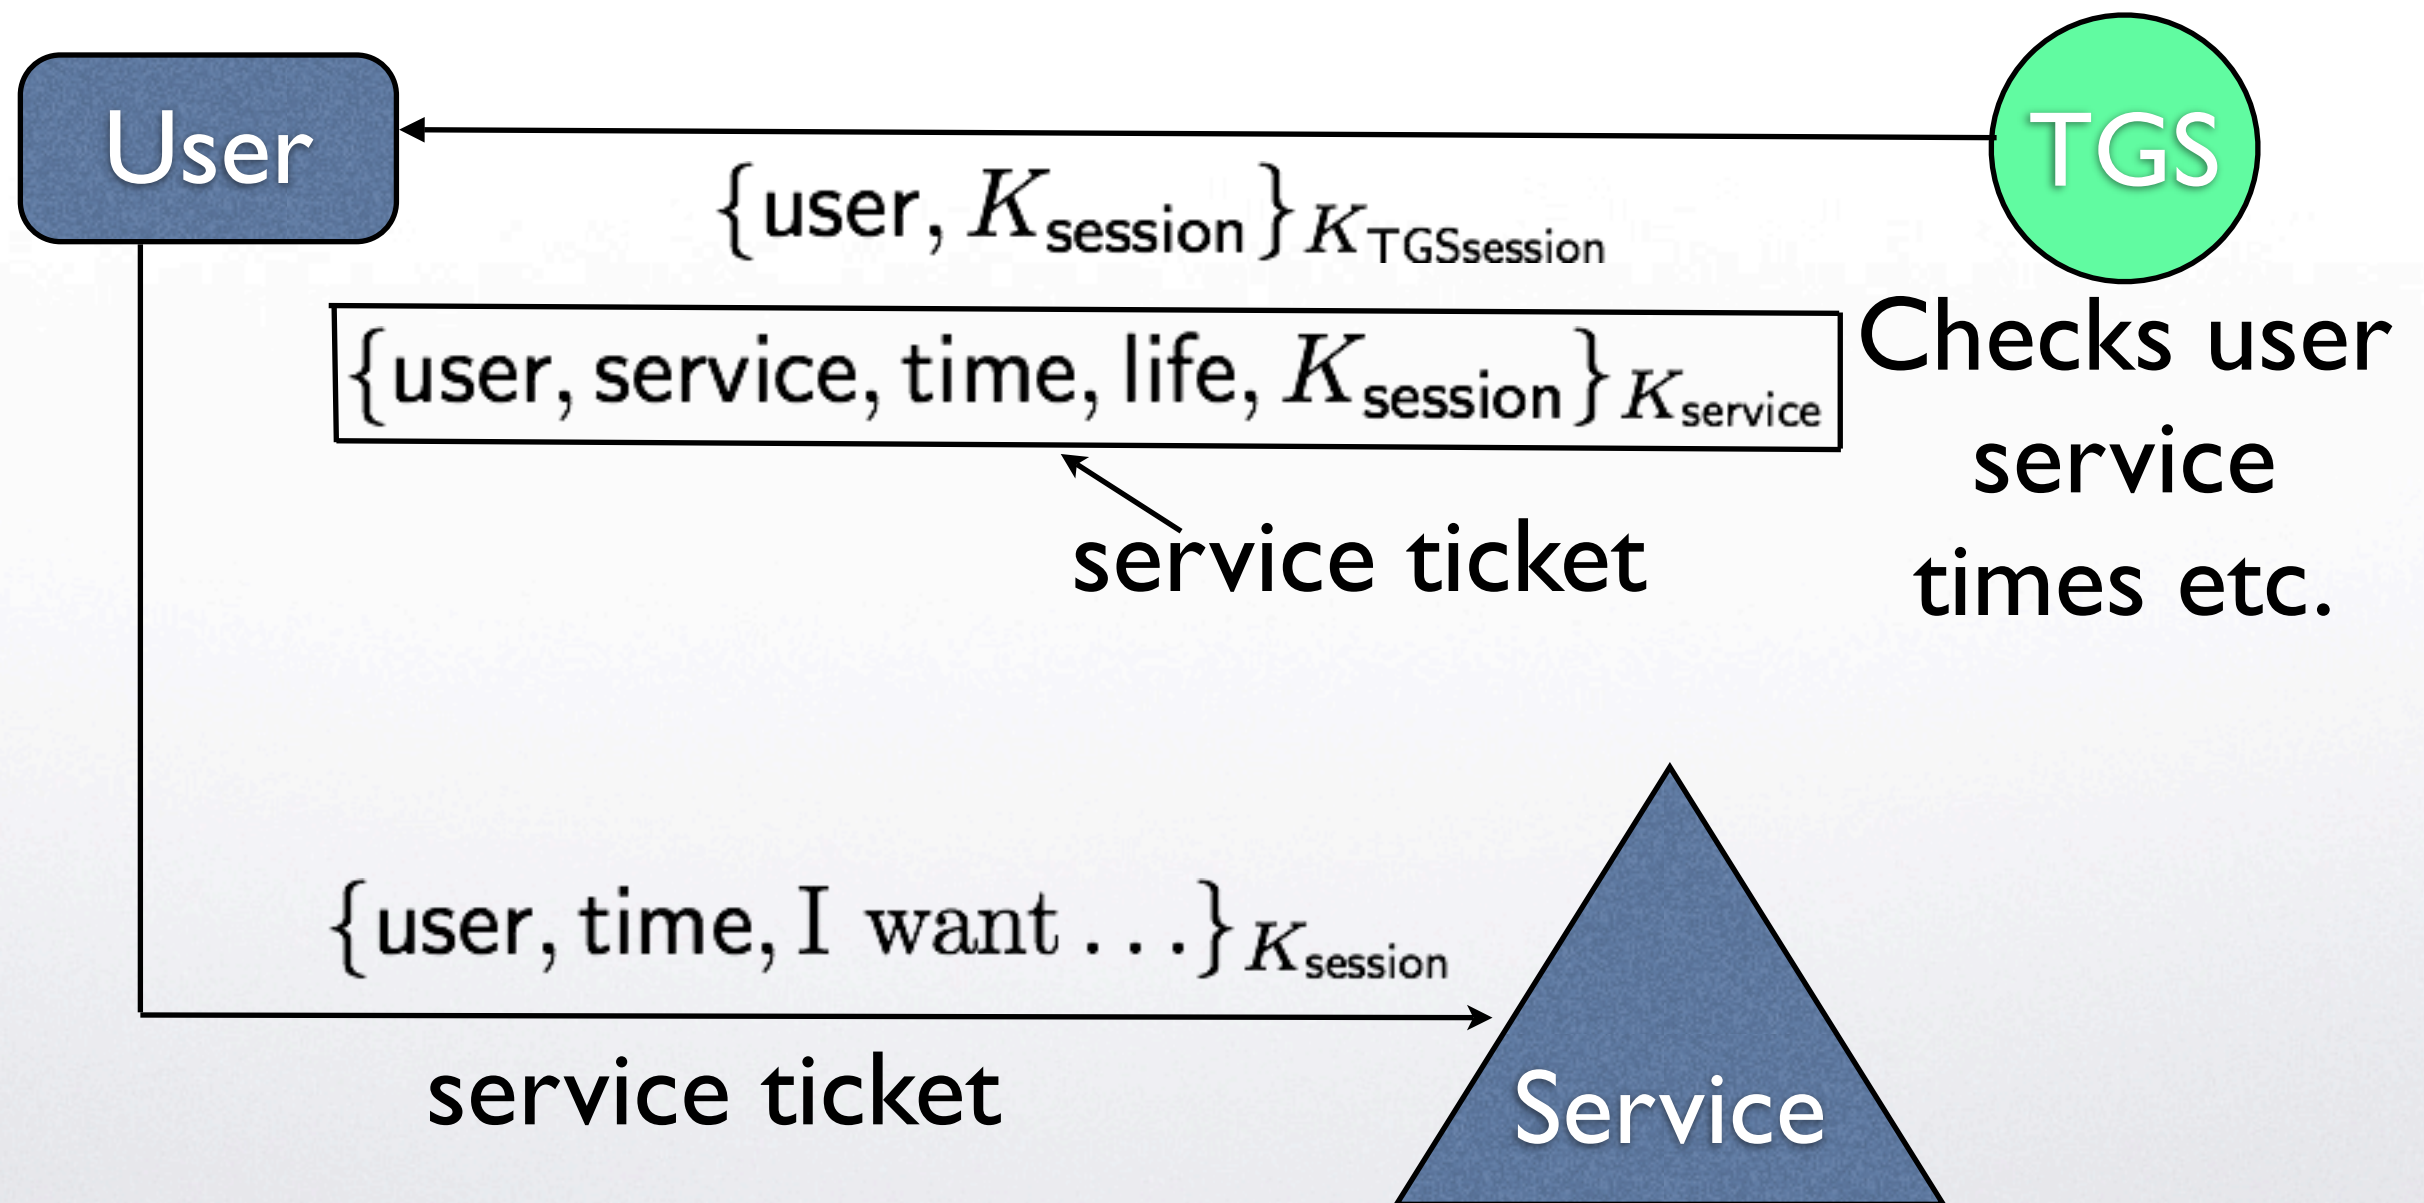


Kerberos, V





Kerberos, VI





Kerberos VII

- Kerberos server knows all user keys and the TGS key. It handles user authentication.
- Ticket Granting Server knows service keys. It handles user requests to access services.
- Kerberos does not need to know about system services. TGS does not need to worry about authenticating users.



Kerberos VIII

- Where do keys come from?
 - user keys are derived from human passwords.
 - service keys are random and stored locally. assumed to be stored securely.



Kerberos IX

- Kerberos advantages:
 - Human passwords are never communicated.
Only on the fly usage by local “login” challenge.
 - Mutual authentication between users and services.
- Kerberos disadvantages:
 - monolithic



Kerberos X

- Windows (all the way since 2000) uses Kerberos for authentication services.
- Possible to install for Linux, Unix.
- Mac-OS X has built-in Kerberos support.